



AuSRoS

Australian School of Robotic Systems

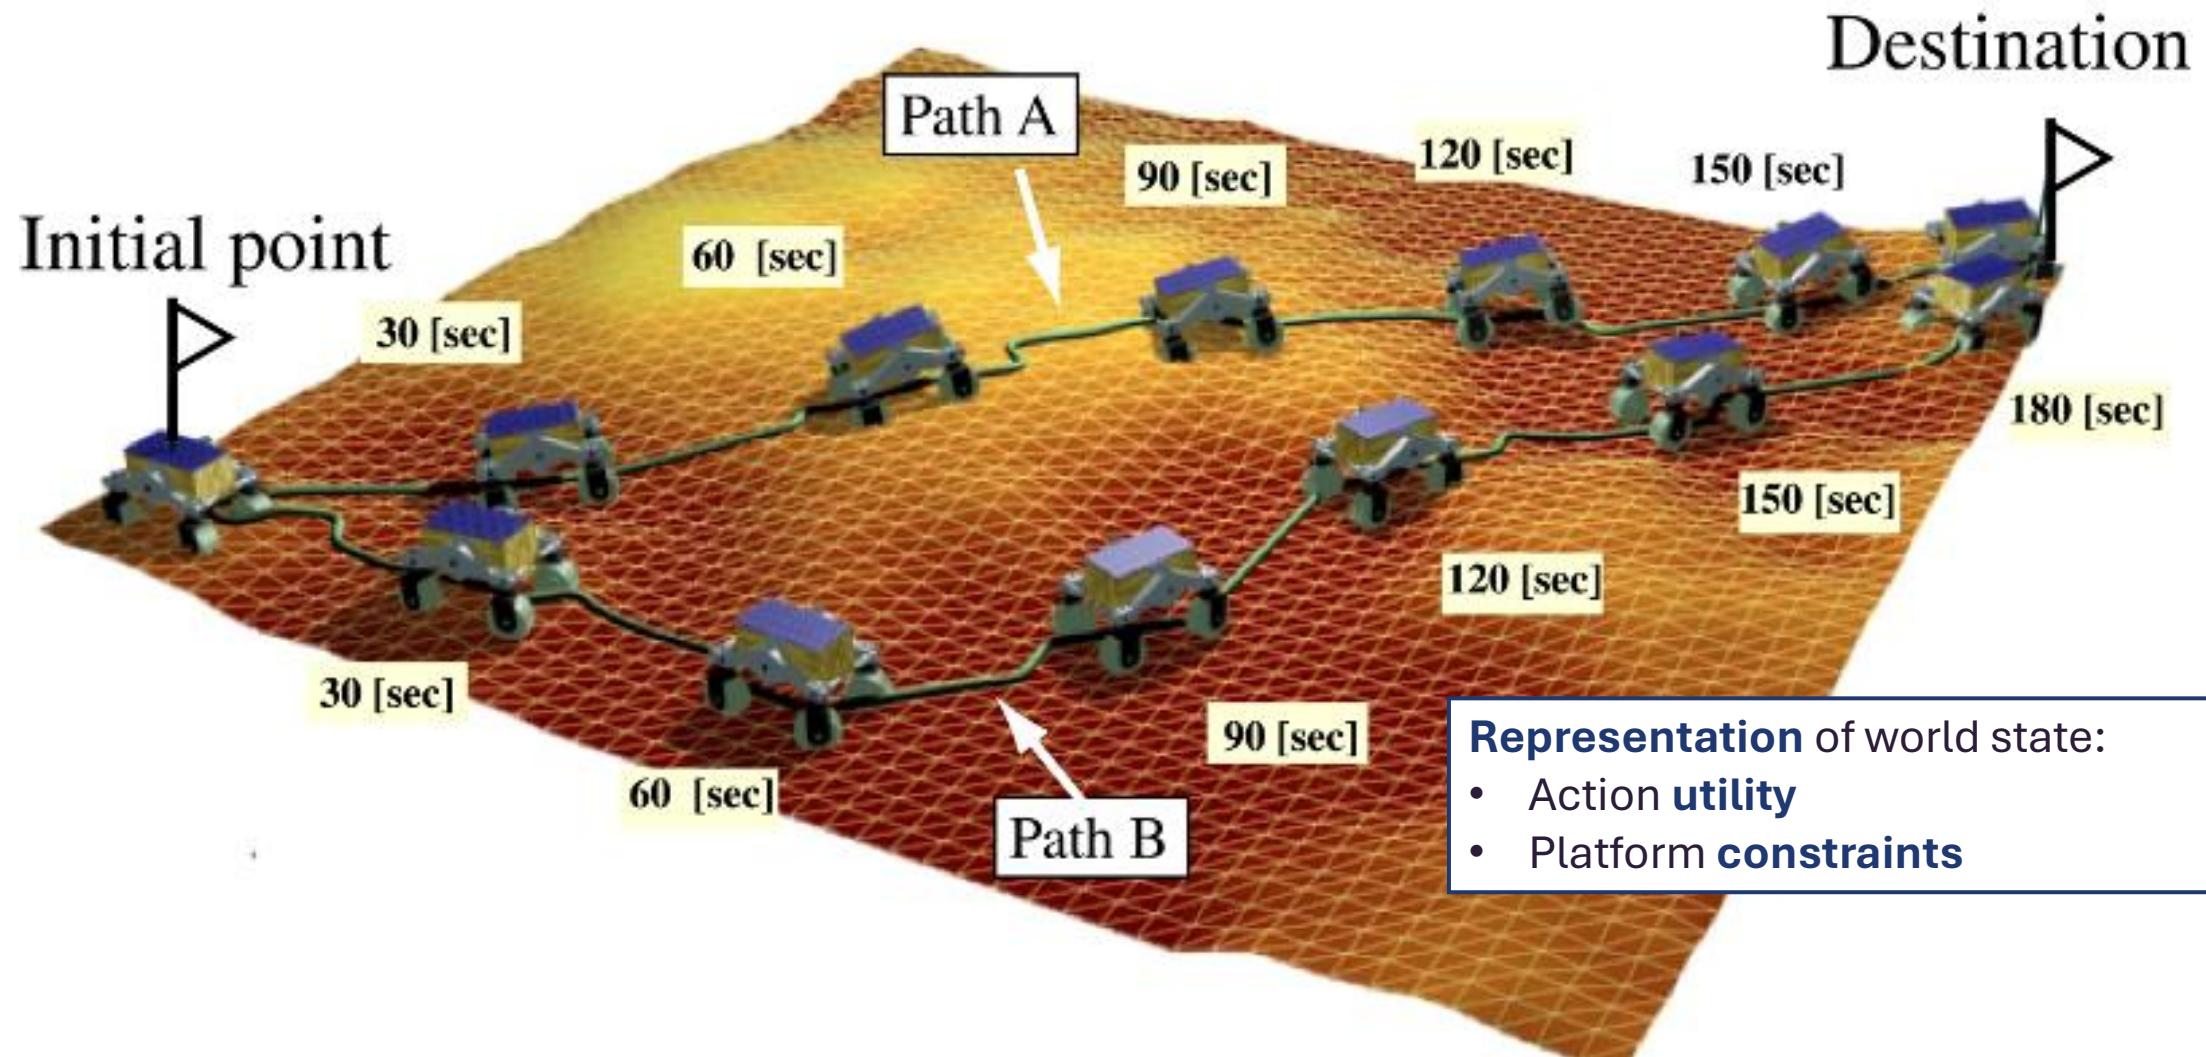
B1: Motion Planning Essentials

A/Prof. Jen Jen Chung

With slides adapted from Dr Nicholas Lawrance

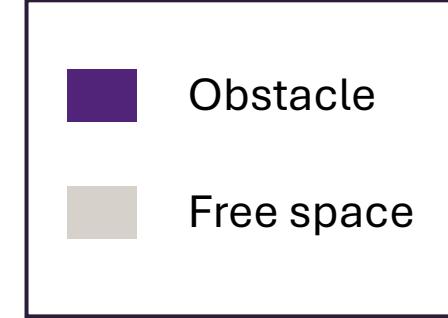
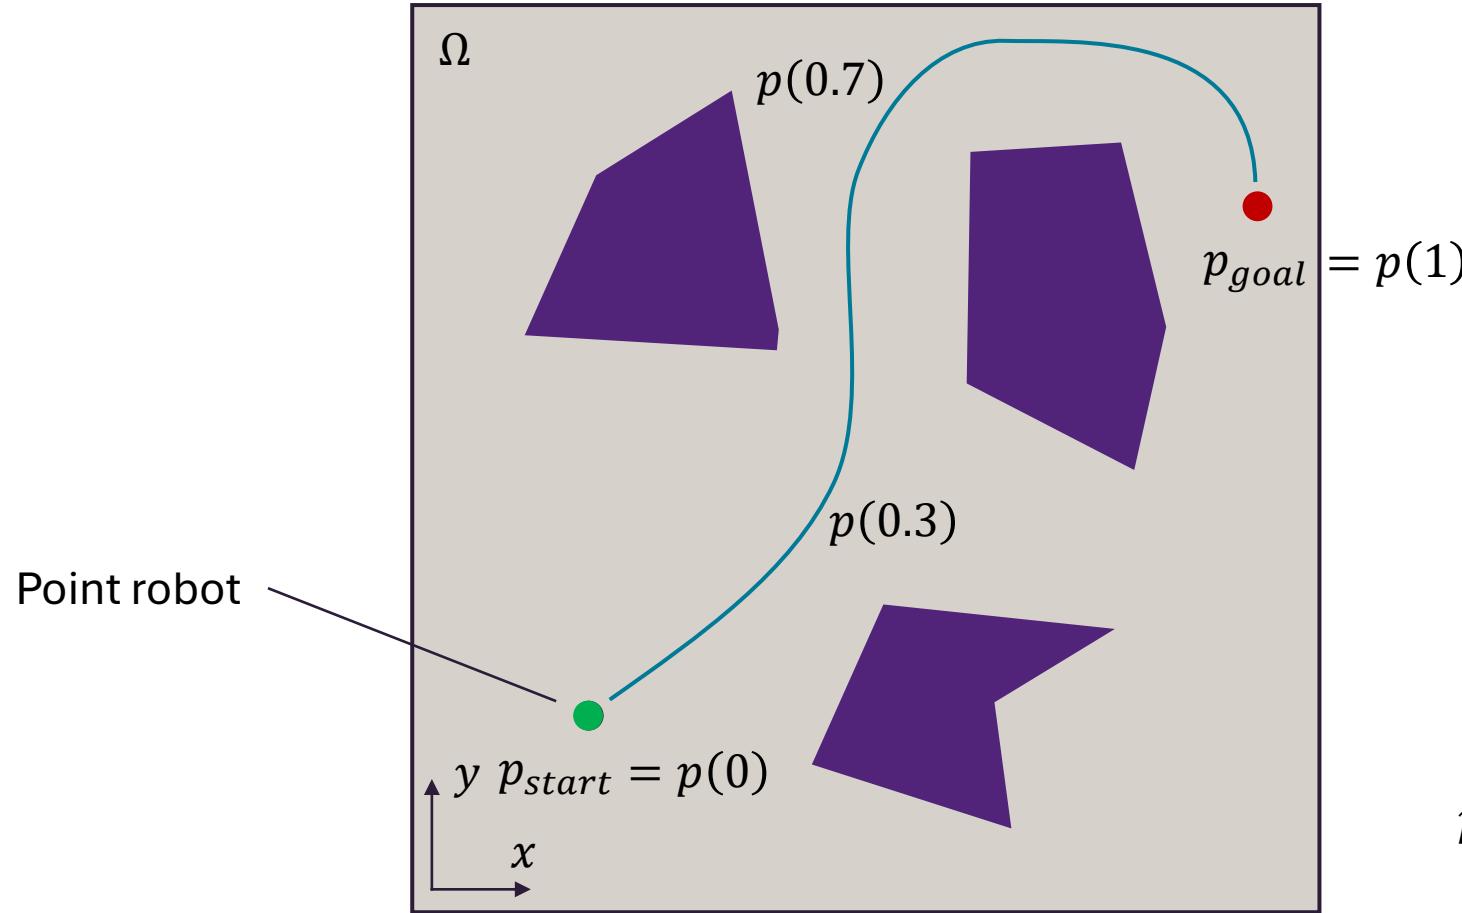


What is Robot Motion Planning?



Motion Planning Representations

Workspace and Paths



$$\omega_{free} = \Omega - \omega_{obs}$$

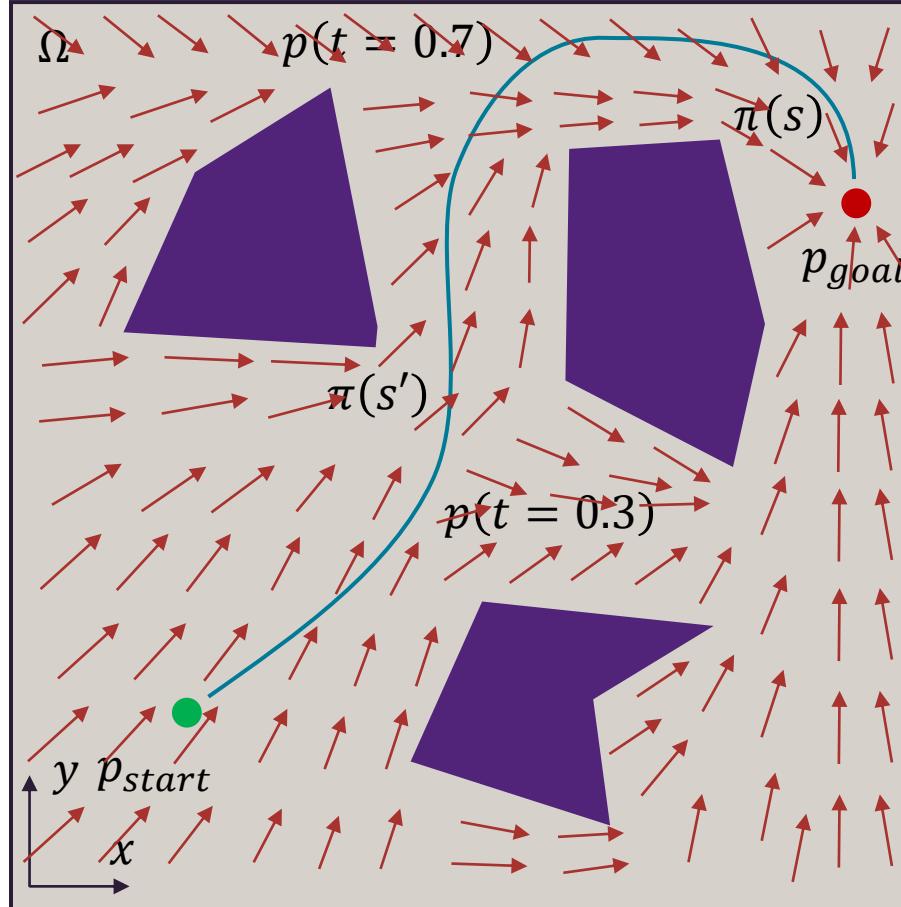
$$p: [0,1] \rightarrow \omega_{free}$$

$$p(0) = p_{start}$$

$$p(t): \forall t \in [0,1], \quad p(t) \in \omega_{free}$$

$$p(1) = p_{goal}$$

Paths, Trajectories and Policies

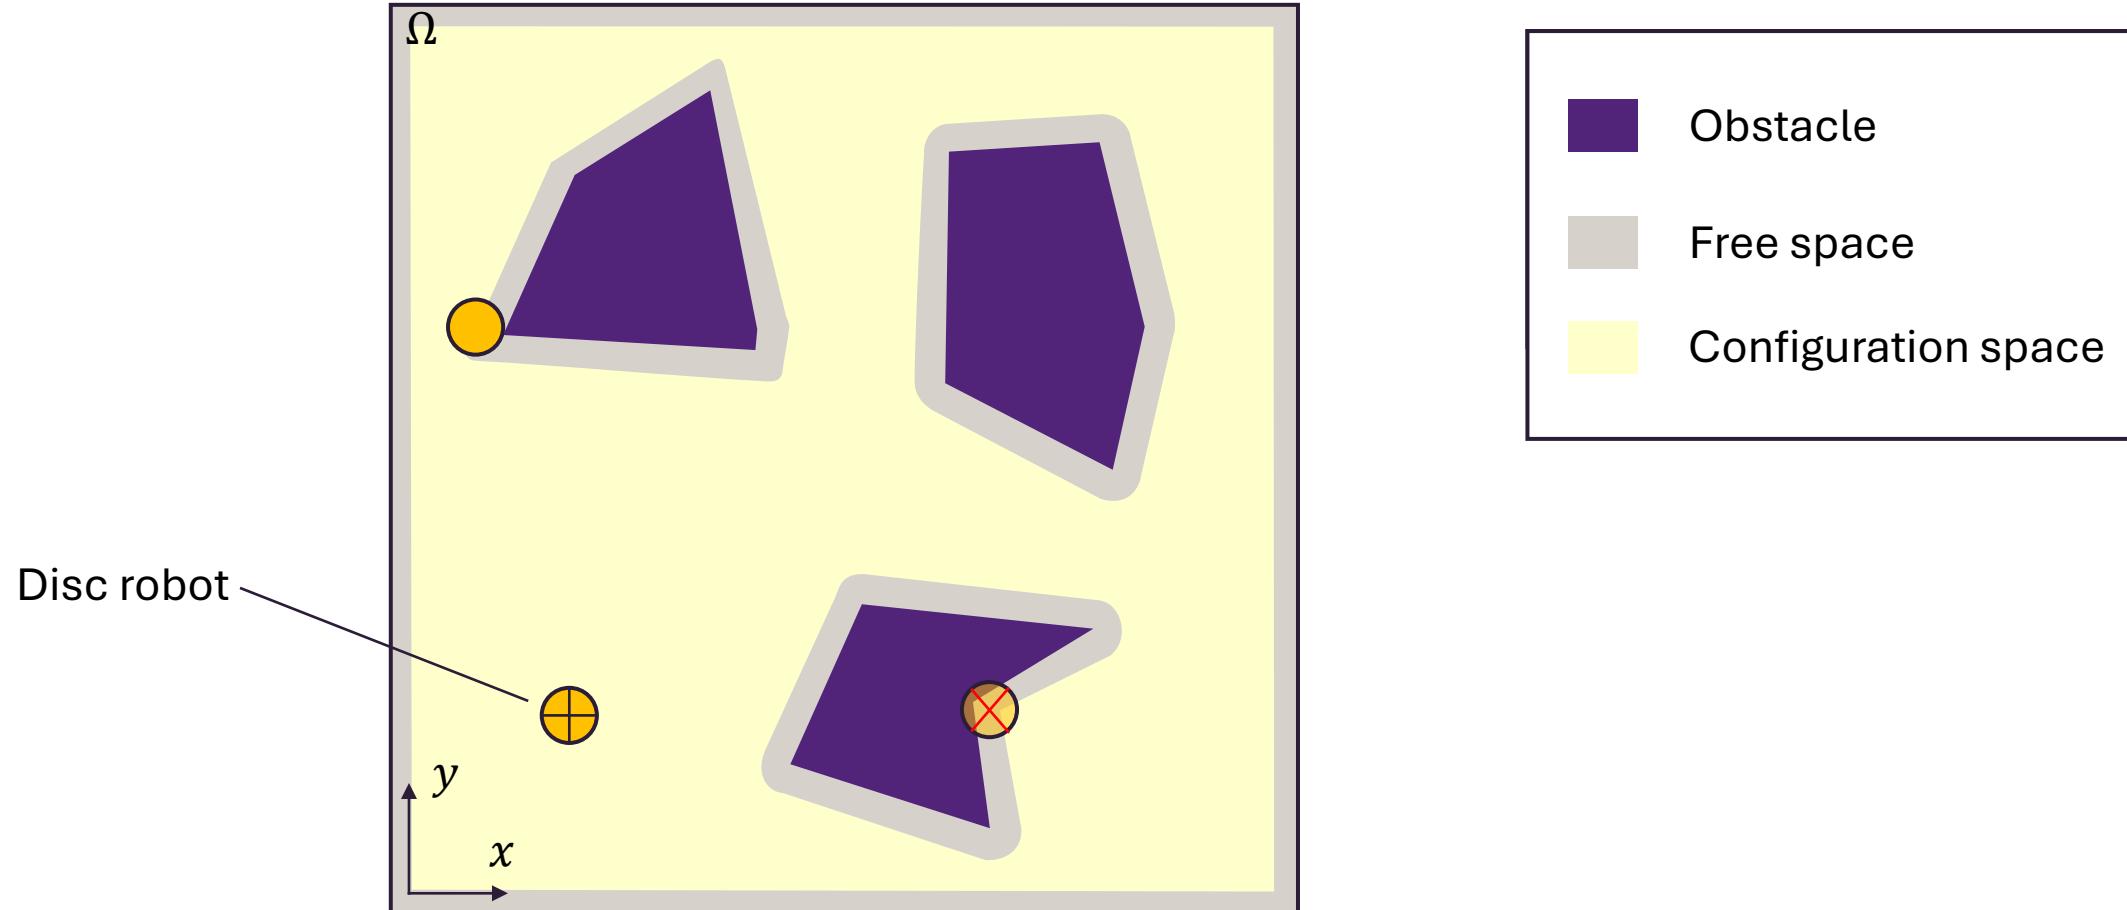


- **Path $p \subset \Omega$:** A set of points (geometric path) that the vehicle will travel along (no time consideration)
- **Trajectory $p(t): T \rightarrow \Omega$:** Path parameterised by **time** – important when considering dynamics constraints (e.g. turning radius as a function of speed)
- **Policy $\pi(s): S \rightarrow A$:** A function that maps (all valid) states to actions

Workspace vs. Configuration Space

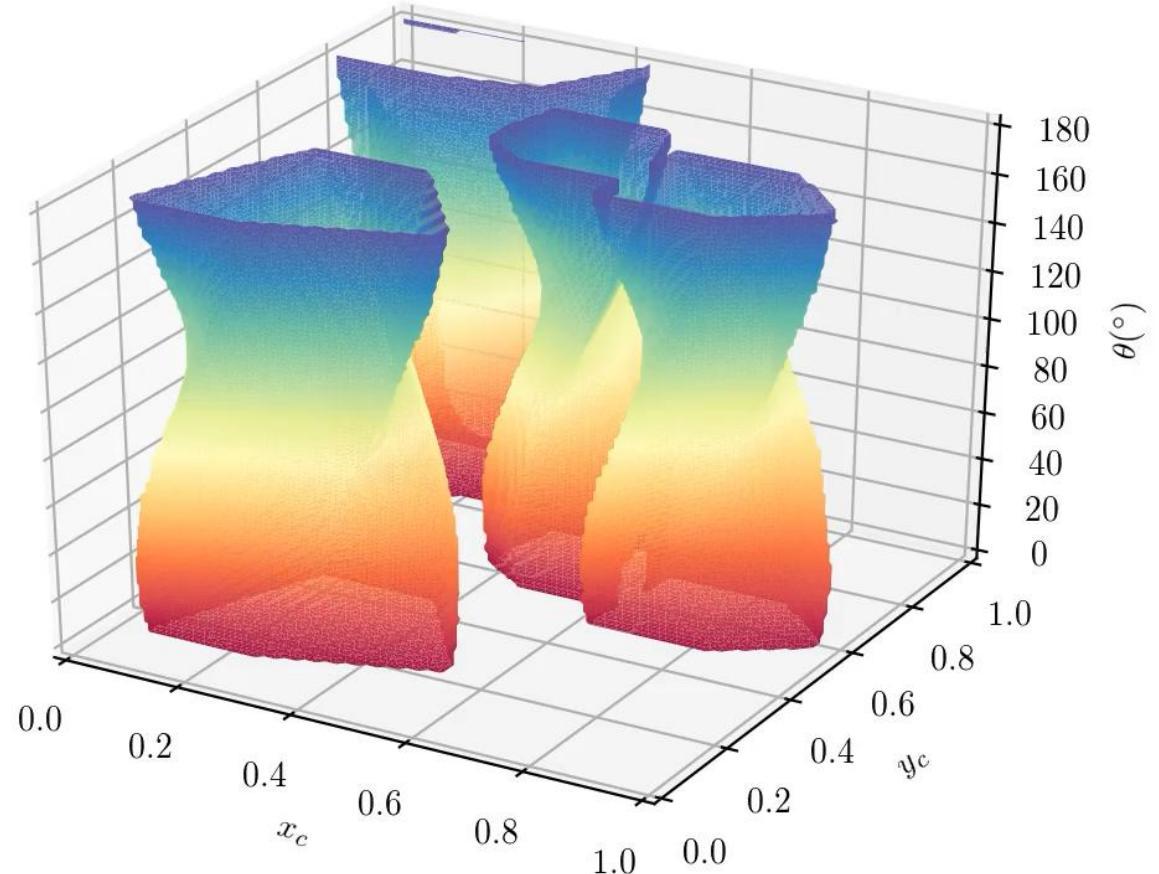
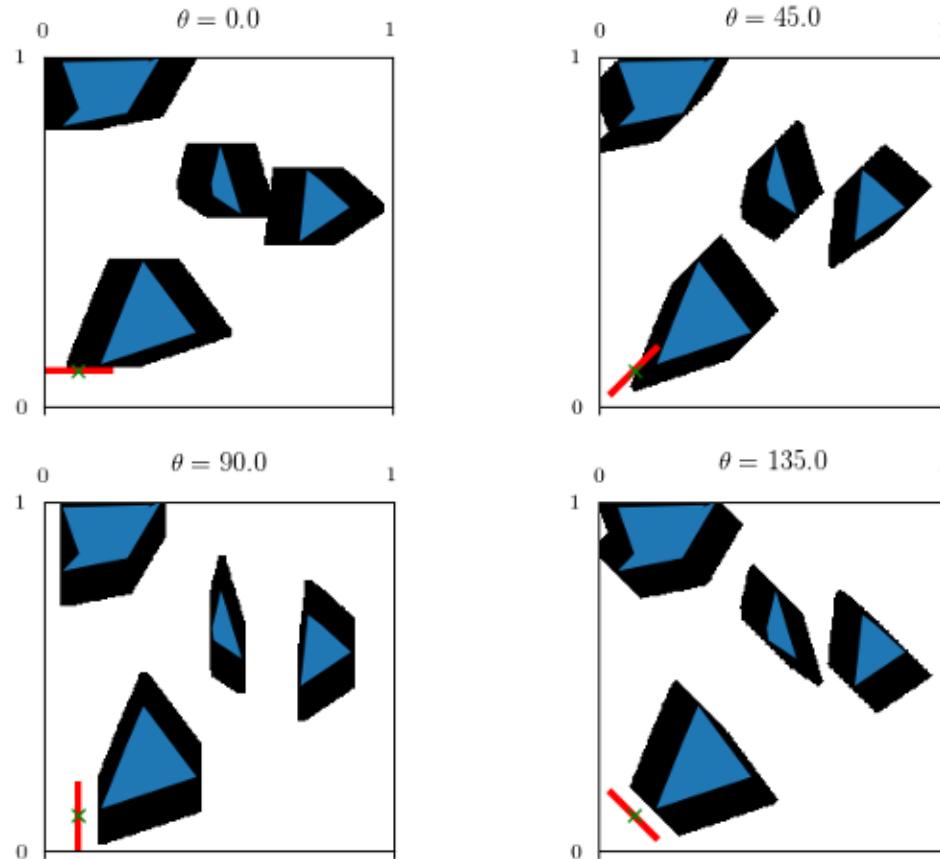
- **Workspace** is often the representation of the world, possibly independent of the robot itself. Often describes some notion of reachability, free vs. occupied
- **Configuration** space describes the full state of the robot in the world (actuator positions, orientation, etc.)

Configuration Space



Configuration Space

A robot without rotational symmetry (x, y, θ)



Configuration Space for Alternative Morphologies

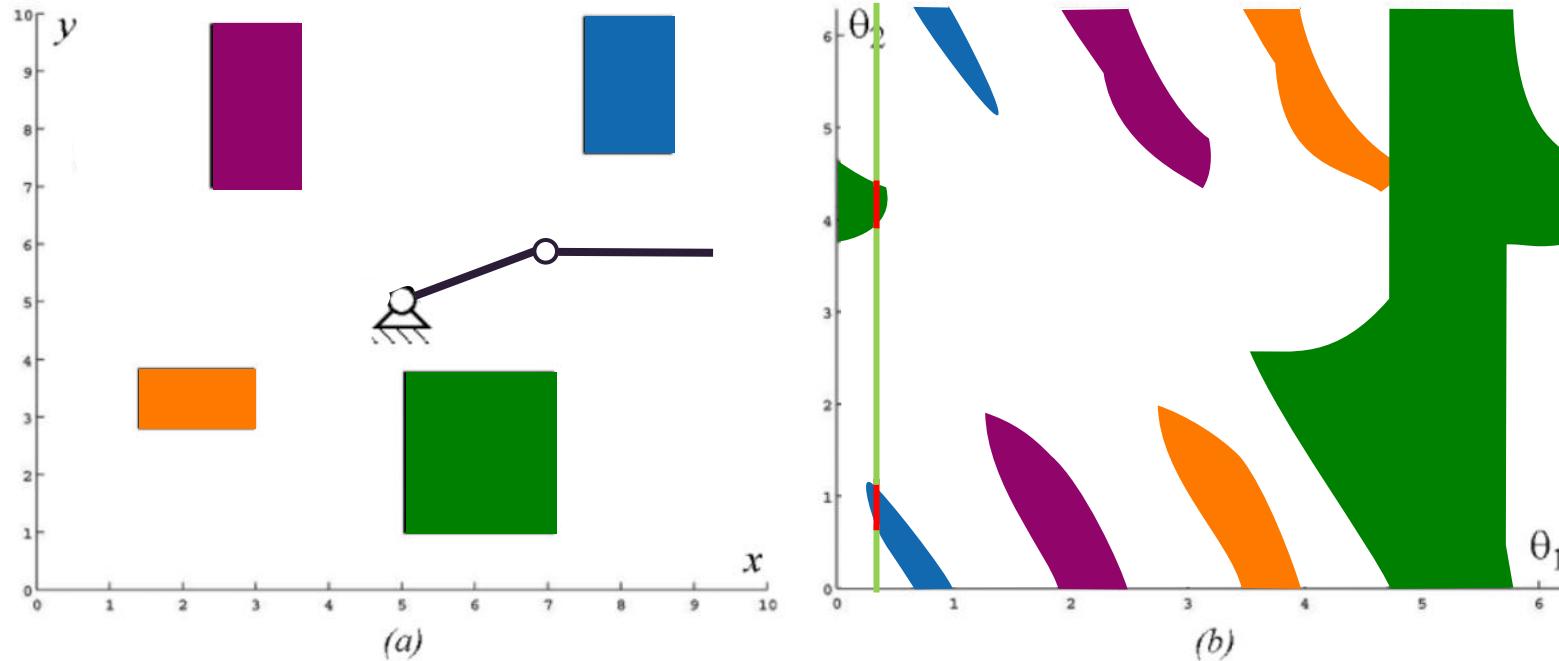
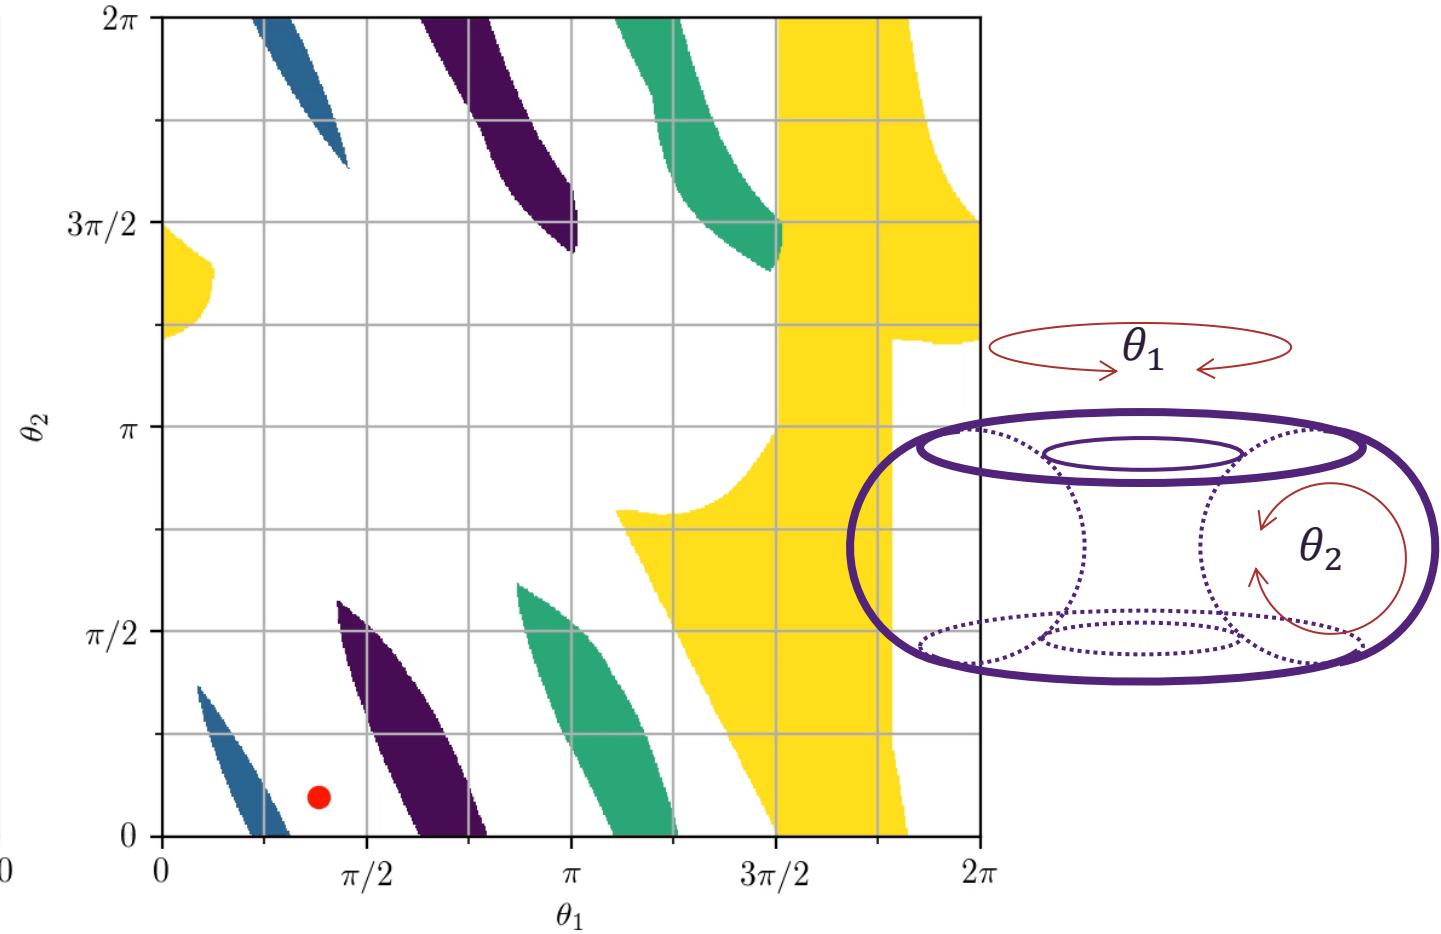
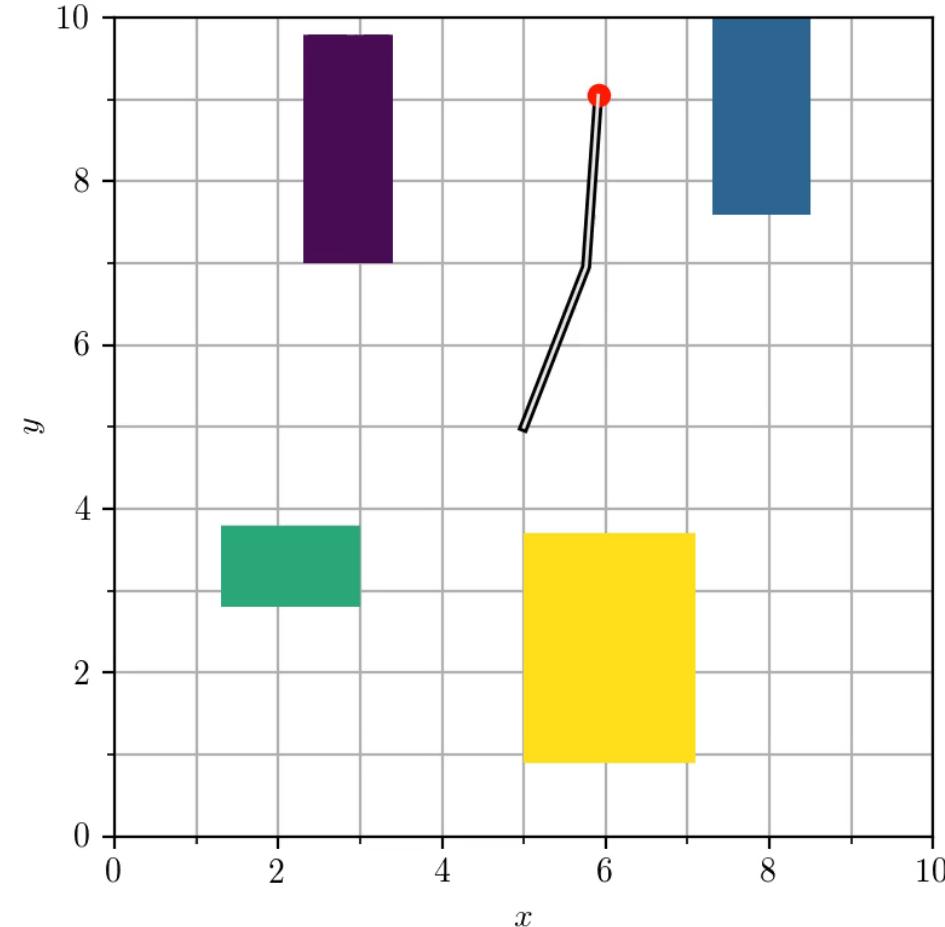


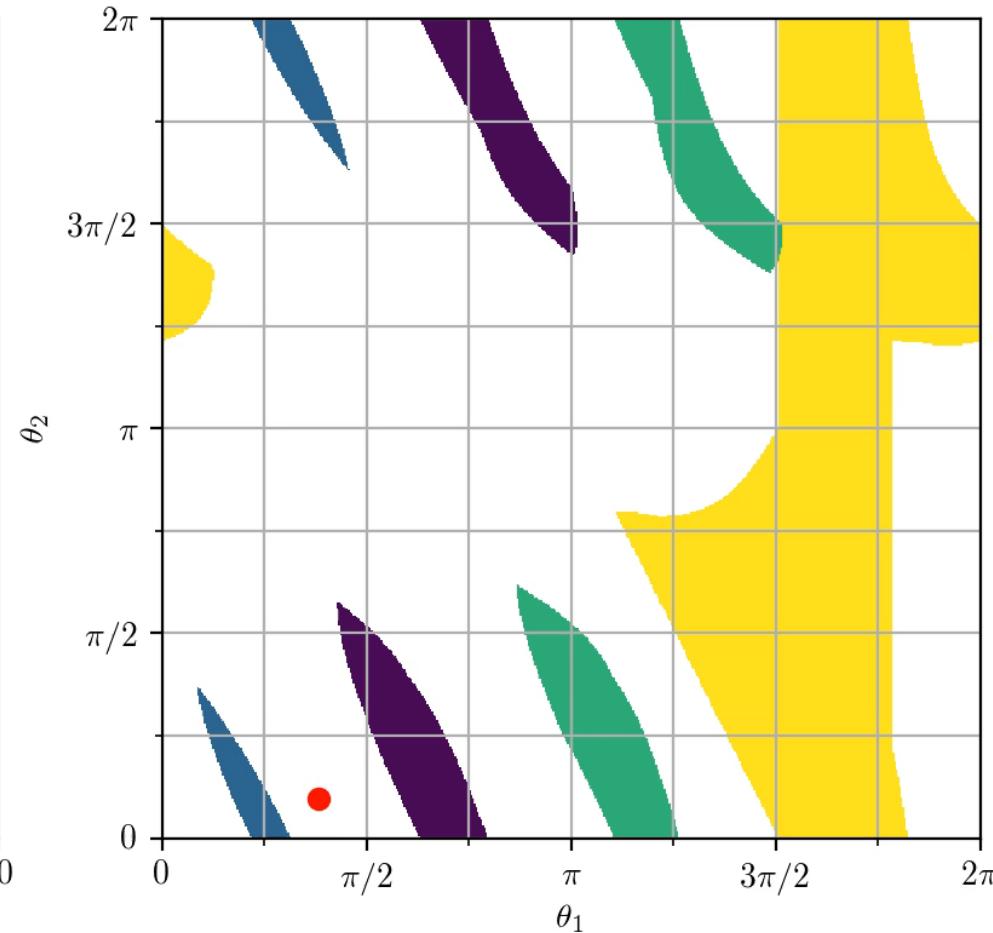
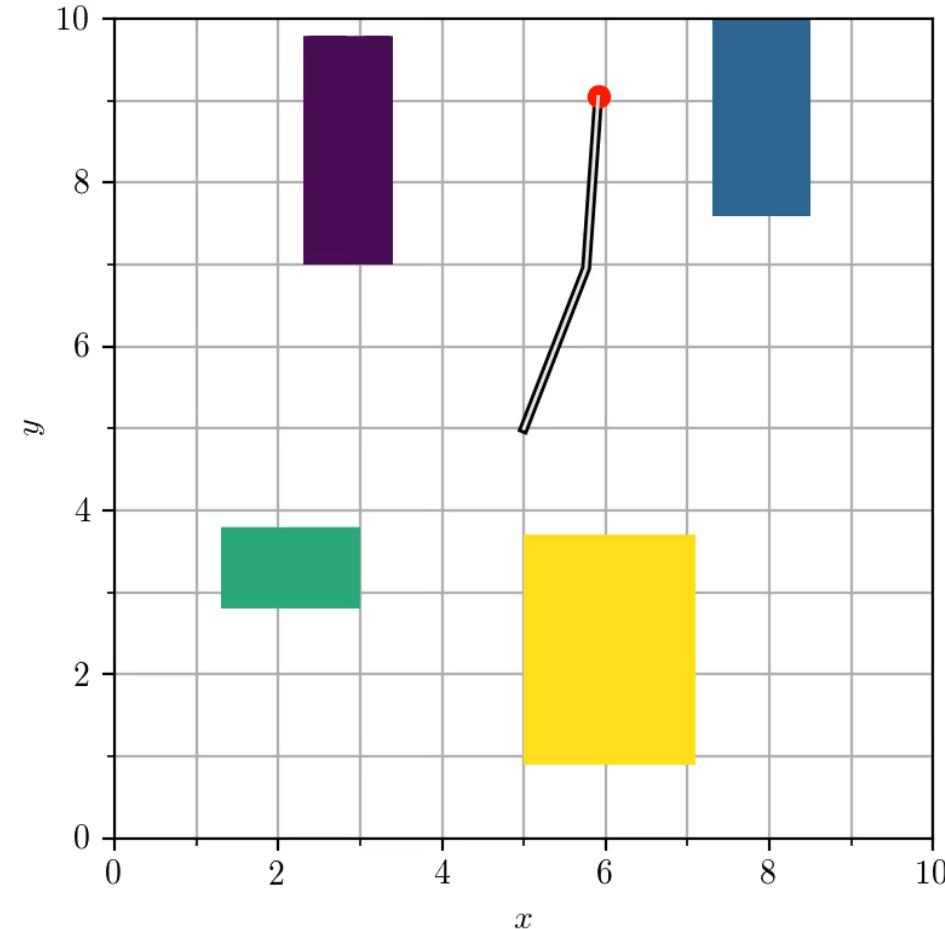
Figure 6.1

Physical space (a) and configuration space (b): (a) A two-link planar robot arm has to move from the configuration *start* to *end*. The motion is thereby constrained by the obstacles 1 to 4. (b) The corresponding configuration space shows the free space in joint coordinates (angle θ_1 and θ_2) and a path that achieves the goal.

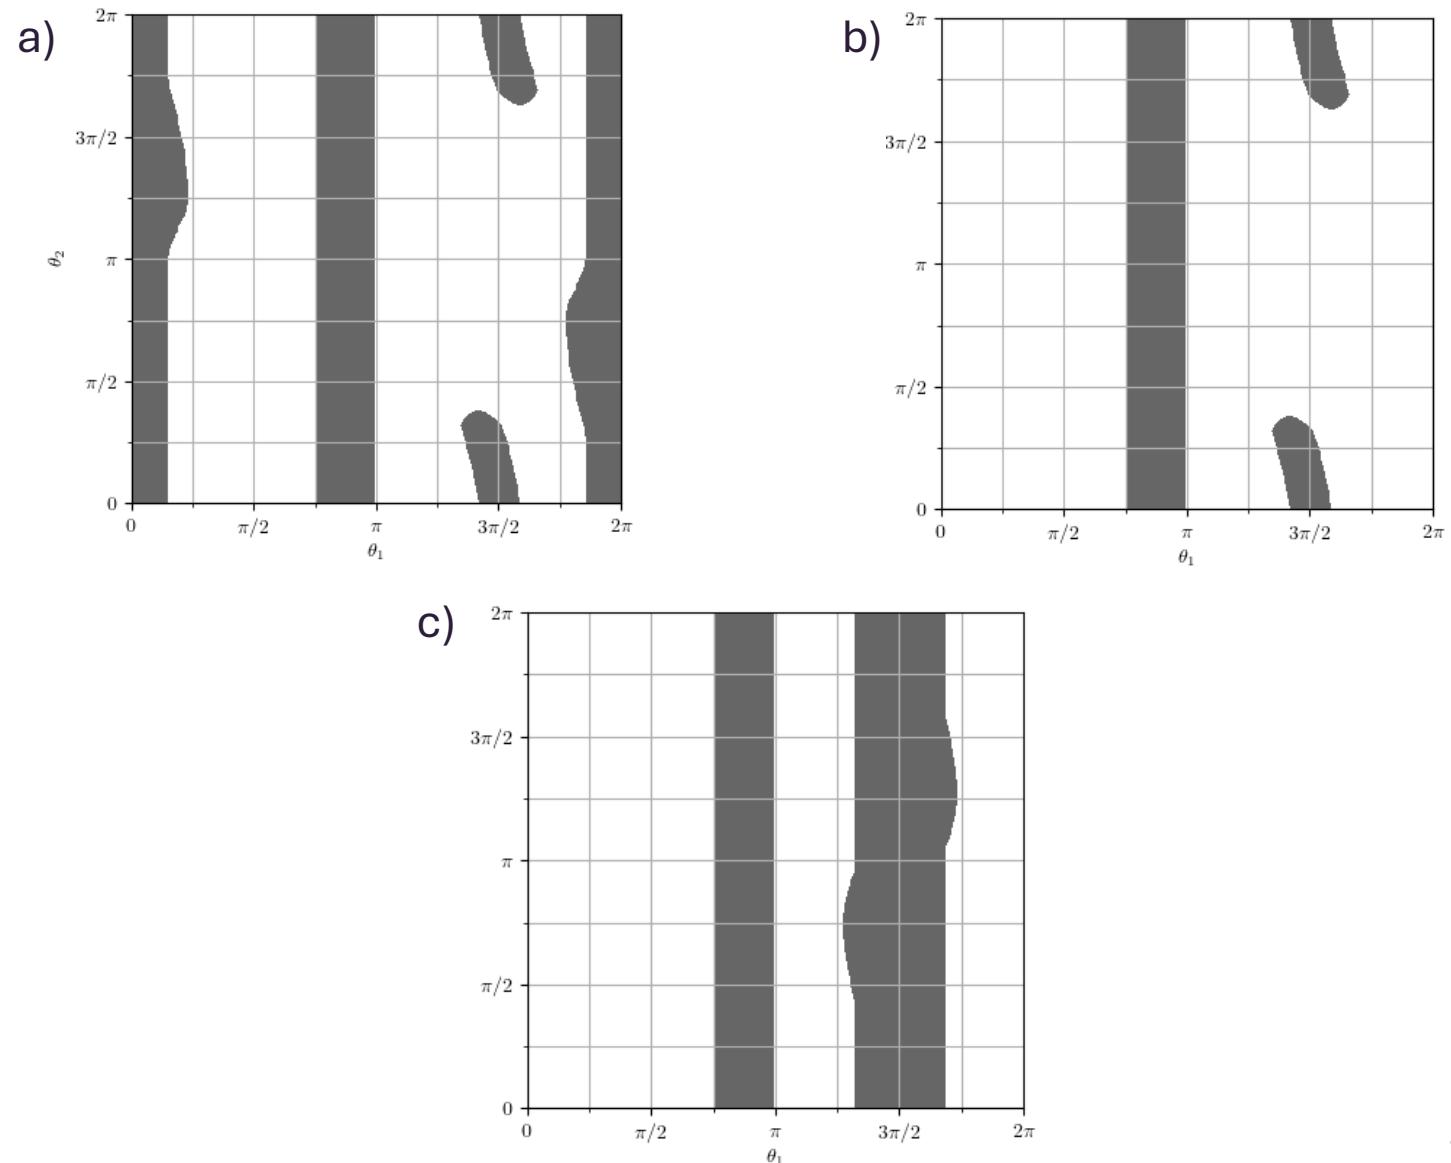
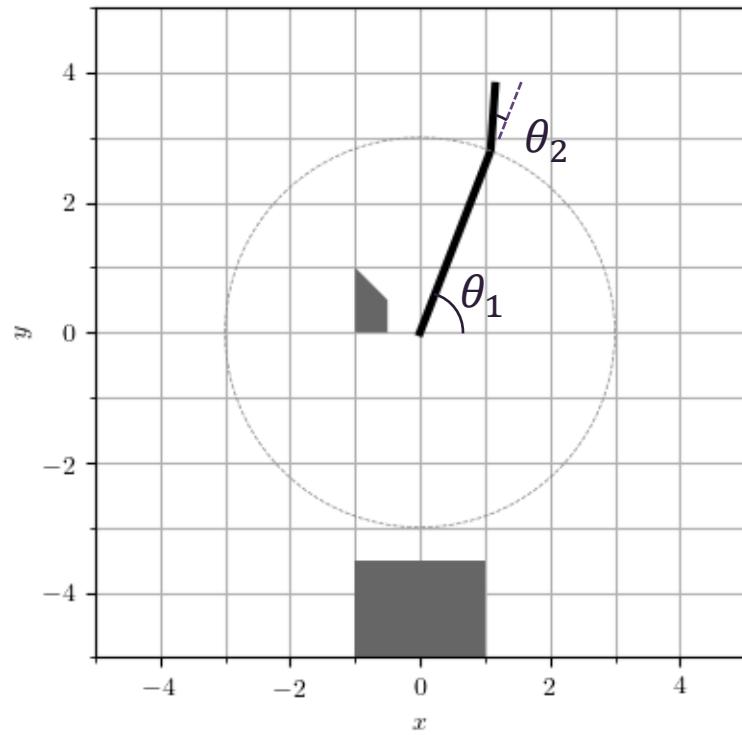
Configuration Space for Alternative Morphologies

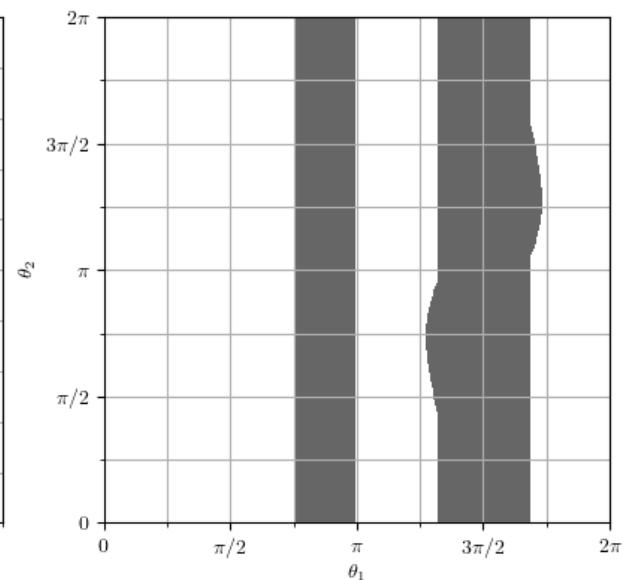
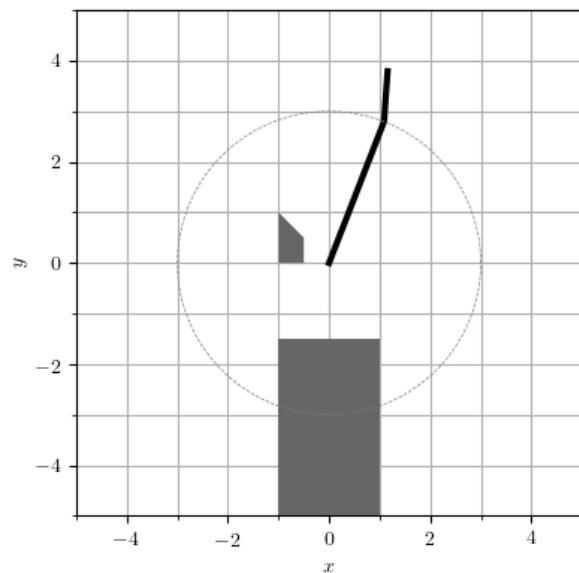
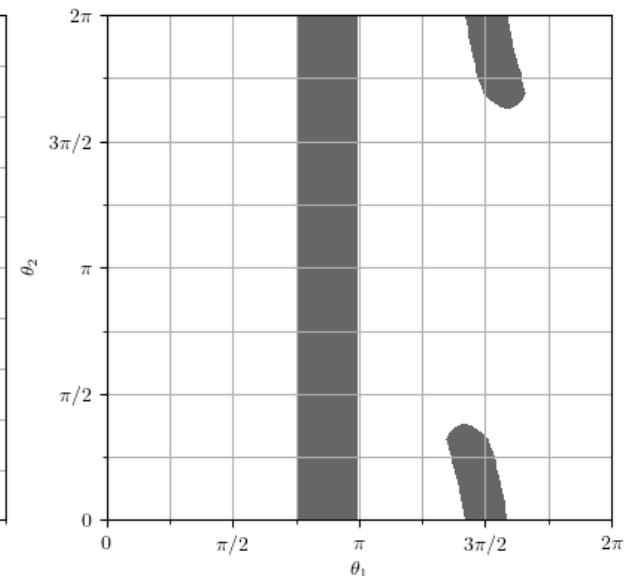
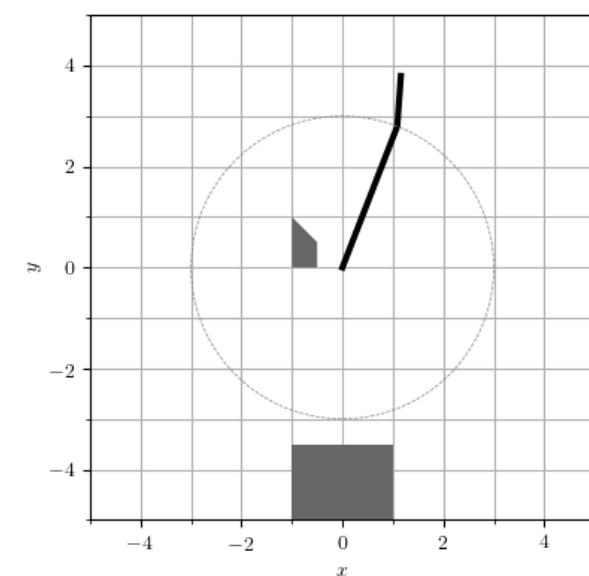
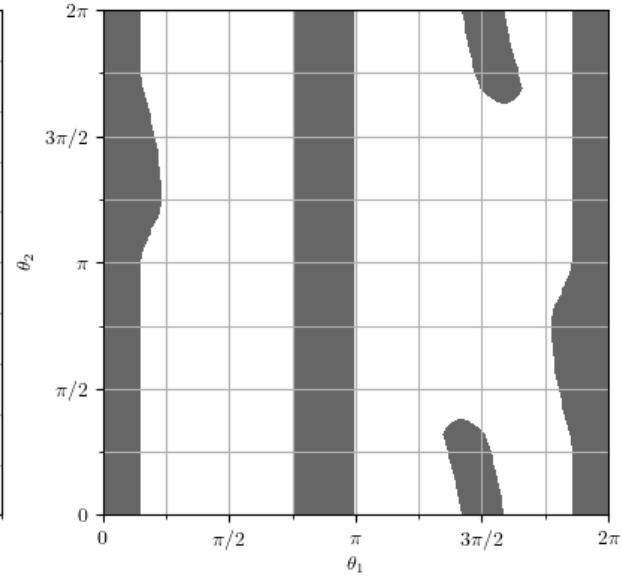
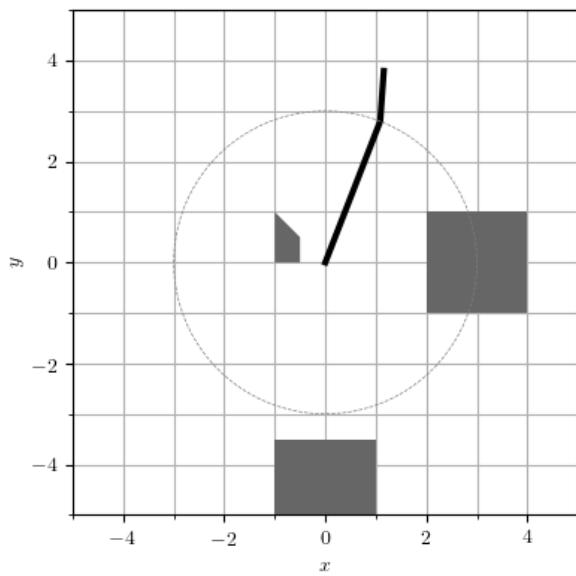


Configuration Space for Alternative Morphologies



Which is the Matching Configuration Space?





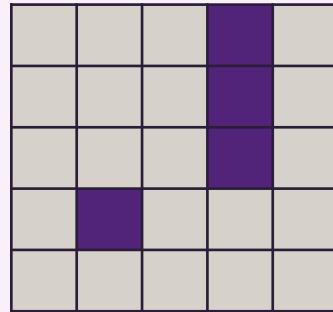
Why use Configuration Space?

- Paths in configuration space correspond to actual state changes of the robot
- Can be easier to solve collision checks and join nearby poses
- Allows a level of abstraction that means solution methods can solve a wider range of problems
- Sometimes helps with wraparound conditions (rotation joints)

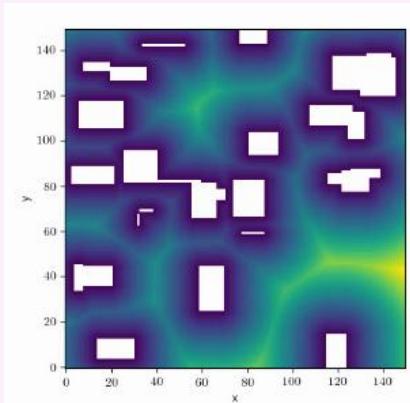
Continuous → Discrete Representations

Structured

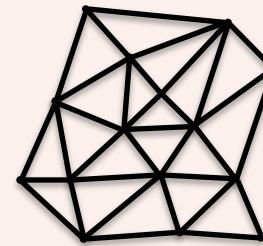
Occupancy grids



Distance fields

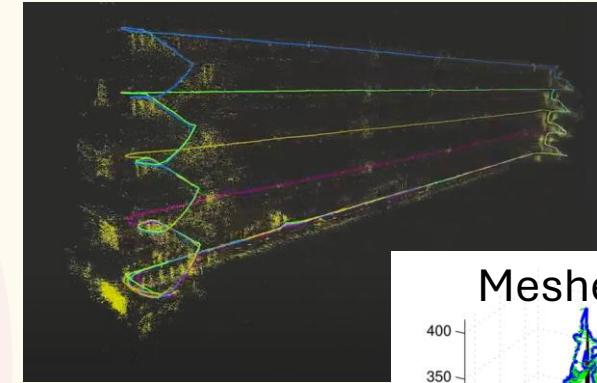


Graphs

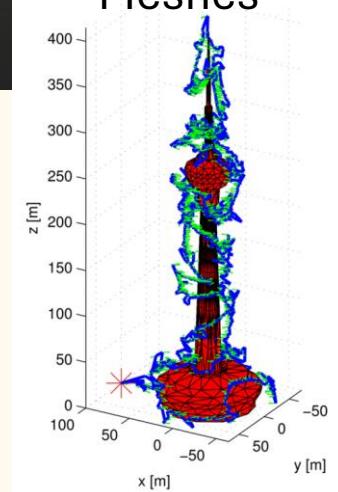


Unstructured

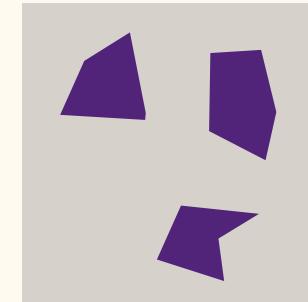
Point clouds



Meshes

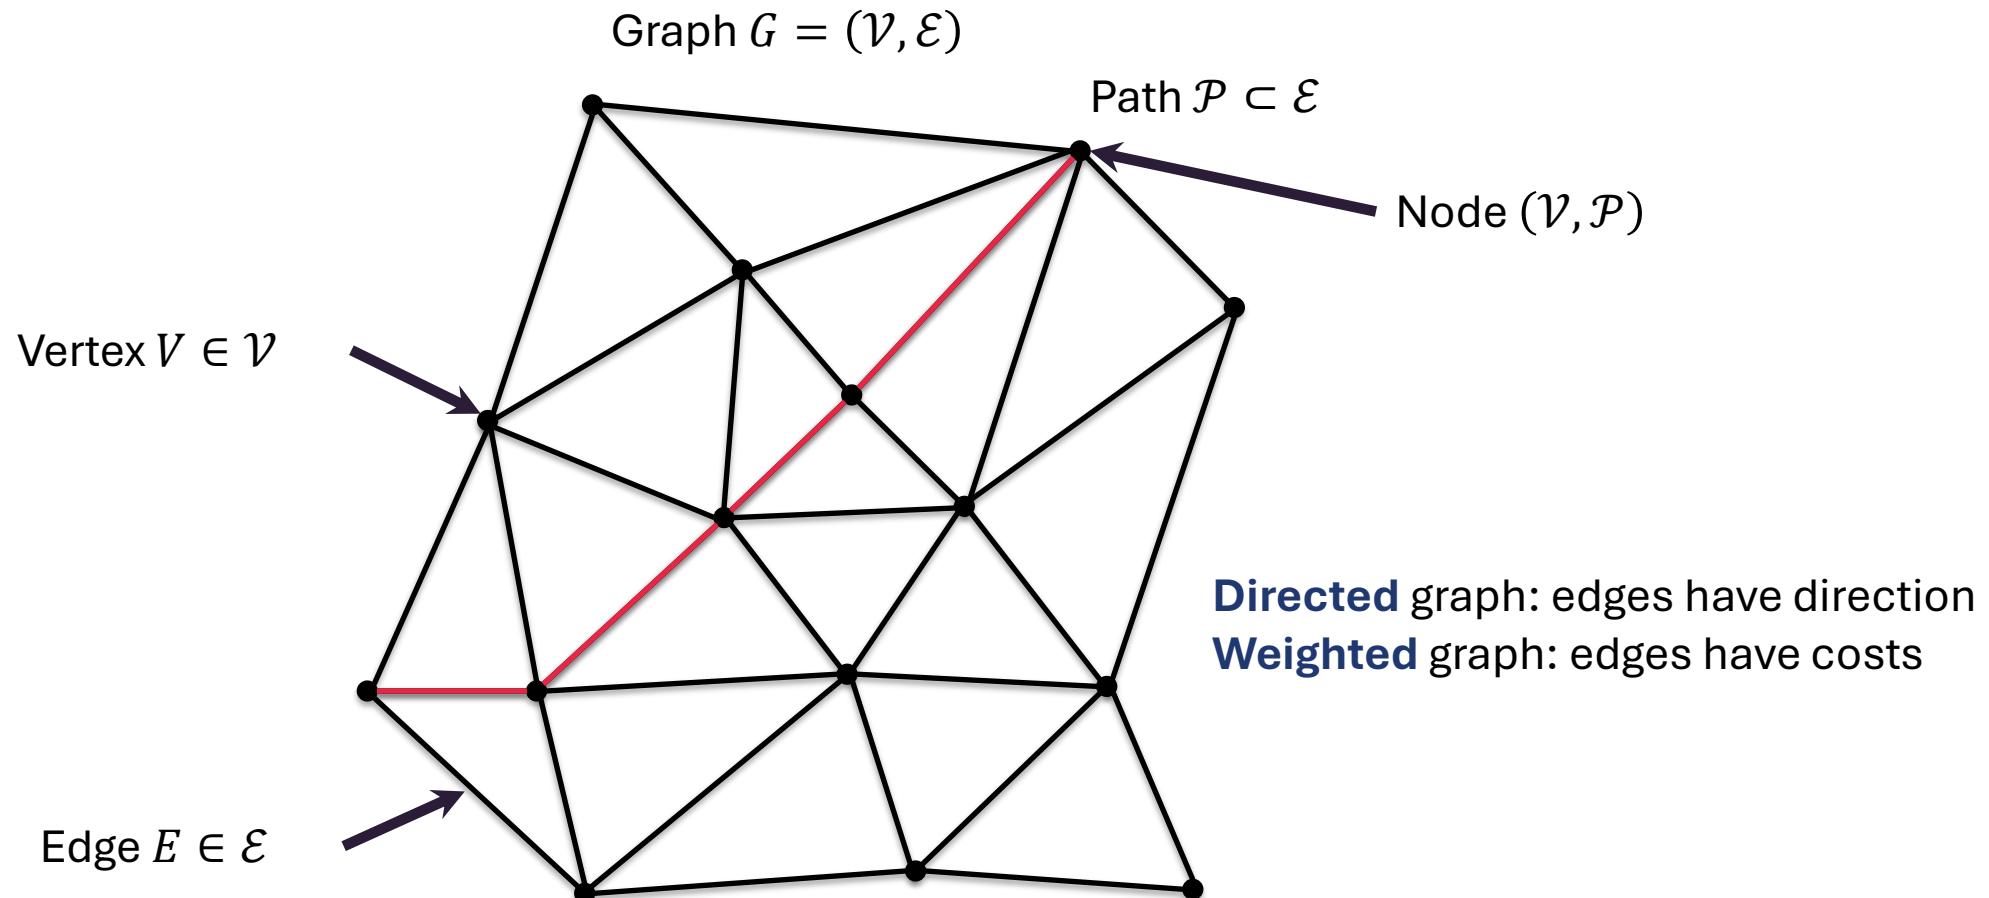


Exact

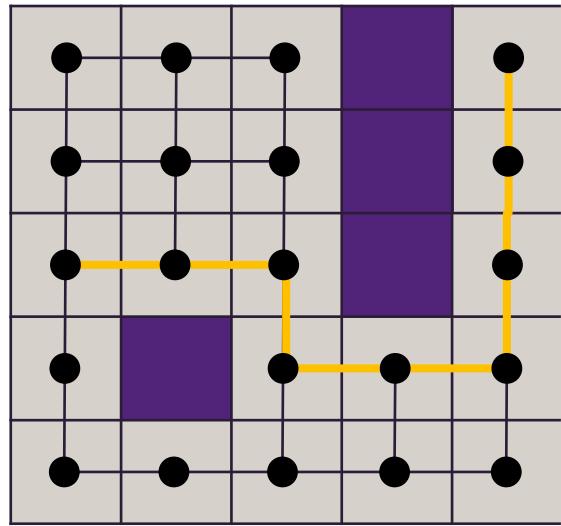


Graph Search Methods

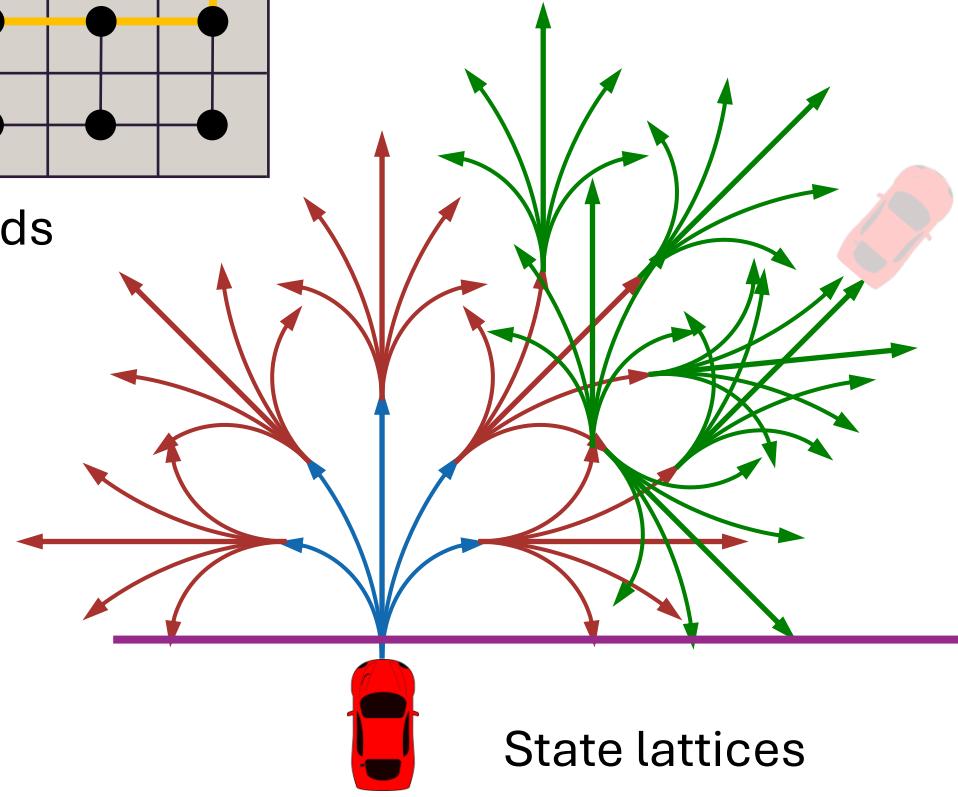
Some Terminology



Graph-based Representations

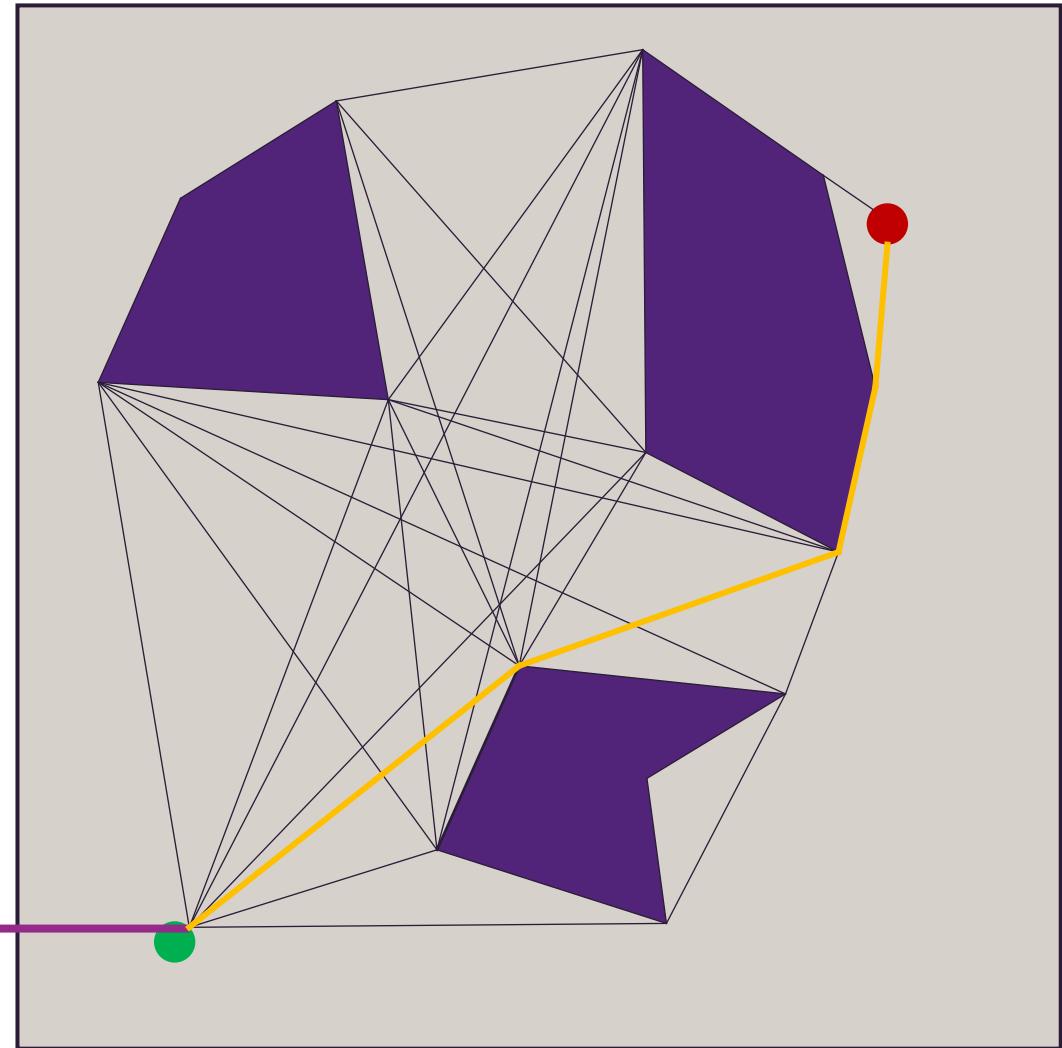


Grids



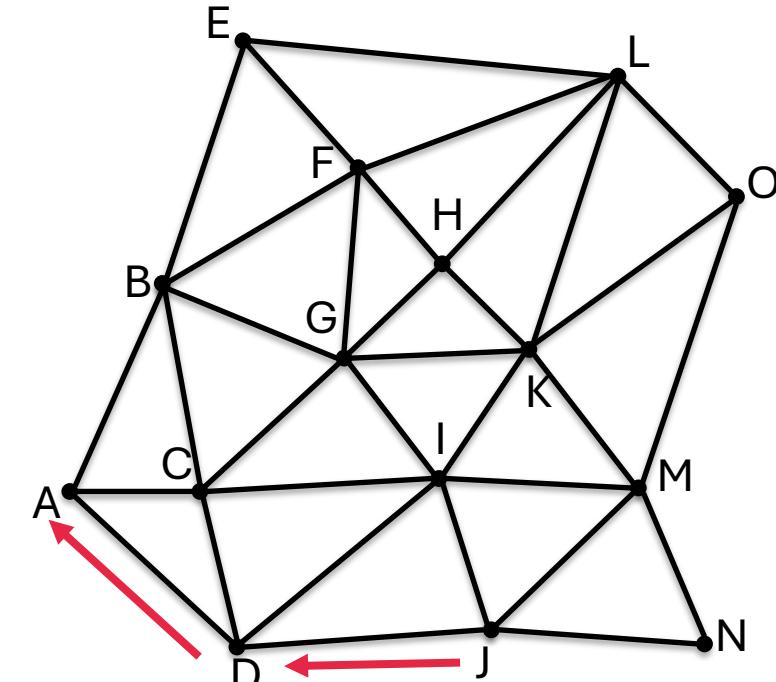
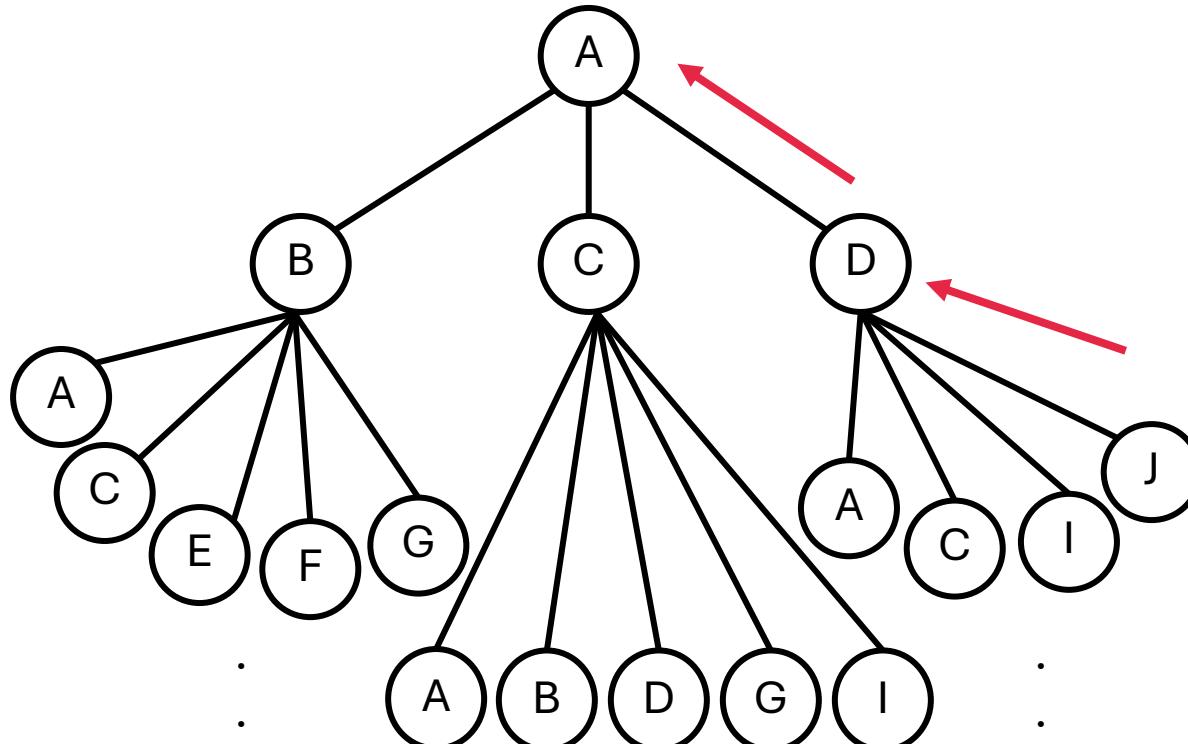
State lattices

Visibility graphs



Graph Search via Search Tree Construction

- Edges may have non-uniform costs (generally only consider positive costs)
- Construct search tree from root node



Basics of Forward Search

- Start from the start, grow tree until you find a solution (path to goal)
- Expanding a node refers to adding children to the tree, pushing them onto the open set
- Try to expand as few tree nodes as possible
- **Open** set maintains a list of frontier (unexpanded) plans
 - Keeps track of what nodes to expand next
 - Often stored as a priority queue
 - For each node in the open list, we know of at least one path to it from the start
- **Closed** (visited) set keeps track of nodes that have been expanded
 - For each node in the closed list, we've already found the lowest-cost path to it from the start

Forward Search Algorithm

FORWARD_SEARCH

- 1 $Q.Insert(x_I)$ and mark x_I as visited
- 2 **while** Q not empty **do**
- 3 $x \leftarrow Q.GetFirst()$
- 4 **if** $x \in X_G$
- 5 **return** SUCCESS
- 6 **forall** $u \in U(x)$
- 7 $x' \leftarrow f(x, u)$
- 8 **if** x' not visited
- 9 Mark x' as visited
- 10 $Q.Insert(x')$
- 11 **else**
- 12 Resolve duplicate x'
- 13 **return** FAILURE

Figure 2.4: A general template for forward search.

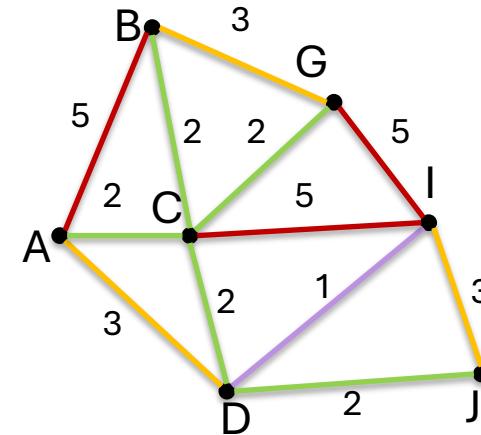
LaValle, Steven M. *Planning algorithms*. Cambridge university press, 2006, p. 33

Dijkstra's Algorithm

- Published by Edsger Dijkstra in 1959
- Basic idea of expanding in order of closest to start (BFS with edge costs)
 - **Open** queue ordered according to currently known best cost to arrive
- One of the most commonly used routing algorithms in graph traversal problems
- Asymptotically the fastest known single-source shortest path algorithm for arbitrary directed graphs

Dijkstra's Algorithm Example

B(0)



FORWARD SEARCH

```

1   Q.Insert( $x_1$ )
2   while  $Q$  not empty do
3        $x \leftarrow Q.GetFirst()$ ,  $R.Insert(x)$ 
4       if  $x \in X_G$ 
5           return SUCCESS
6       forall  $u \in U(x)$ 
7            $x' \leftarrow f(x, u)$ 
8           if  $x' \notin Q$  or  $R$ 
9               Mark  $x'$  as visited
10               $Q.Insert(x')$ 
11           else
12               Resolve duplicate  $x'$ 
13   return FAILURE
  
```

Open (Q):

Node	Cost	Parent

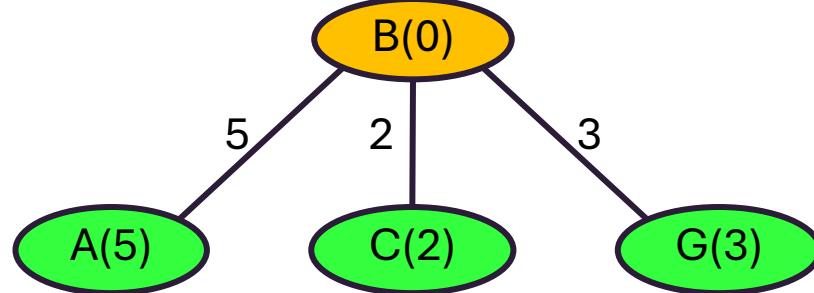
Closed (R):

Node	Cost	Parent

Our Dijkstra queue will be ordered by cost to arrive:

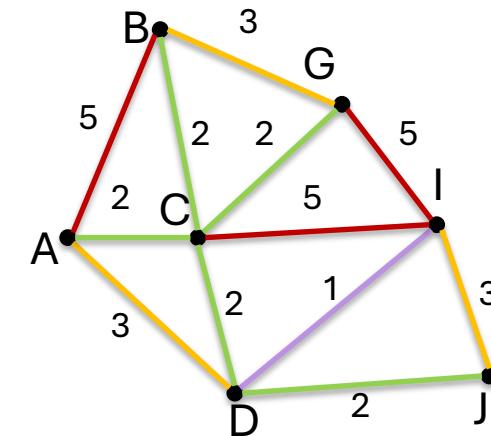
- push ($Q.Insert$) by cost
- pop ($Q.GetFirst$) from the front, and add it to the closed list

Dijkstra's Algorithm Example



FORWARD_SEARCH

- 1 $Q.Insert(x_1)$
- 2 **while** Q not empty **do**
- 3 $x \leftarrow Q.GetFirst()$, $R.Insert(x)$
- 4 **if** $x \in X_G$
- 5 **return** SUCCESS
- 6 **forall** $u \in U(x)$
- 7 $x' \leftarrow f(x, u)$
- 8 **if** $x' \notin Q$ **or** R
- 9 ~~Mark x' as visited~~
- 10 $Q.Insert(x')$
- 11 **else**
- 12 Resolve duplicate x'
- 13 **return** FAILURE



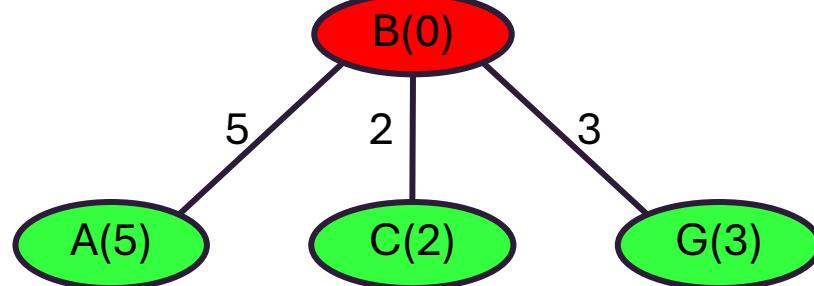
Open (Q):

Node	Cost	Parent
C	2	B
G	3	B
A	5	B

Closed (R):

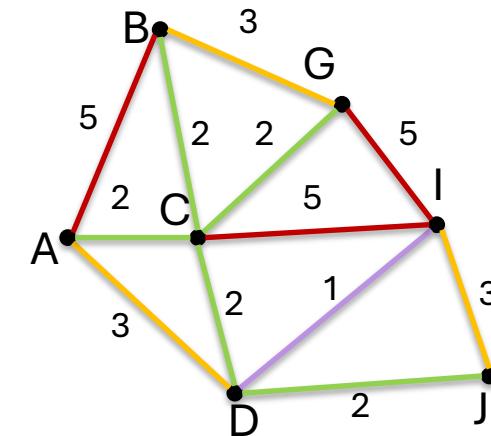
Node	Cost	Parent
B	0	

Dijkstra's Algorithm Example



```

FORWARD_SEARCH
1   Q.Insert( $x_1$ )
2   while  $Q$  not empty do
3        $x \leftarrow Q.GetFirst()$ ,  $R.Insert(x)$ 
4       if  $x \in X_G$ 
5           return SUCCESS
6       forall  $u \in U(x)$ 
7            $x' \leftarrow f(x, u)$ 
8           if  $x' \notin Q$  or  $R$ 
9               Mark  $x'$  as visited
10               $Q.Insert(x')$ 
11          else
12              Resolve duplicate  $x'$ 
13  return FAILURE
  
```



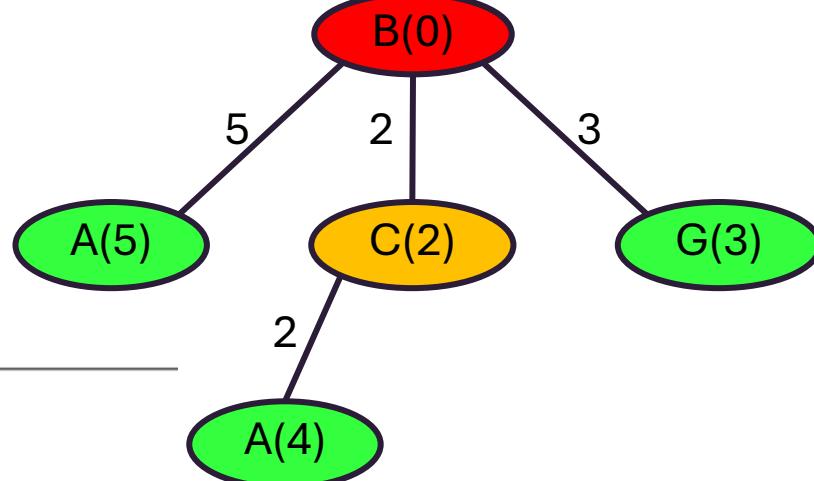
Open (Q):

Node	Cost	Parent
G	3	B
A	5	B

Closed (R):

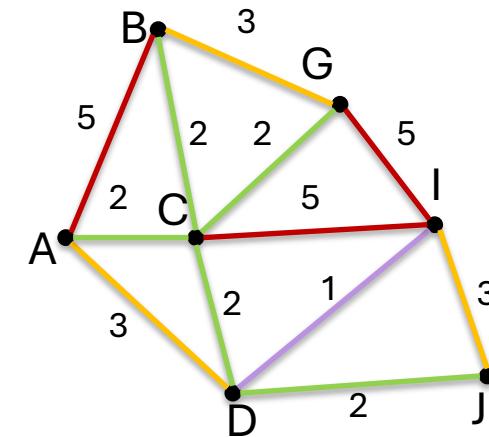
Node	Cost	Parent
B	0	
C	2	B

Dijkstra's Algorithm Example



```

FORWARD_SEARCH
1   Q.Insert( $x_I$ )
2   while  $Q$  not empty do
3        $x \leftarrow Q.GetFirst()$ ,  $R.Insert(x)$ 
4       if  $x \in X_G$ 
5           return SUCCESS
6       forall  $u \in U(x)$ 
7            $x' \leftarrow f(x, u)$ 
8           if  $x' \notin Q$  or  $R$ 
9               Mark  $x'$  as visited
10               $Q.Insert(x')$ 
11           else
12               Resolve duplicate  $x'$ 
13   return FAILURE
  
```



Open (Q):

Node	Cost	Parent
G	3	B
A	4	C

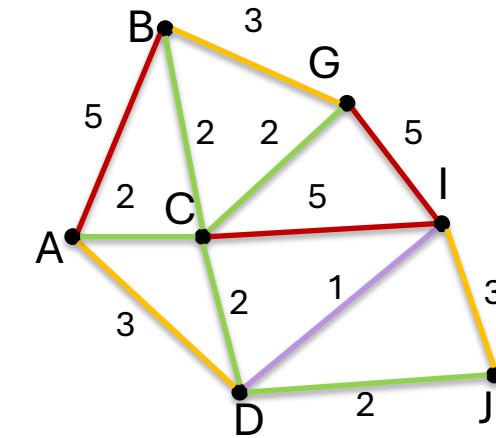
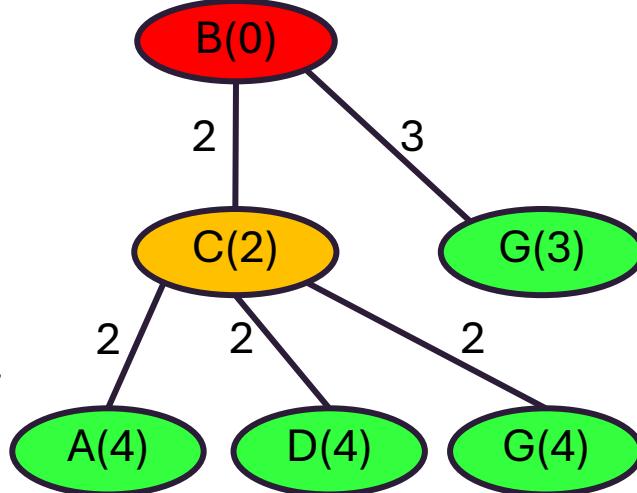
Closed (R):

Node	Cost	Parent
B	0	
C	2	B

Dijkstra's Algorithm Example

```

FORWARD_SEARCH
1   Q.Insert( $x_I$ )
2   while  $Q$  not empty do
3        $x \leftarrow Q.GetFirst()$ , R.Insert( $x$ )
4       if  $x \in X_G$ 
5           return SUCCESS
6       forall  $u \in U(x)$ 
7            $x' \leftarrow f(x, u)$ 
8           if  $x' \notin Q$  or R
9               Mark  $x'$  as visited
10              Q.Insert( $x'$ )
11           else
12               Resolve duplicate  $x'$ 
13   return FAILURE
  
```



Open (Q):

Node	Cost	Parent
G	3	B
A	4	C
D	4	C

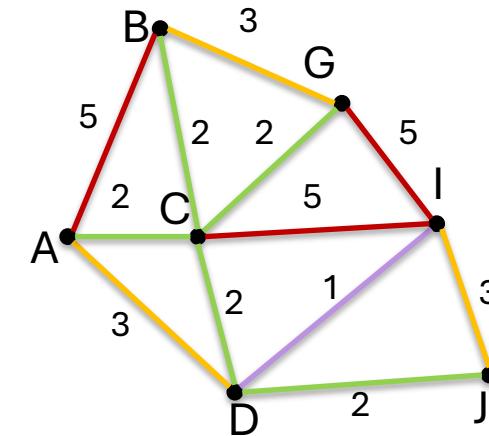
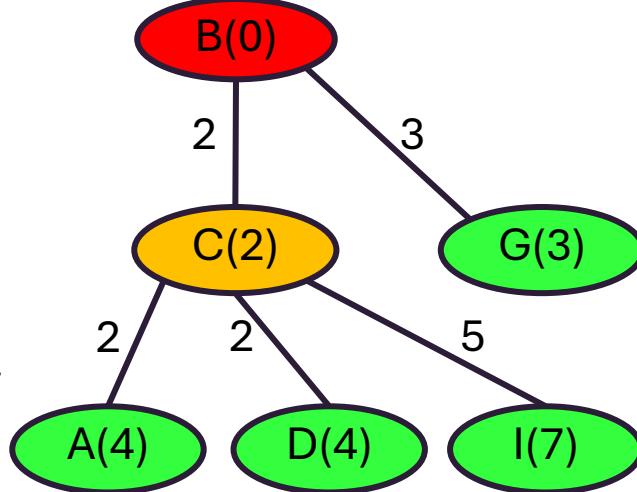
Closed (R):

Node	Cost	Parent
B	0	
C	2	B

Dijkstra's Algorithm Example

```

FORWARD_SEARCH
1   Q.Insert( $x_I$ )
2   while  $Q$  not empty do
3        $x \leftarrow Q.GetFirst()$ , R.Insert( $x$ )
4       if  $x \in X_G$ 
5           return SUCCESS
6       forall  $u \in U(x)$ 
7            $x' \leftarrow f(x, u)$ 
8           if  $x' \notin Q$  or R
9               Mark  $x'$  as visited
10              Q.Insert( $x'$ )
11           else
12               Resolve duplicate  $x'$ 
13   return FAILURE
  
```



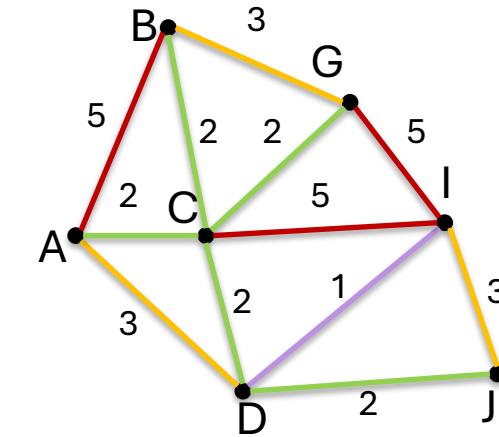
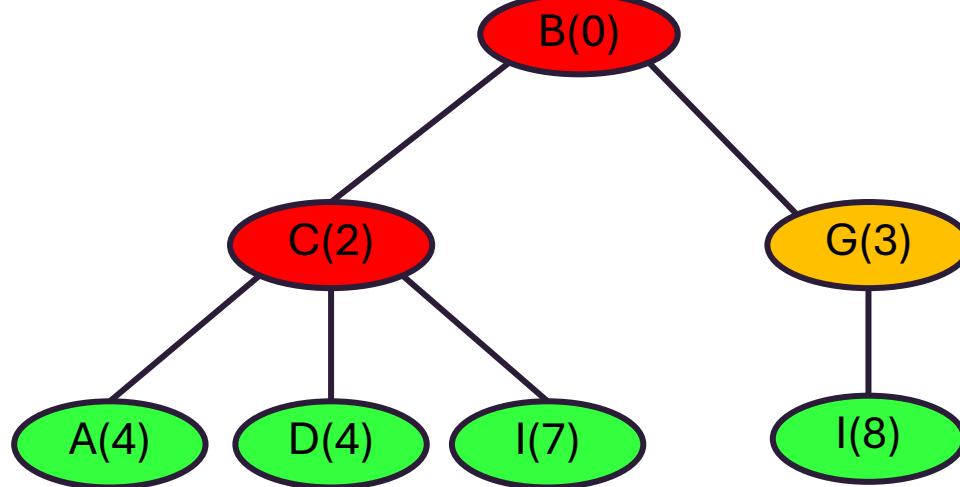
Open (Q):

Node	Cost	Parent
G	3	B
A	4	C
D	4	C
I	7	C

Closed (R):

Node	Cost	Parent
B	0	
C	2	B

Dijkstra's Algorithm Example



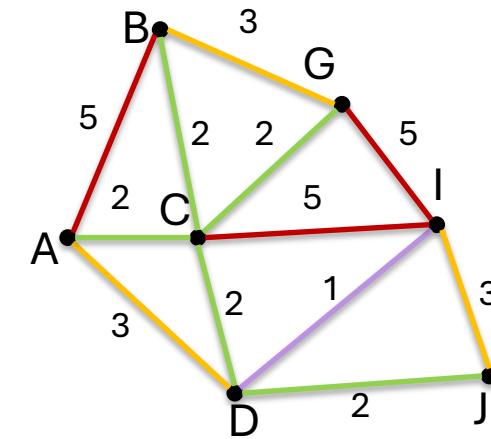
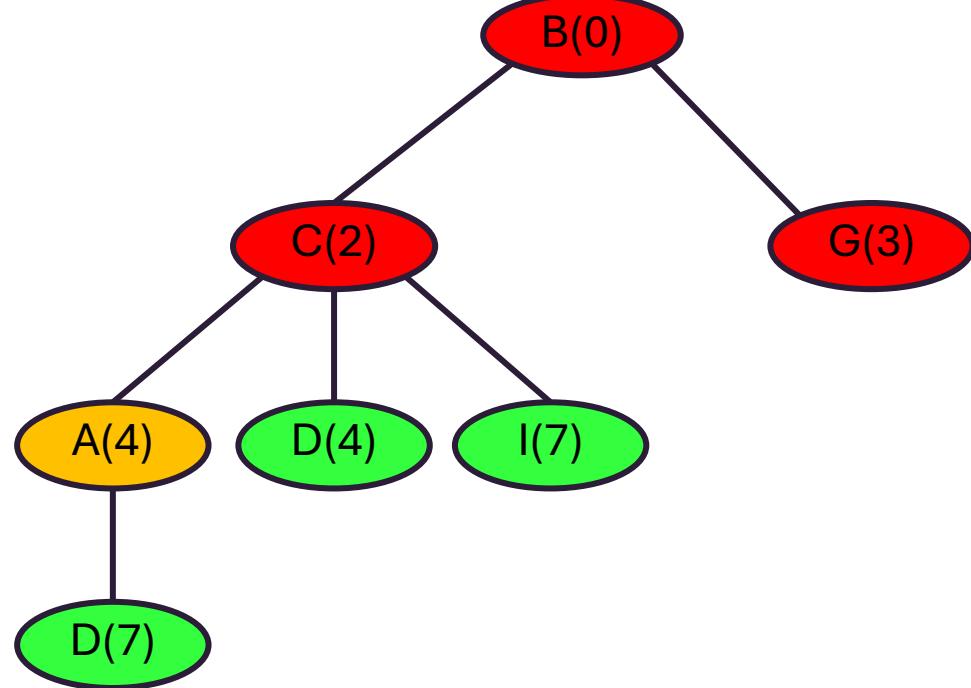
Open (Q):

Node	Cost	Parent
A	4	C
D	4	C
I	7	C

Closed (R):

Node	Cost	Parent
B	0	
C	2	B
G	3	B

Dijkstra's Algorithm Example



Open (Q):

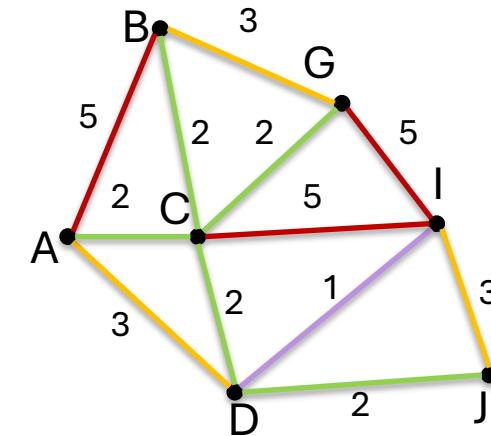
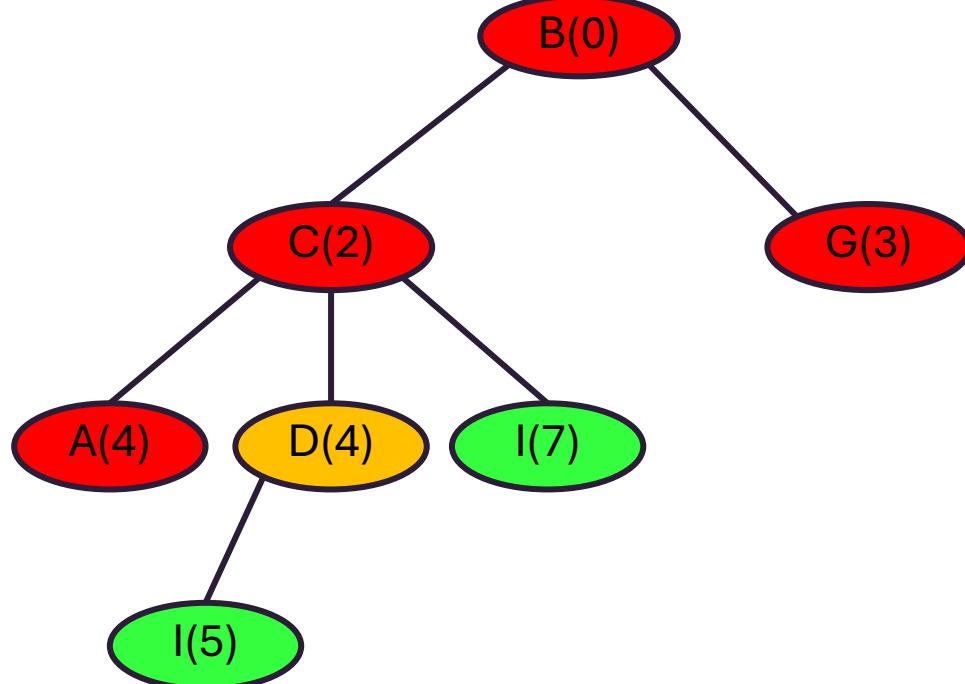
Node	Cost	Parent
D	4	C
I	7	C

Closed (R):

Node	Cost	Parent
D	4	C
I	7	C

Node	Cost	Parent
B	0	
C	2	B
G	3	B
A	4	C

Dijkstra's Algorithm Example



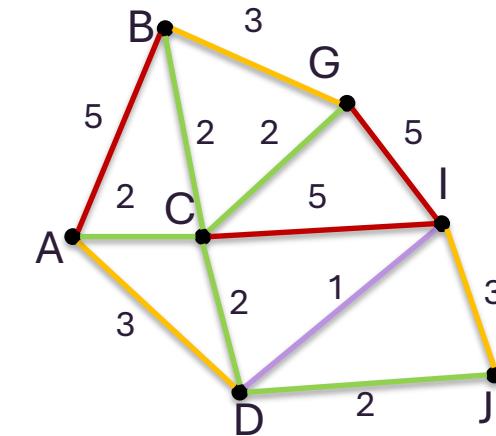
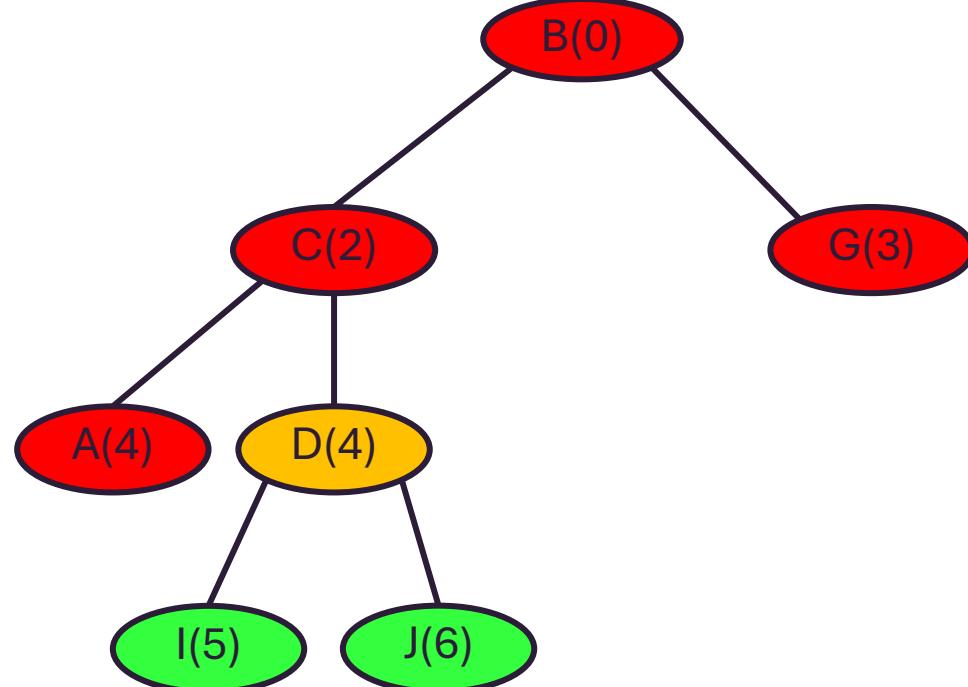
Open (Q):

Node	Cost	Parent
I	5	D
I	7	C

Closed (R):

Node	Cost	Parent
B	0	
C	2	B
G	3	B
A	4	C
D	4	C

Dijkstra's Algorithm Example



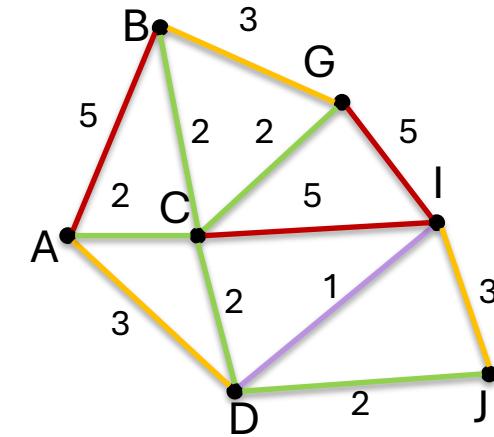
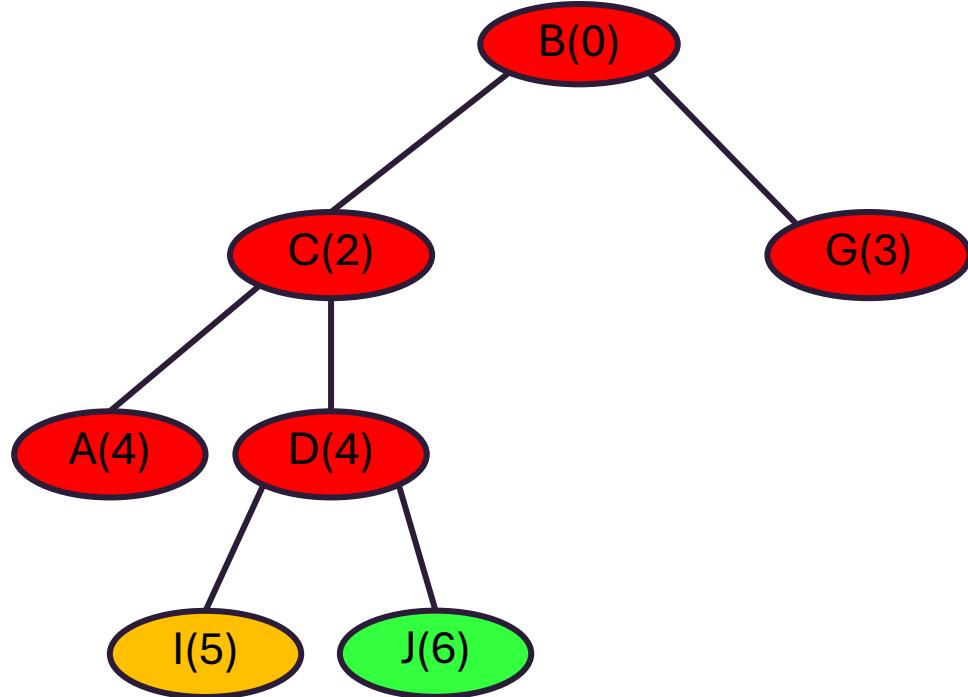
Open (Q):

Node	Cost	Parent
I	5	D
J	6	D

Closed (R):

Node	Cost	Parent
B	0	
C	2	B
G	3	B
A	4	C
D	4	C

Dijkstra's Algorithm Example



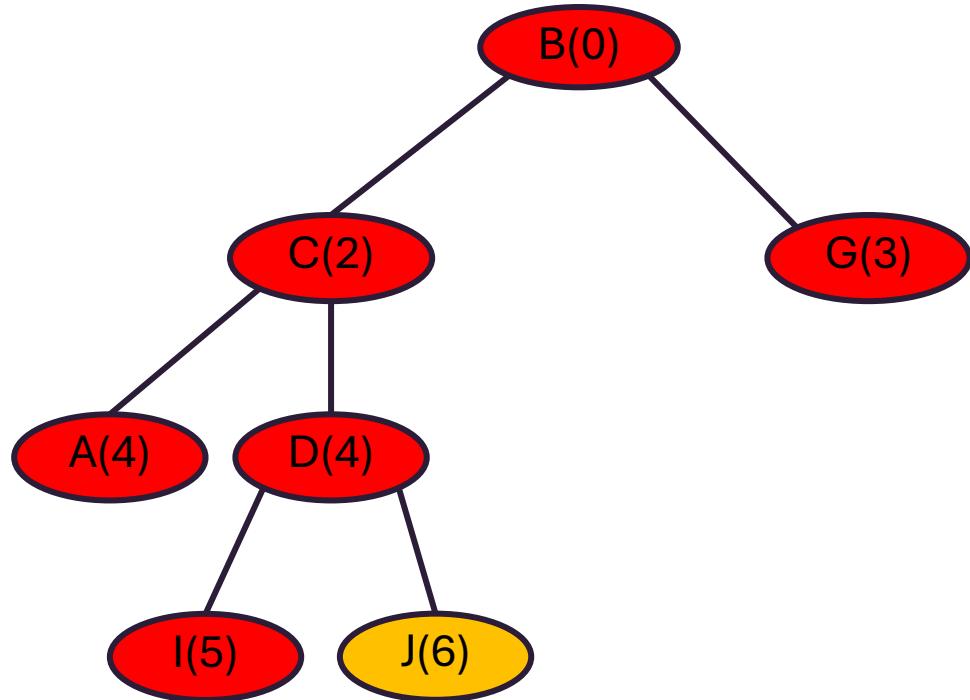
Open (Q):

Node	Cost	Parent
J	6	D

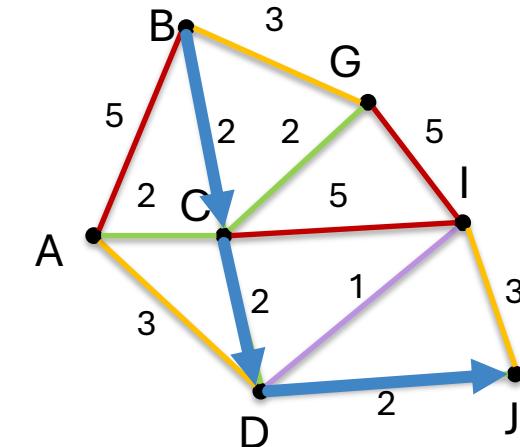
Closed (R):

Node	Cost	Parent
B	0	
C	2	B
G	3	B
A	4	C
D	4	C
I	5	D

Dijkstra's Algorithm Example



Final path solution: B → C → D → J
with path cost 6



Frontier (Q):

Node	Cost	Parent

Closed

Node	Cost	Parent
B	0	
C	2	B
G	3	B
A	4	C
D	4	C
I	5	D
J	6	D

A* Heuristic Search

- Heuristic:

- Any optimistic estimate of how close a state is to a goal
- Designed for a particular search problem
- Examples: Manhattan distance, Euclidean distance

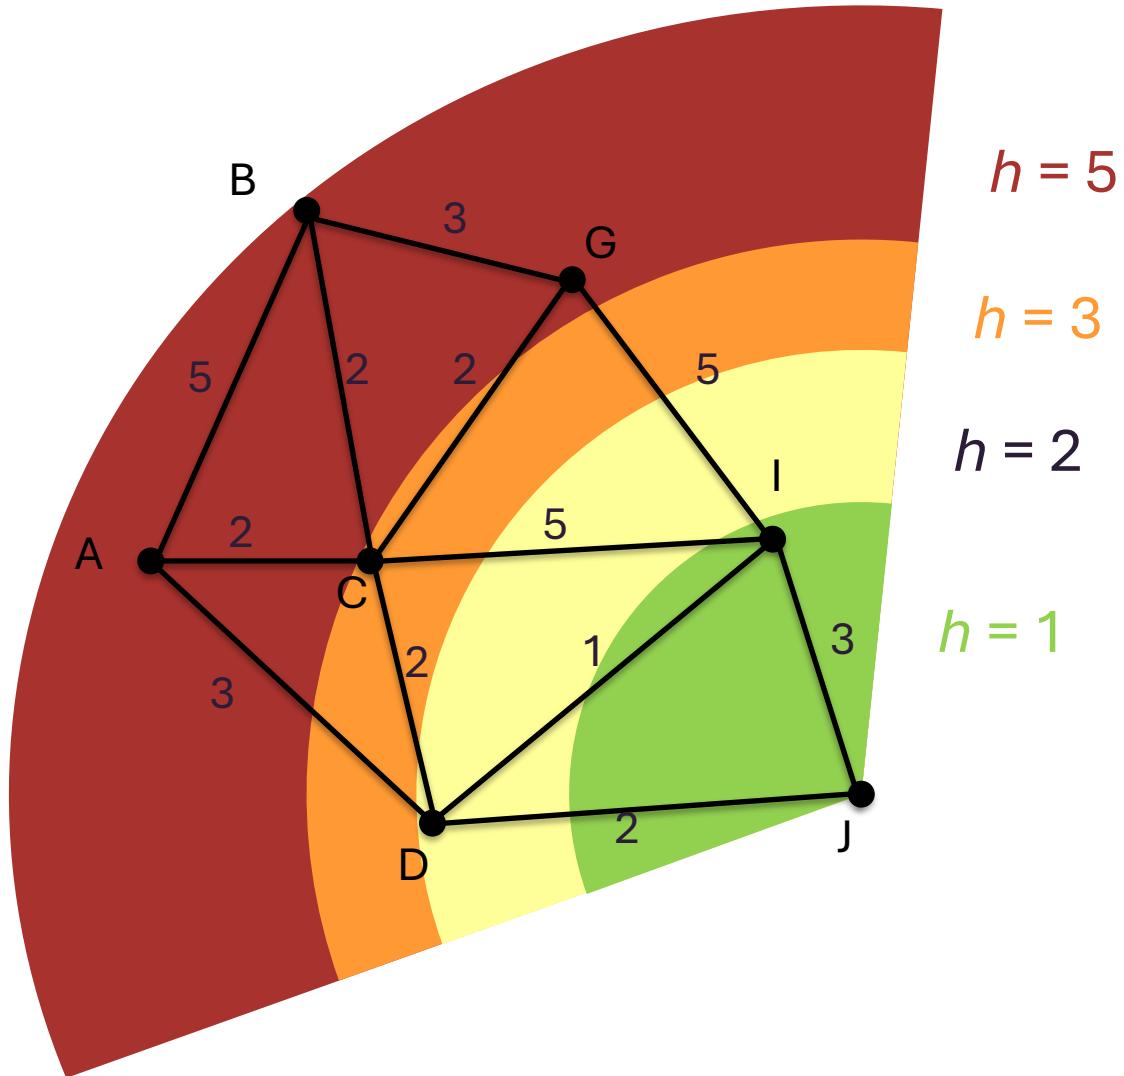
- A* priority:

$$f(n) = g(n) + h(n)$$

Cost to arrive Heuristic cost to goal

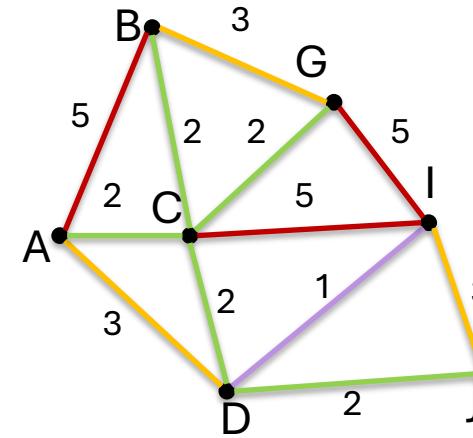


A* Heuristic



A* Algorithm Example

B(0)



Our A* queue will be ordered by cost to arrive + heuristic:

- push ($Q.Insert$) by A* priority, $f(n)$
- pop ($Q.GetFirst$) from the front, and add it to the closed list

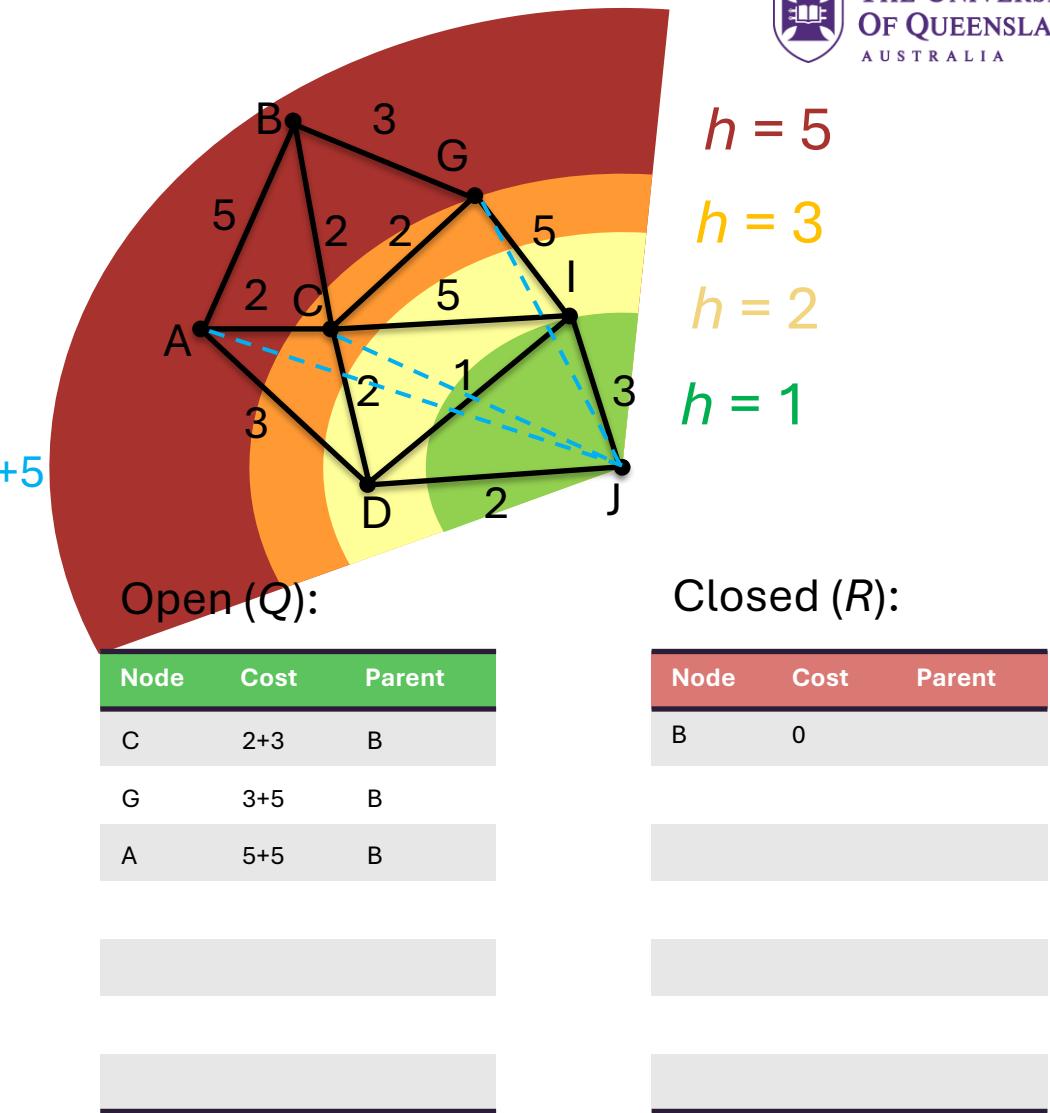
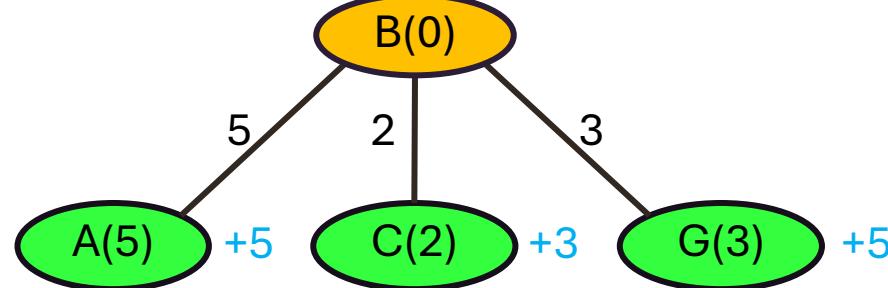
Open (Q):

Node	Cost	Parent

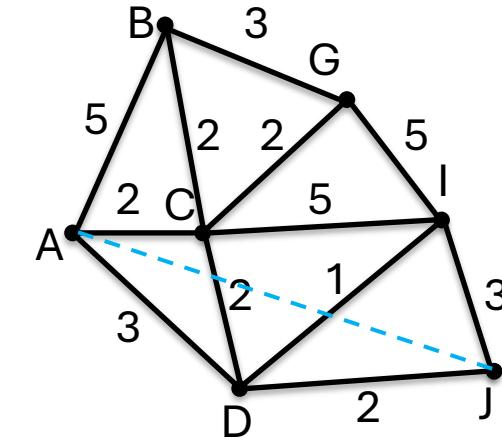
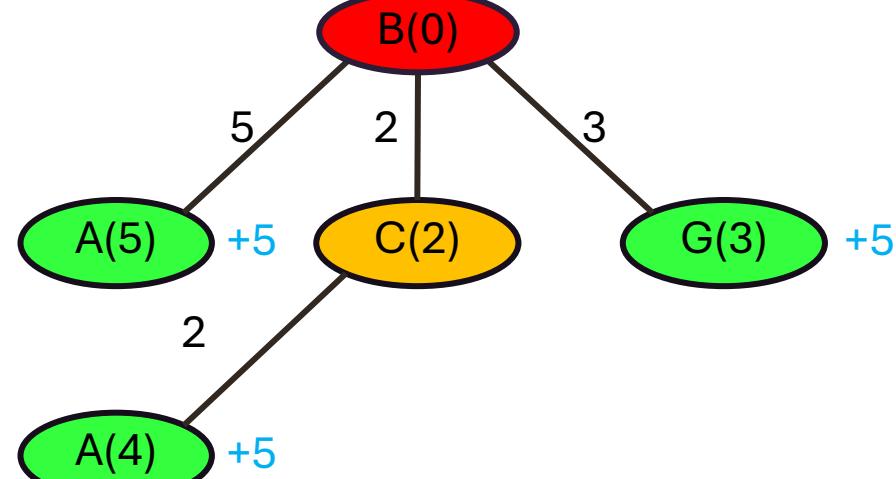
Closed (R):

Node	Cost	Parent
B	0	

A* Algorithm Example



A* Algorithm Example



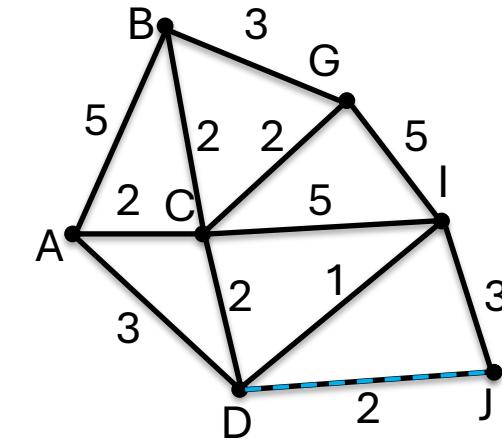
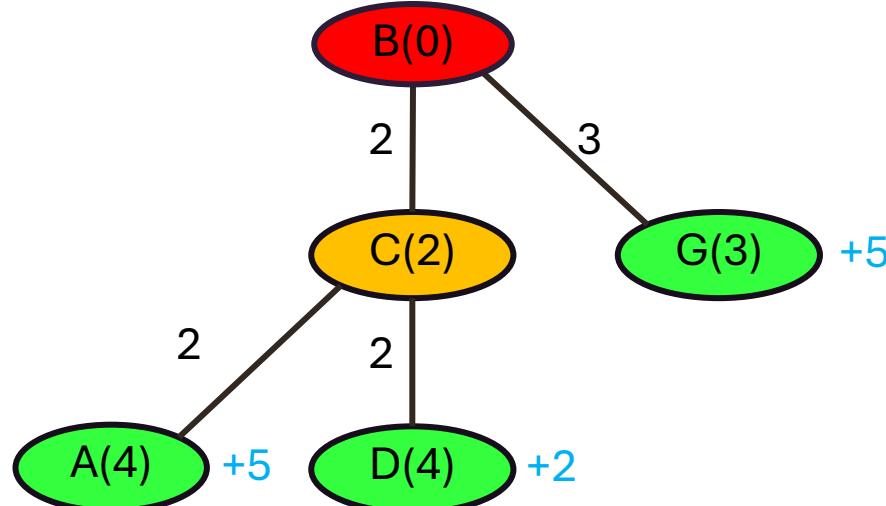
Open (Q):

Node	Cost	Parent
G	3+5	B
A	4+5	C

Closed (R):

Node	Cost	Parent
B	0	
C	2	B

A* Algorithm Example



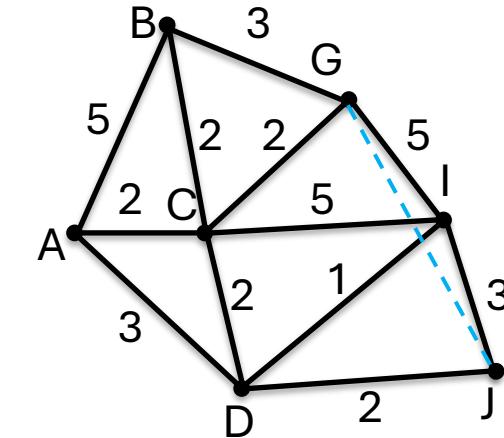
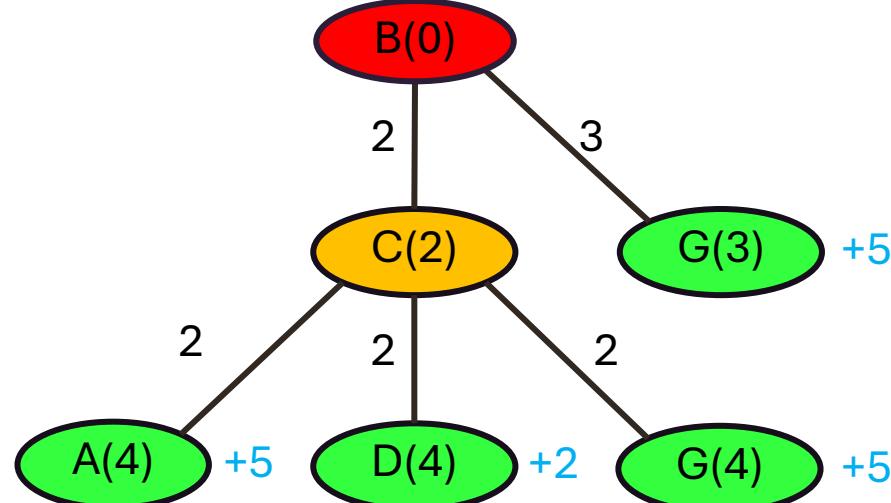
Open (Q):

Node	Cost	Parent
D	4+2	B
A	4+5	C

Closed (R):

Node	Cost	Parent
B	0	
C	2	B

A* Algorithm Example



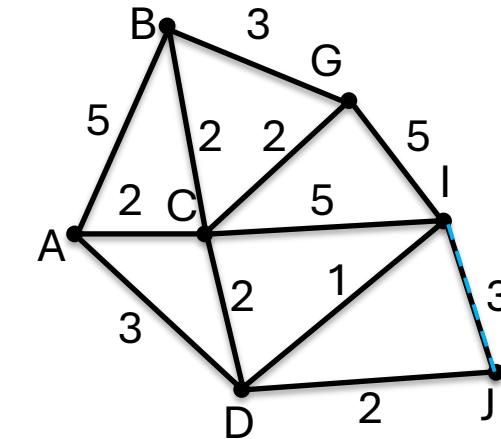
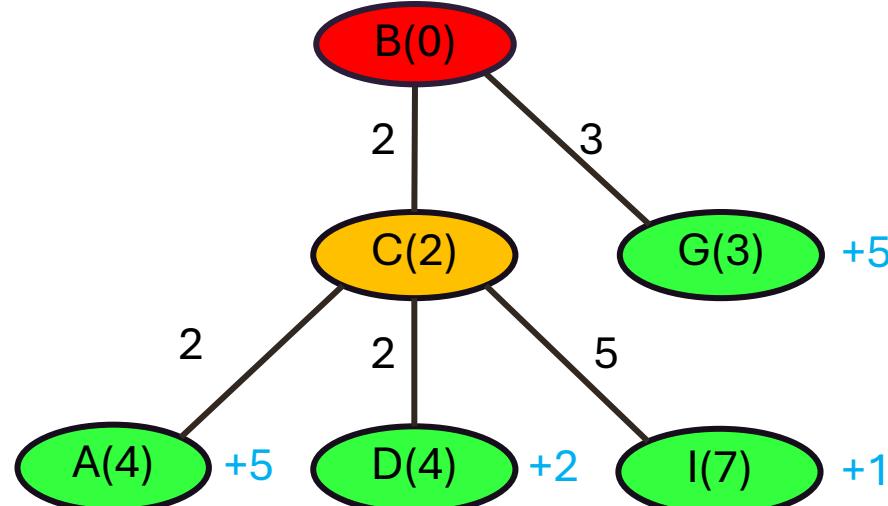
Open (Q):

Node	Cost	Parent
D	4+2	C
G	3+5	B
A	4+5	C

Closed (R):

Node	Cost	Parent
B	0	
C	2	B

A* Algorithm Example



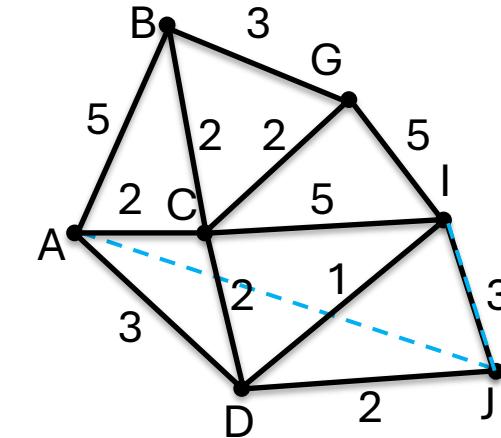
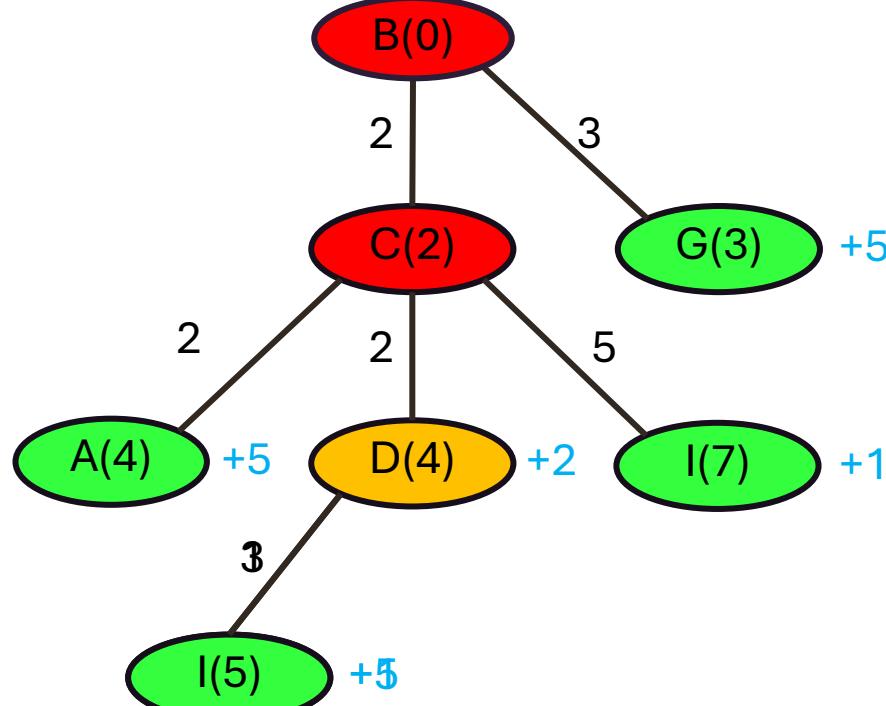
Open (Q):

Node	Cost	Parent
D	$4+2$	C
G	$3+5$	B
A	$4+5$	C

Closed (R):

Node	Cost	Parent
B	0	
C	2	B

A* Algorithm Example



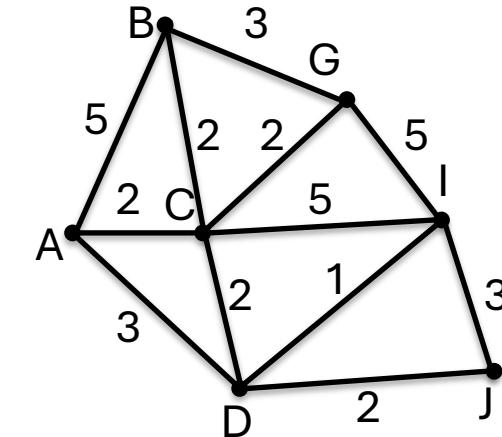
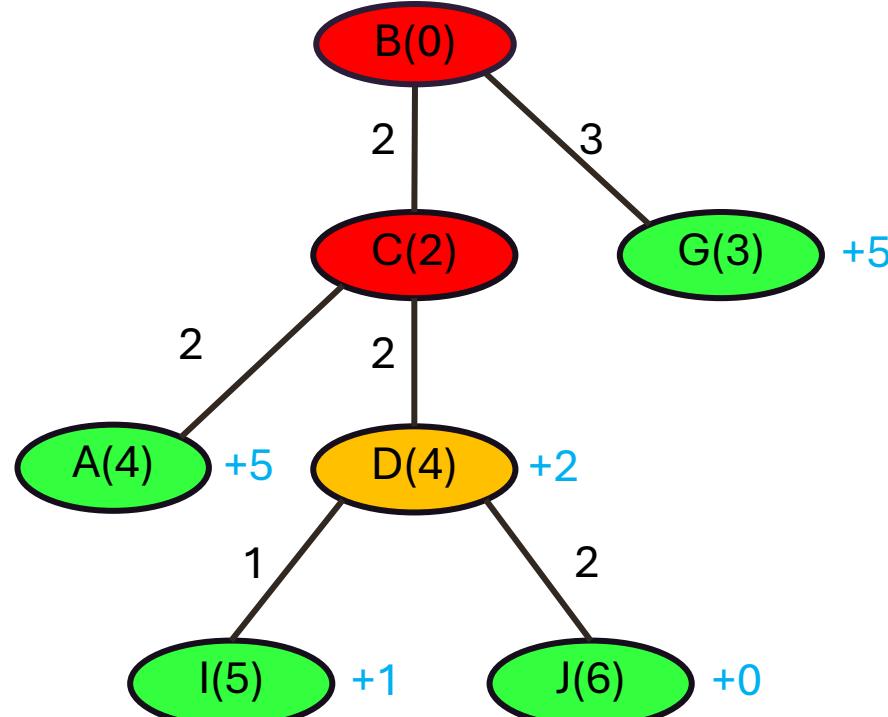
Open (Q):

Node	Cost	Parent
I	5+1	D
I	7+1	C
G	3+5	B
A	4+5	C

Closed (R):

Node	Cost	Parent
B	0	
C	2	B
D	4	C

A* Algorithm Example



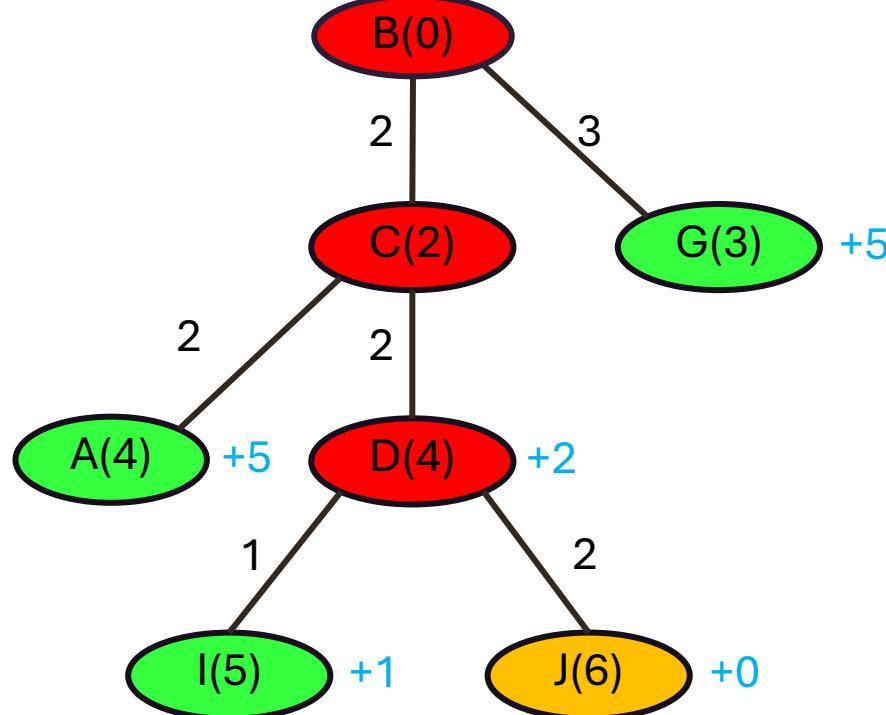
Open (Q):

Node	Cost	Parent
I	$6+0$	D
G	$3+5$	B
A	$4+5$	C

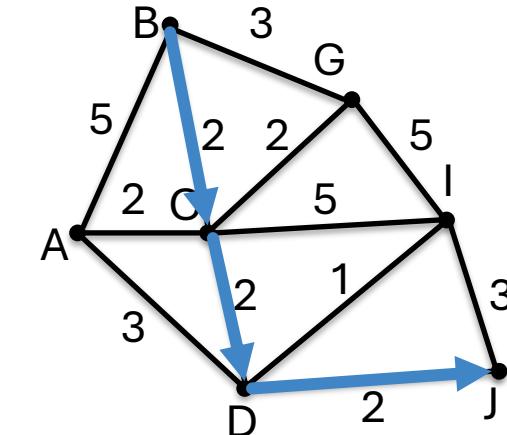
Closed (R):

Node	Cost	Parent
B	0	
C	2	B
D	4	C

A* Algorithm Example



Final path solution: B → C → D → J
with path cost 6



Open (Q):

Node	Cost	Parent
I	$5+1$	D
G	$3+5$	B
A	$4+5$	C

Closed (R):

Node	Cost	Parent
B	0	
C	2	B
D	4	C
J	6	D

A* Heuristic

- The heuristic must be **admissible**

- It never overestimates the cost

$$h(x) \leq d(x, goal)$$

True cost to goal



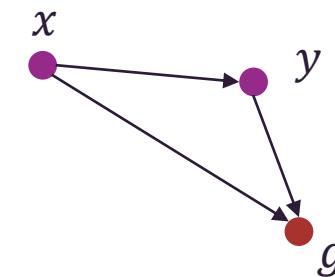
- The heuristic must be **consistent**

- For any pair of adjacent nodes x and y , where $d(x,y)$ is the cost of edge between them

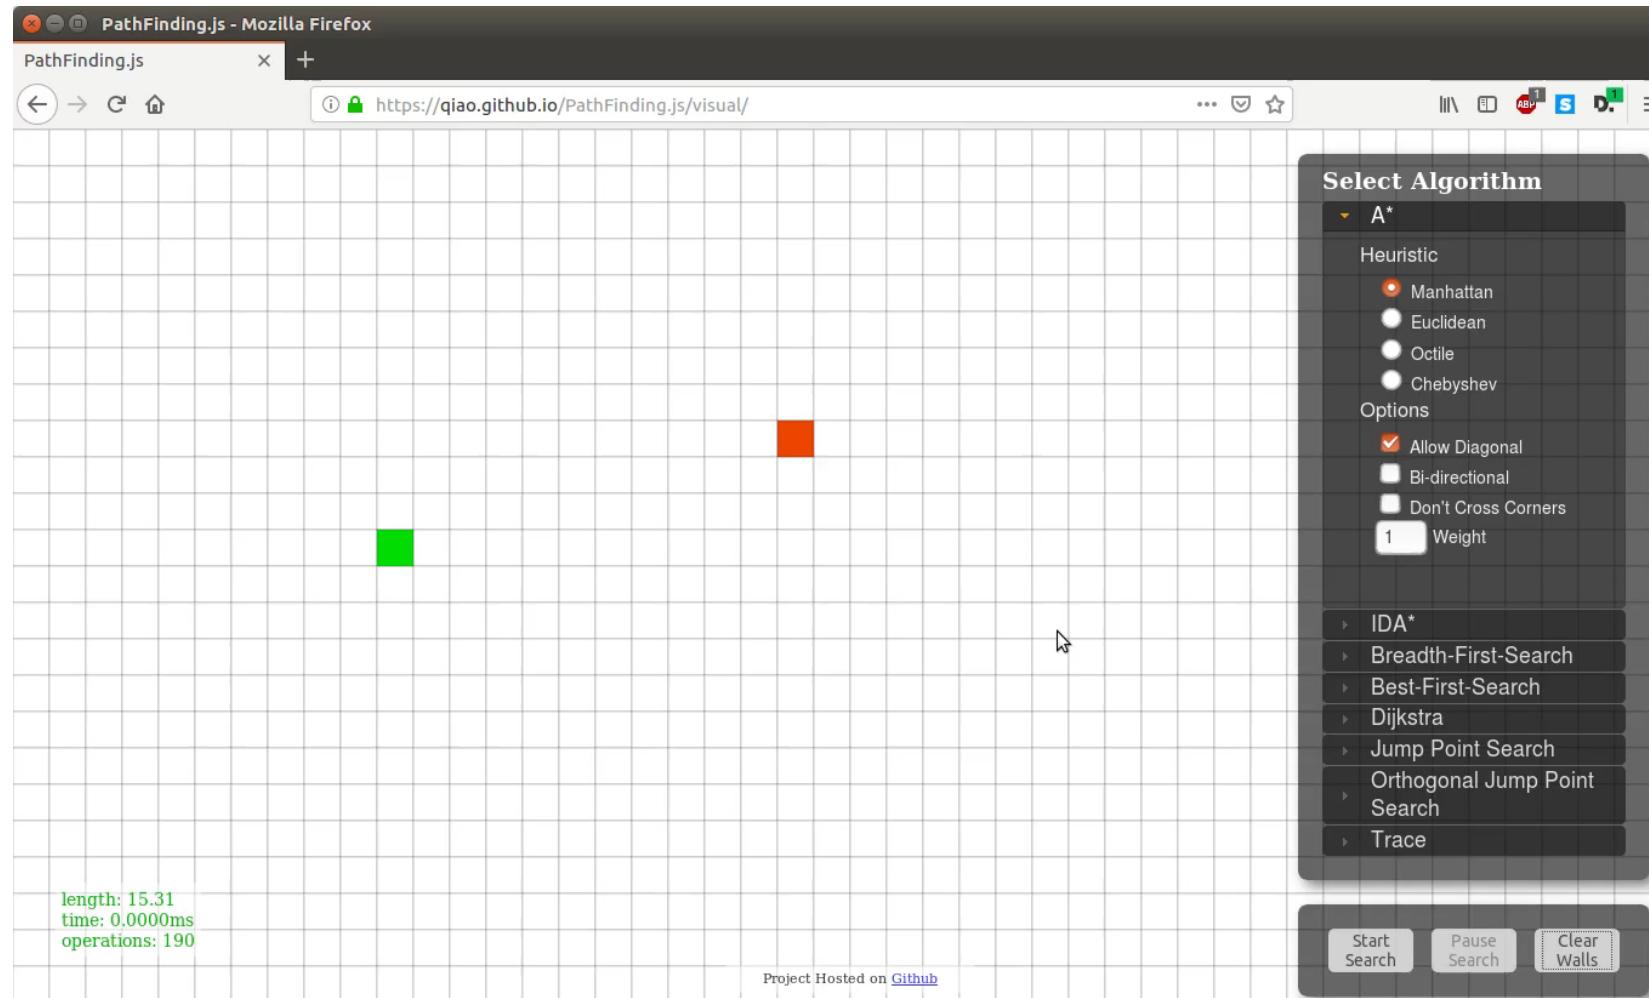
$$h(x) \leq d(x, y) + h(y)$$

- Typical valid heuristics:

- Euclidean distance
- Manhattan distance
- Zero (Dijkstra's algorithm)

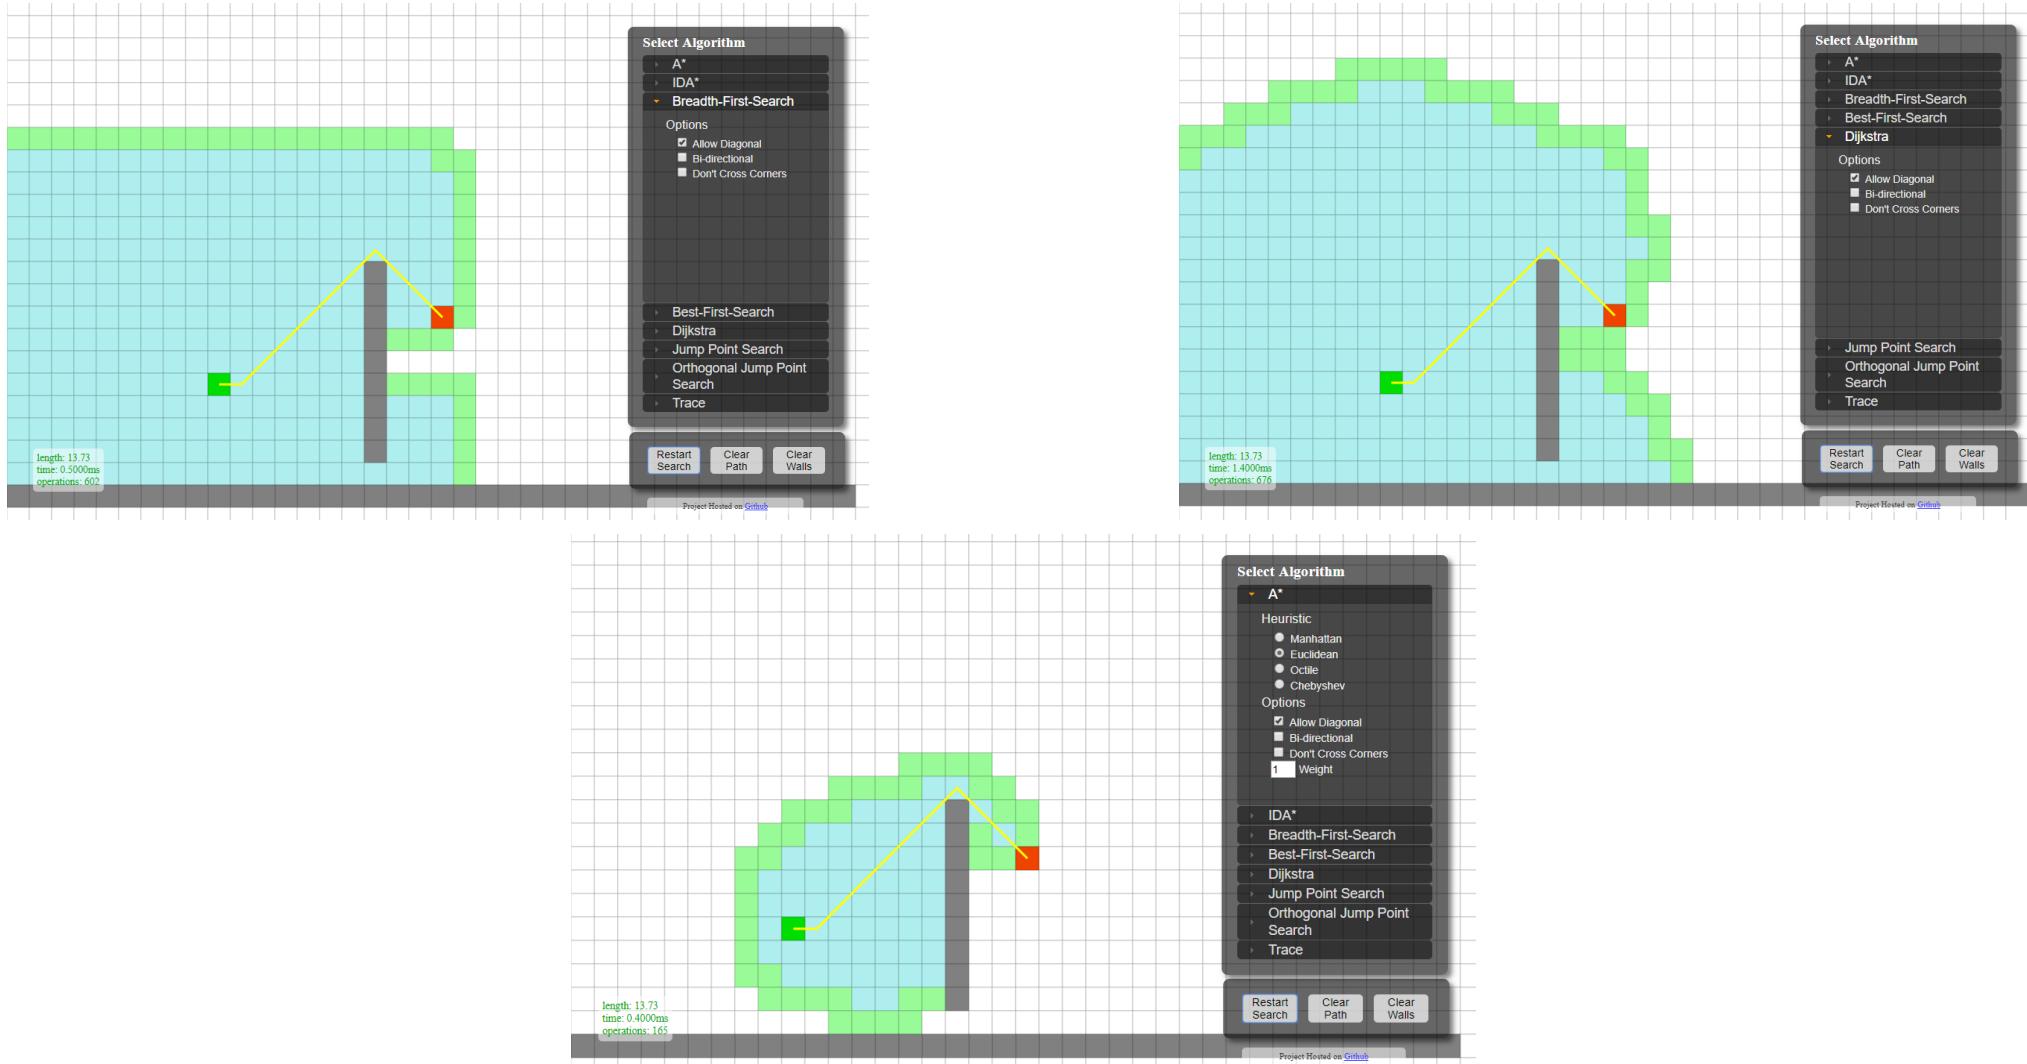


Example: Graph Search



<https://qiao.github.io/PathFinding.js/visual/>

Example: Graph Search

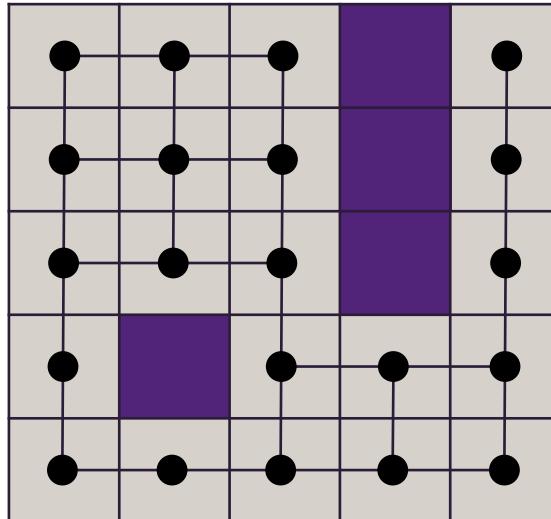


<https://qiao.github.io/PathFinding.js/visual/>

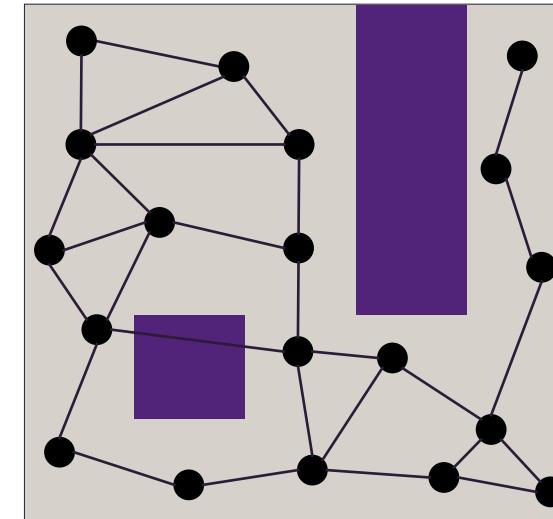
Sampling-based Methods

How to Build a Valid Graph

Explicitly use obstacles and grid



Use collision checks

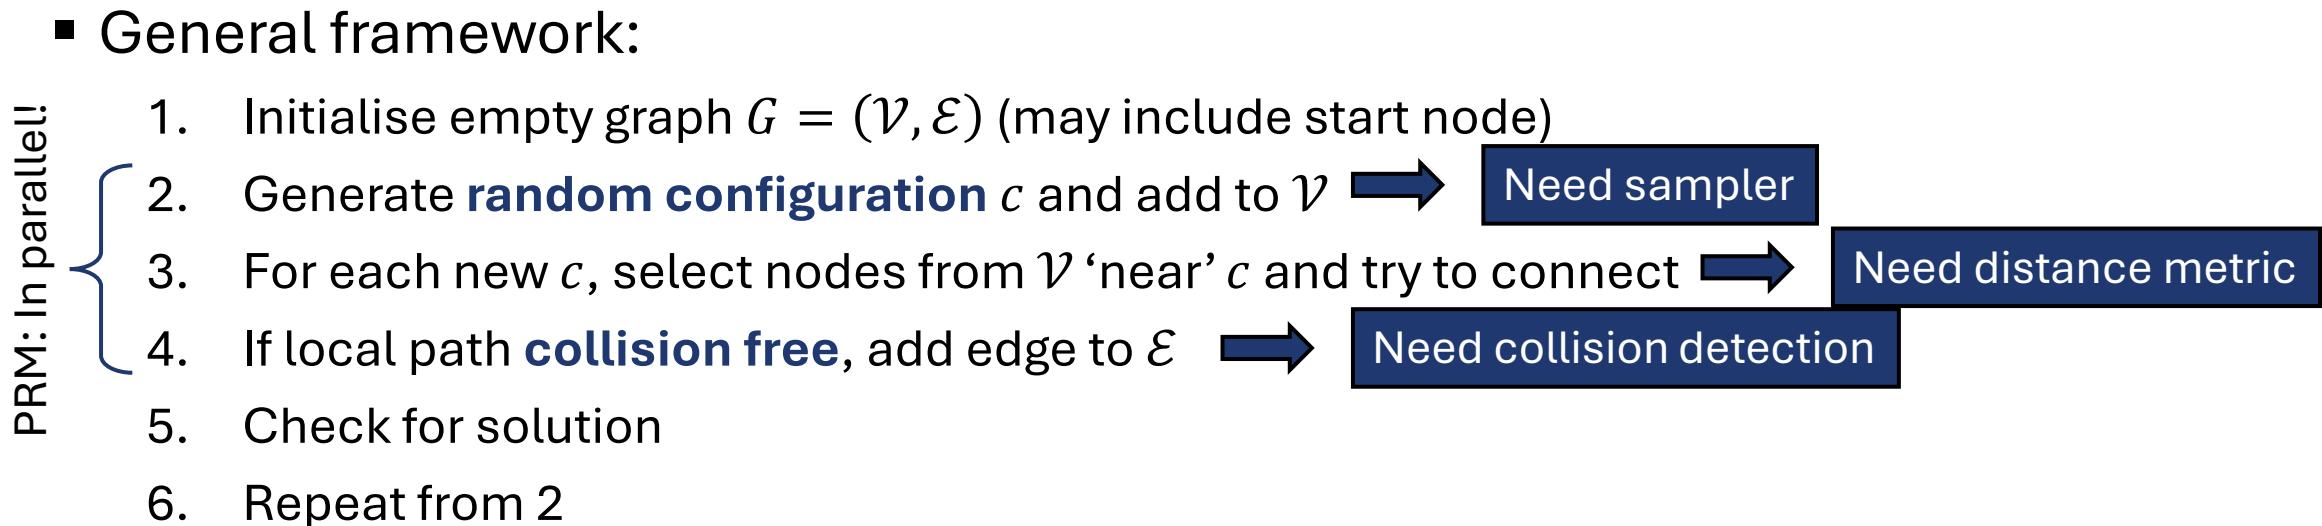


- Grid resolution trades off representation precision vs. search time
- Connectivity defined by grid shape
- Suffers from *curse of dimensionality*

- Representation precision determined by sampling distribution
- Connectivity defined by explicit collision check

Sampling-based Motion Planning

- Probabilistic Roadmaps (PRMs, 1996)
 - Construct roadmap, then search
- Rapidly Exploring Random Trees (RRTs, 1998)
 - Online roadmap construction and search

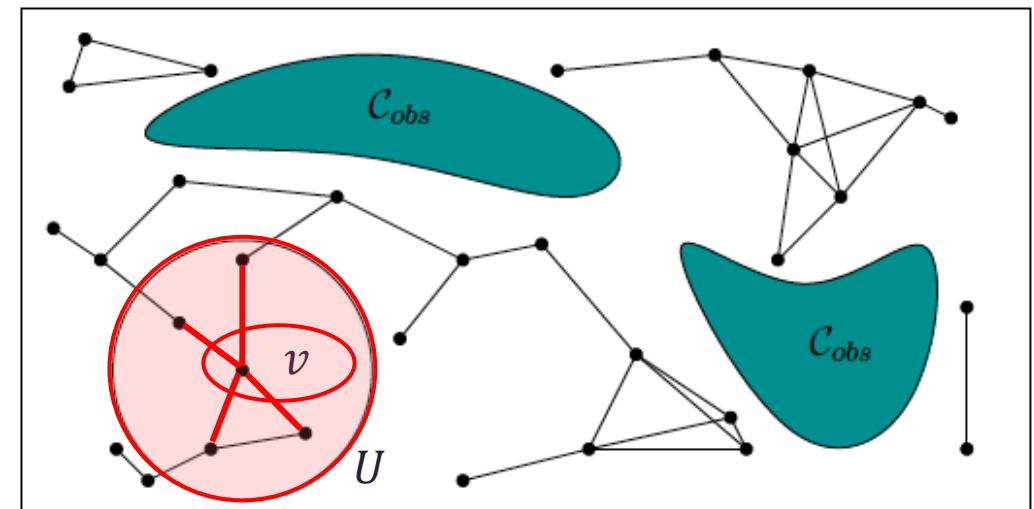
- General framework:
 1. Initialise empty graph $G = (\mathcal{V}, \mathcal{E})$ (may include start node)
 2. Generate **random configuration** c and add to \mathcal{V} → Need sampler
 3. For each new c , select nodes from \mathcal{V} ‘near’ c and try to connect → Need distance metric
 4. If local path **collision free**, add edge to \mathcal{E} → Need collision detection
 5. Check for solution
 6. Repeat from 2
- PRM: In parallel!
- 

PRM: Building the Roadmap

Algorithm 2: sPRM.

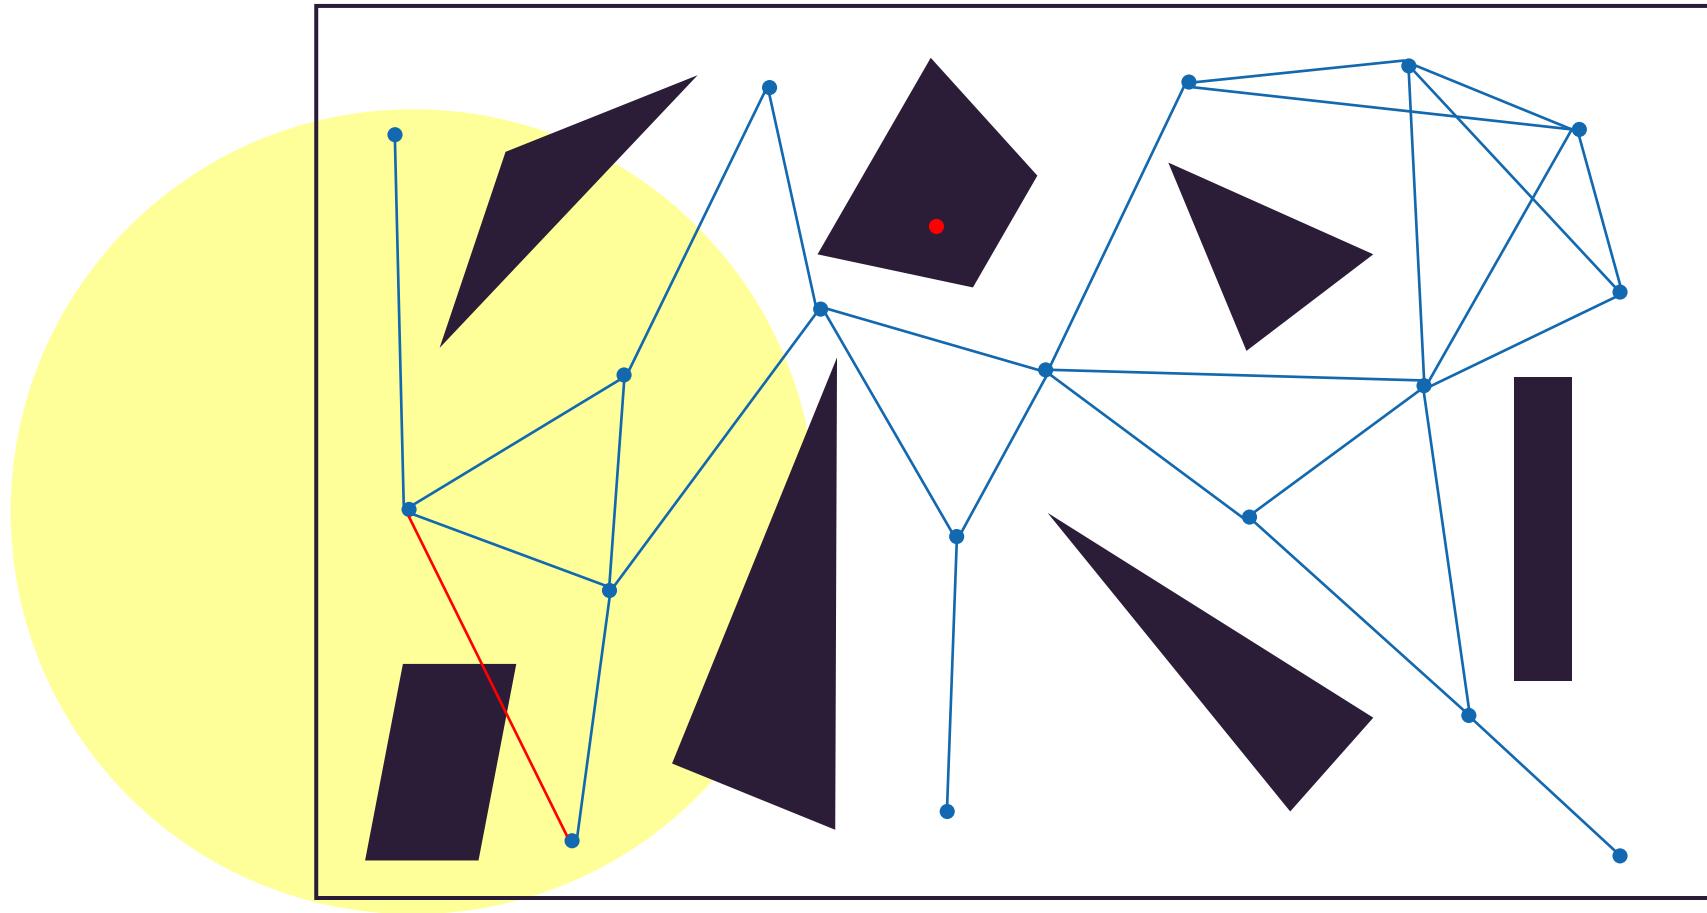
```

1  $V \leftarrow \{x_{\text{init}}\} \cup \{\text{SampleFree}_i\}_{i=1,\dots,n}; E \leftarrow \emptyset;$ 
2 foreach  $v \in V$  do
3    $U \leftarrow \text{Near}(G = (V, E), v, r) \setminus \{v\};$ 
4   foreach  $u \in U$  do
5     if CollisionFree( $v, u$ ) then
6        $E \leftarrow E \cup \{(v, u), (u, v)\}$ 
6 return  $G = (V, E);$ 
  
```

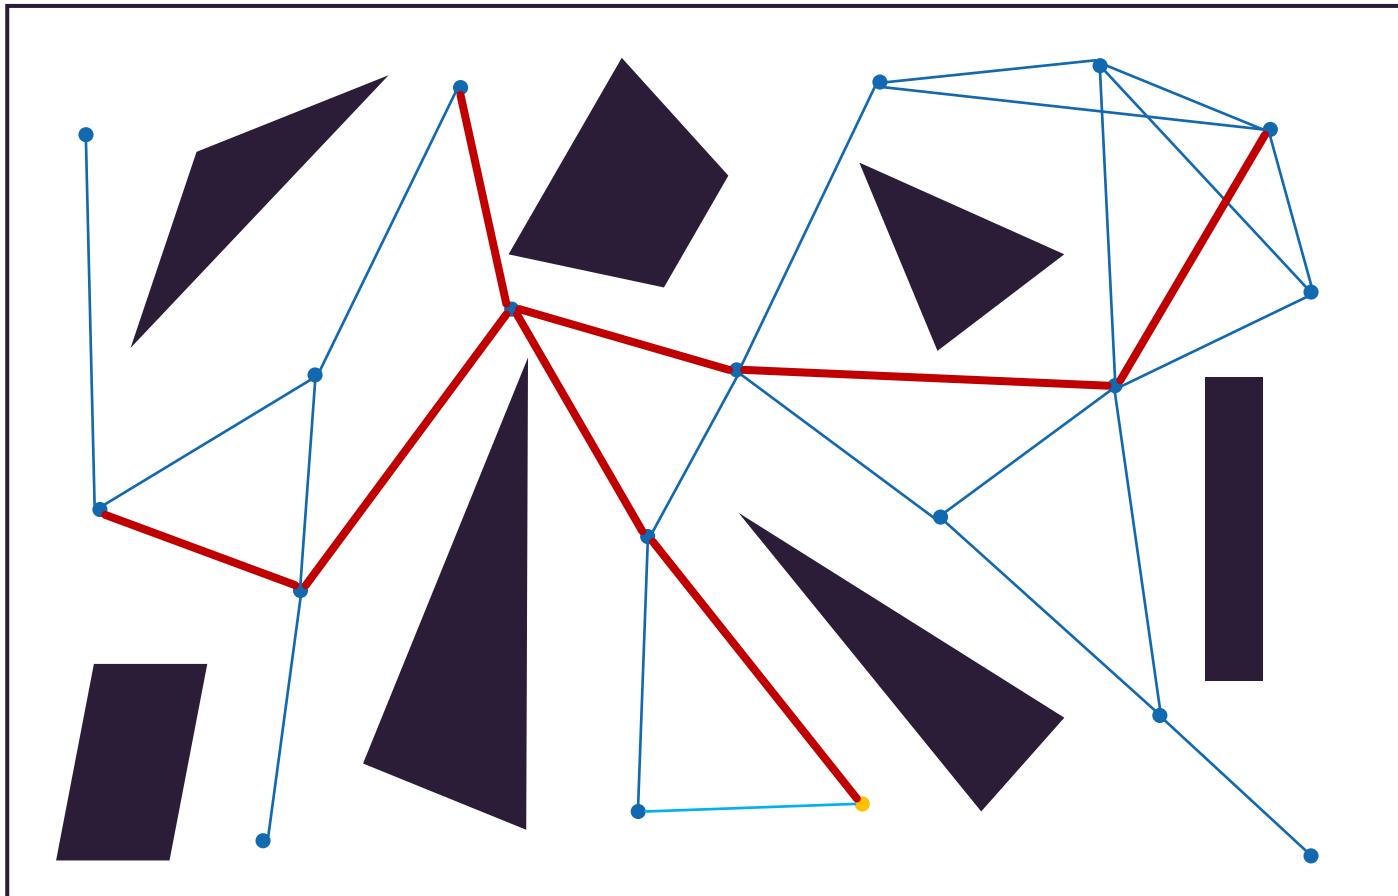


S. Karaman and E. Frazzoli, "Sampling-based algorithms for optimal motion planning." *The International Journal of Robotics Research* 30.7 (2011): 846-894.

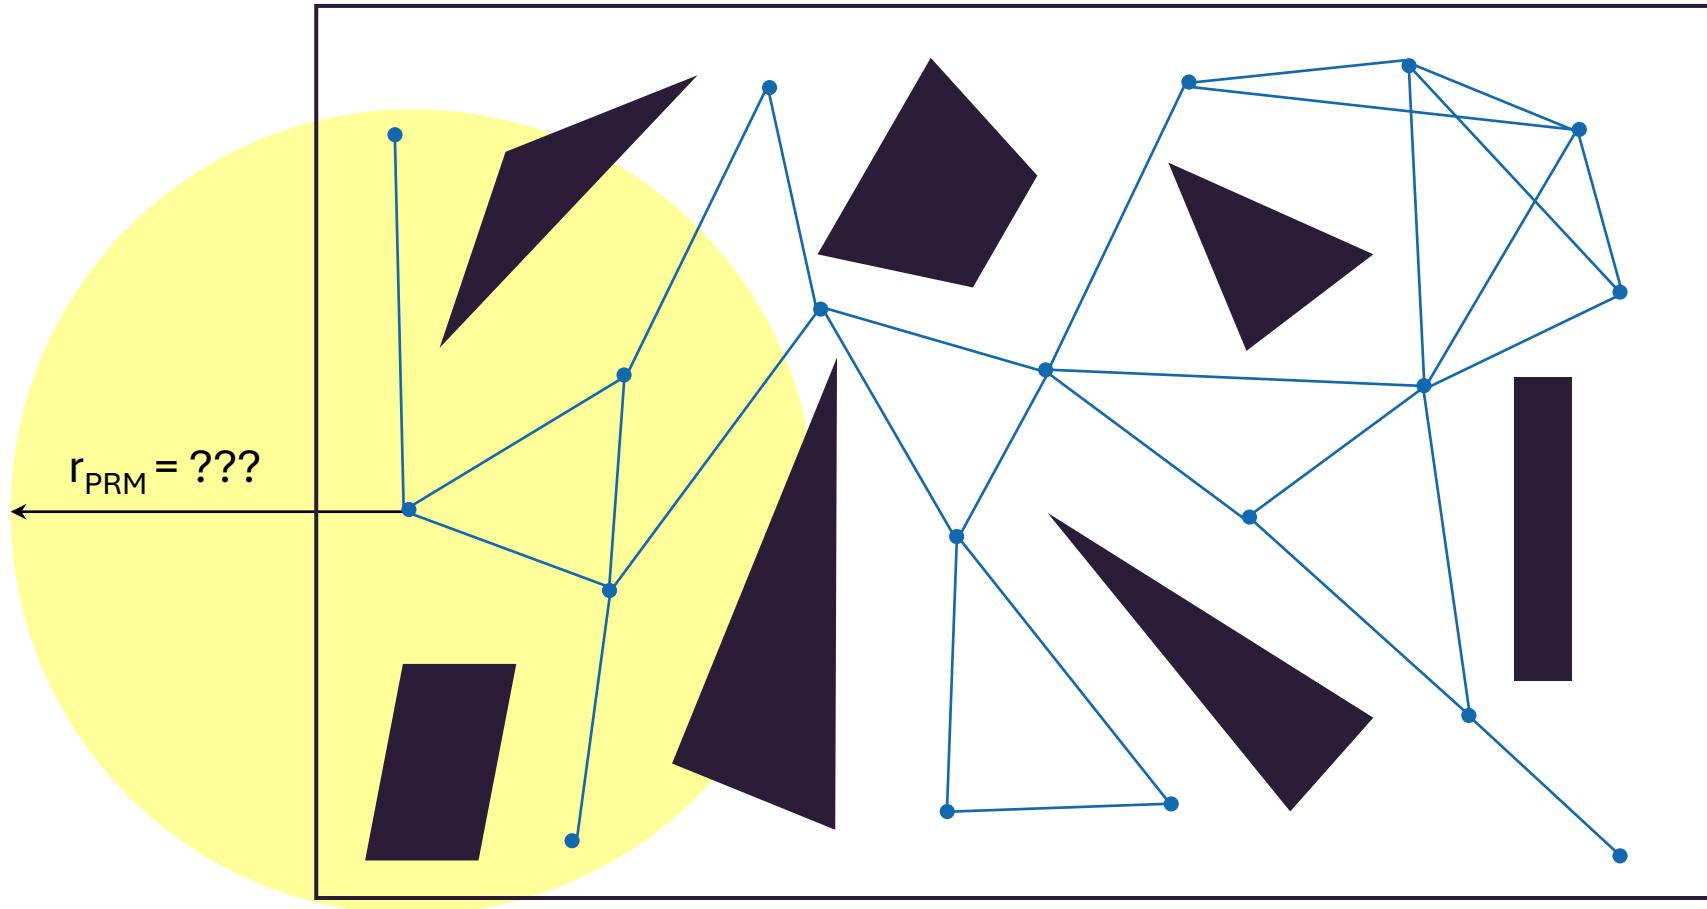
Example: PRM



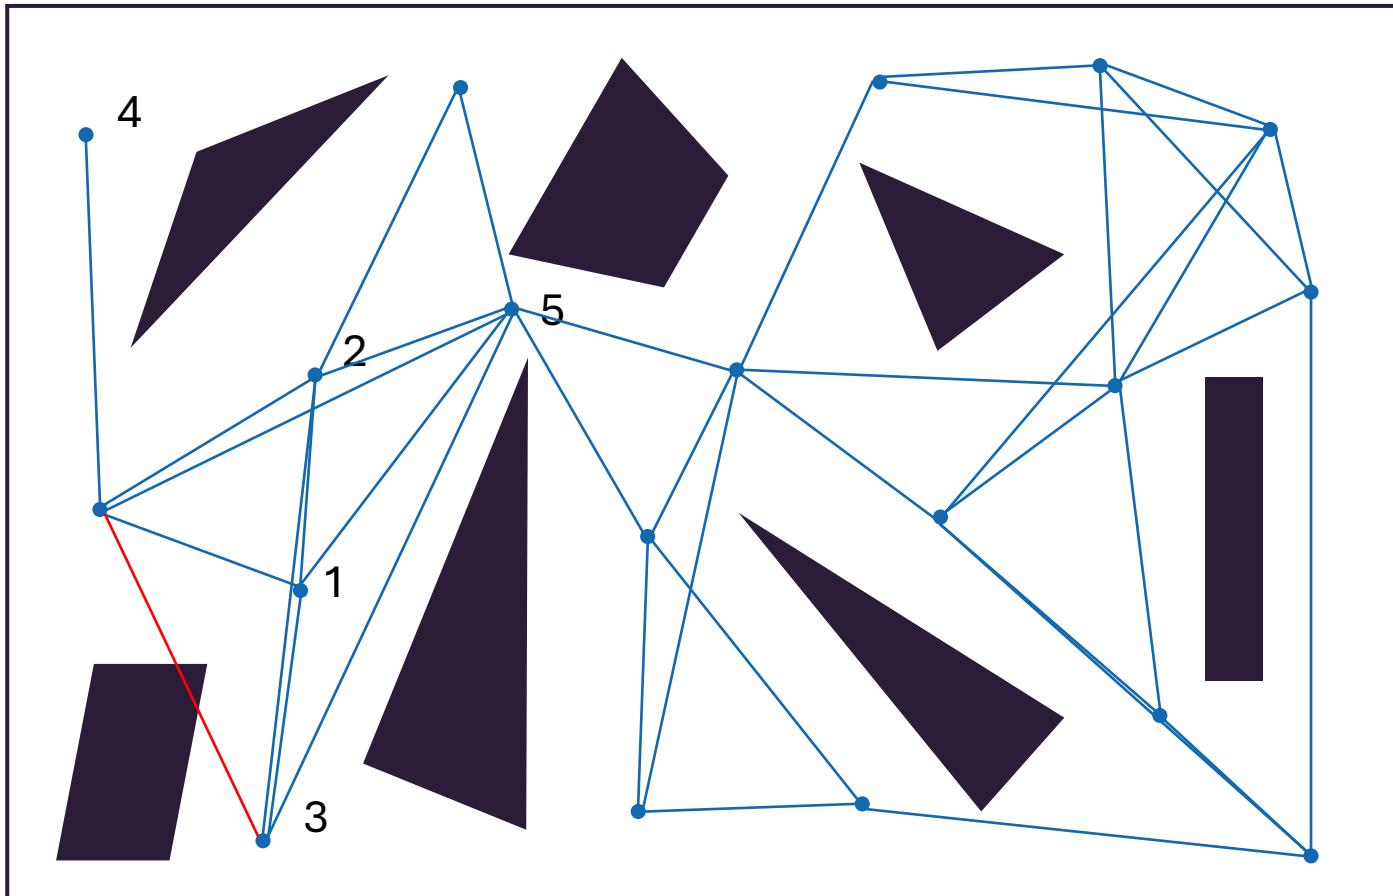
Example: PRM



PRM: How to Connect?



PRM: How to Connect?



Connect to k nearest
neighbours

PRM Summary

- Graph constructed by connecting random configurations
- Designed for **multiple queries**
- Underlying assumption is that it's worth performing **substantial pre-computation** on a given environment to enable multiple planning queries
- However, often only interested in an **efficient single query** without pre-computations
 - E.g. in dynamic environments, different robot models
- Idea: combine **exploration and search** into a single method
 - Motivates tree-based planning approaches

Rapidly Exploring Random Trees

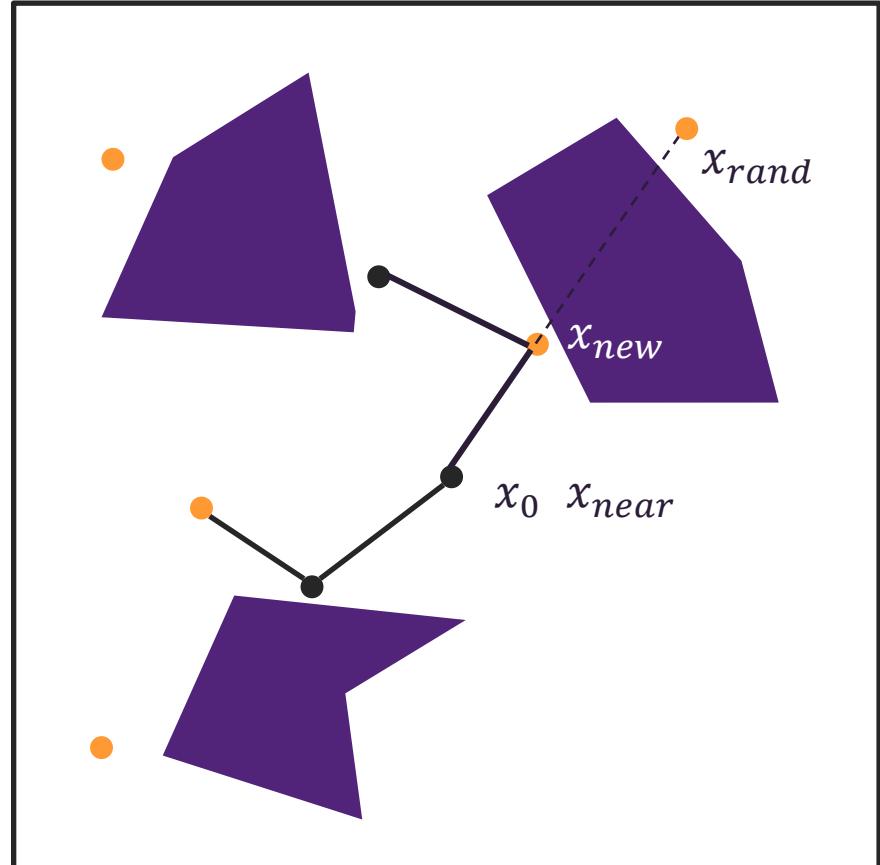
- Instead of sampling all our target states at once, start from initial configuration and **incrementally** grow outwards from there, forming a tree of connected, dynamically feasible local paths
- Terminate when we reach a configuration in a goal **region**
- For **single-query searches**

RRTs

Algorithm 1 RRT (space filling) - Adapted from [1]

Require: Initial pose x_0 ,

```
1:  $V \leftarrow \{x_0\}$ ,  $E \leftarrow \emptyset$ 
2: for  $i = 1, \dots, n$  do
3:    $x_{rand} \leftarrow \text{SAMPLEFREESPACE}()$ 
4:    $x_{near} \leftarrow \text{FINDNEAREST}(V, x_{rand})$ 
5:    $x_{new} \leftarrow \text{STEER}(x_{rand}, x_{near})$ 
6:   if OBSTACLEFREE( $x_{near}, x_{new}$ ) then
7:      $V \leftarrow V \cup \{x_{new}\}$ 
8:      $E \leftarrow E \cup \{(x_{near}, x_{new})\}$ 
9: return  $\mathcal{G} = (V, E)$ 
```



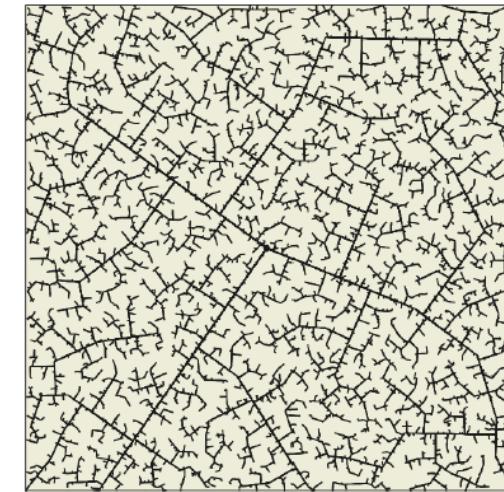
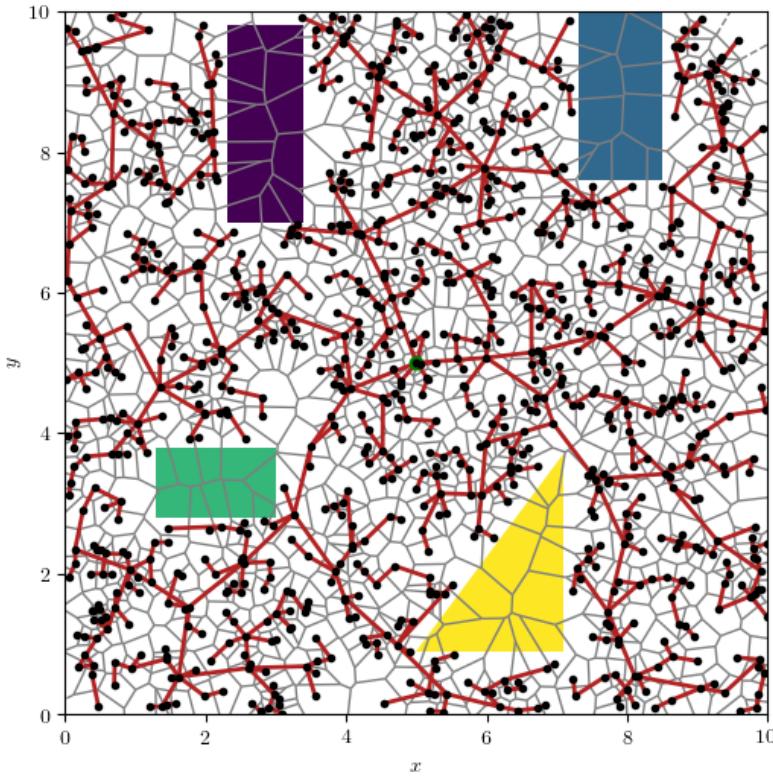
S. M. LaValle, "Rapidly-exploring random trees: A new tool for path planning." *Technical Report. Computer Science Department, Iowa State University (TR 98-11)*. (1998)

S. Karaman and E. Frazzoli, "Sampling-based algorithms for optimal motion planning." *The International Journal of Robotics Research* 30.7 (2011): 846-894.

RRTs

- Incrementally explore the C-space

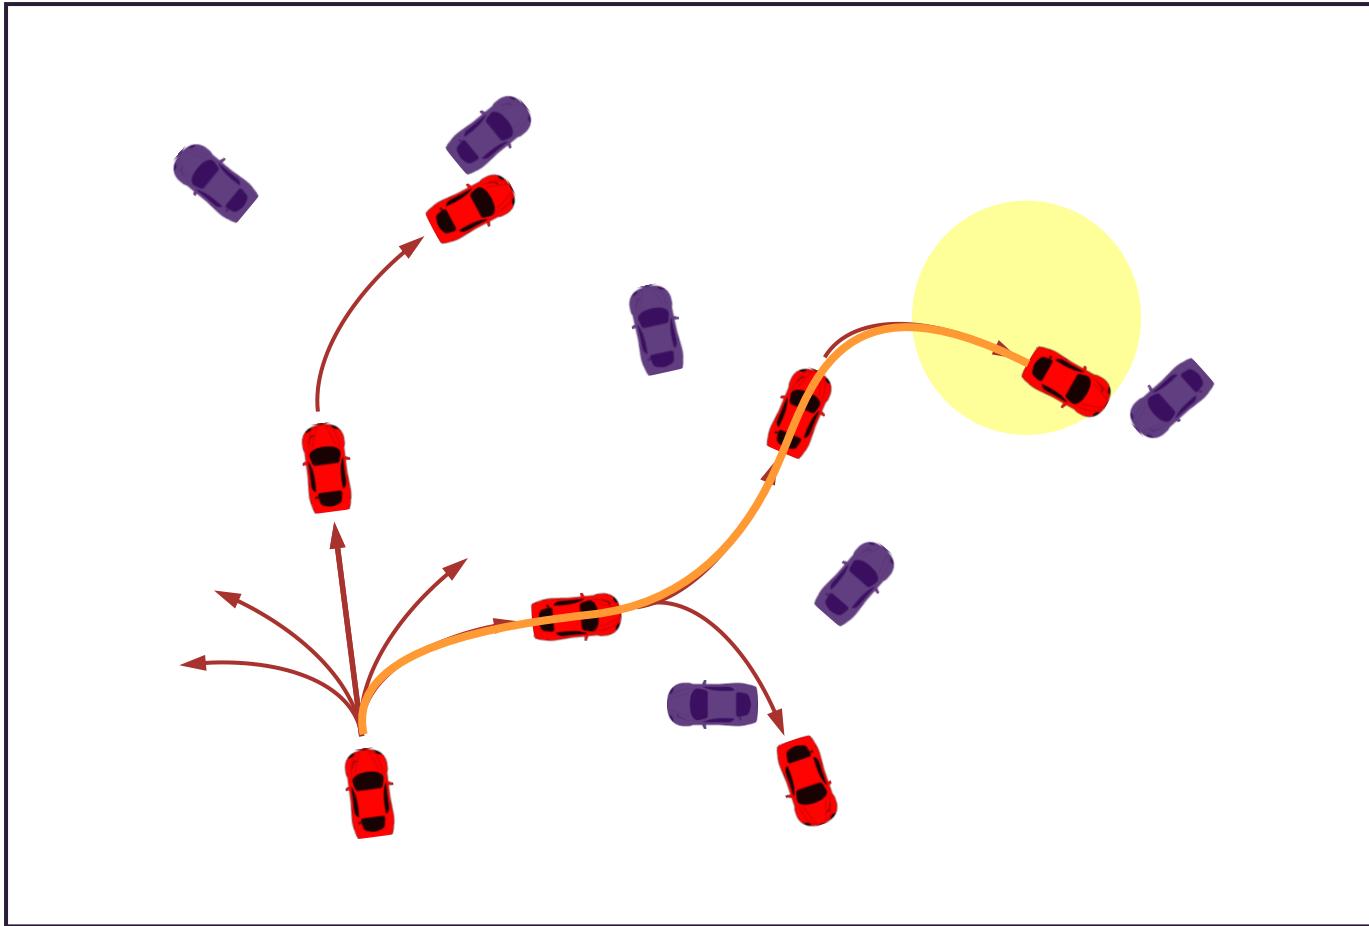
- In early iterations, the RRT quickly reaches unexplored parts (Voronoi diagram)
- However, RRT is dense in the limit → will eventually reach all points in the space



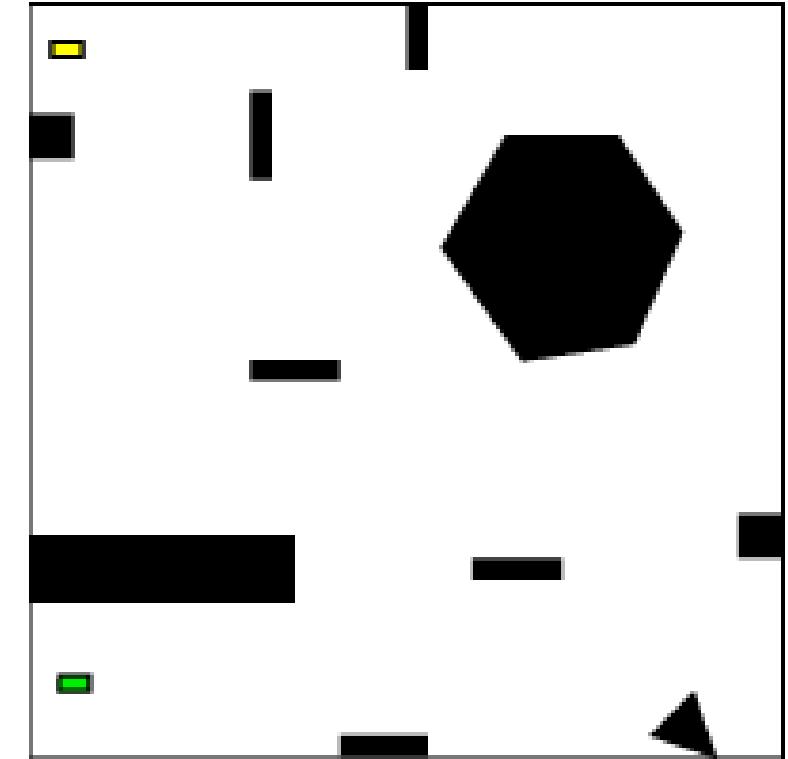
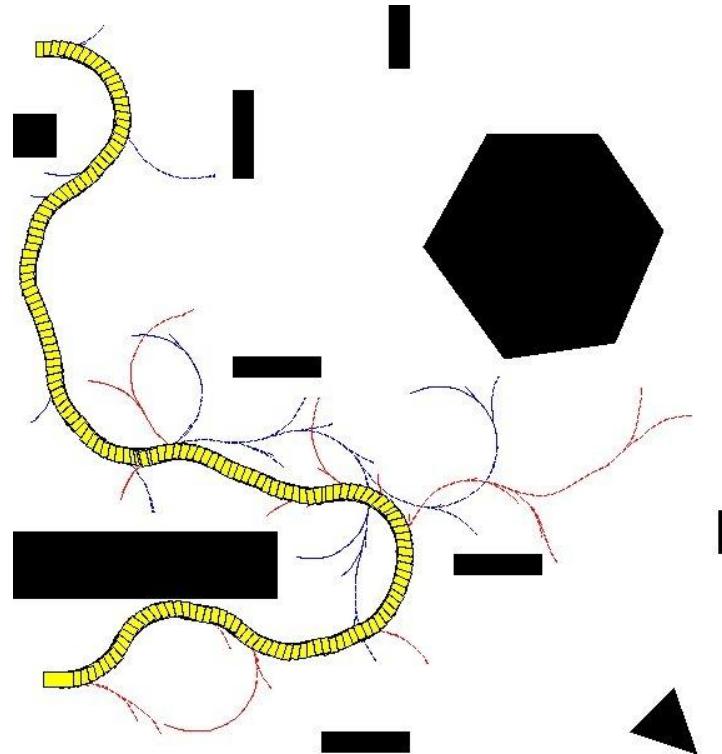
Kinodynamic Planning

- Specifically designed to handle non-holonomic constraints and dynamic constraints (kinodynamic planning)
- Rather than solve for control inputs to exactly connect two points in C-space, just ‘steer’ in the right direction
 - Sample controls to generate **feasible** local paths and reach **approximate, intermediate** configuration that gets ‘close’ to the target

RRT Steering Function



RRT Examples



Steve LaValle (<http://msl.cs.uiuc.edu/rrt/gallery.html>)

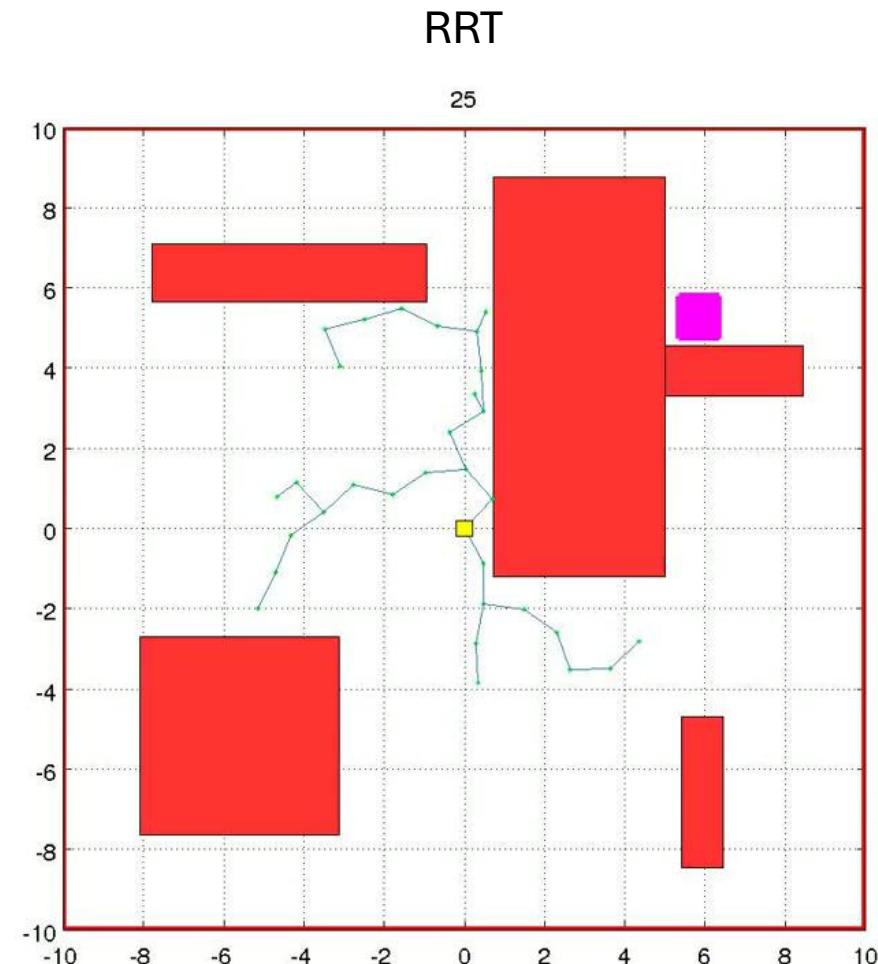
Theoretical Guarantees

- Early on, PRM (and RRT) were known to be probabilistically complete, i.e.

“If there exists a solution, it will be found with probability $\rightarrow 1$ as the number of samples $\rightarrow \infty$ ”

- Will the RRT find the shortest path?

NO! They almost surely converge to a non-optimal solution



<https://www.youtube.com/watch?v=FAFw8DoKvik>

A Star Rises

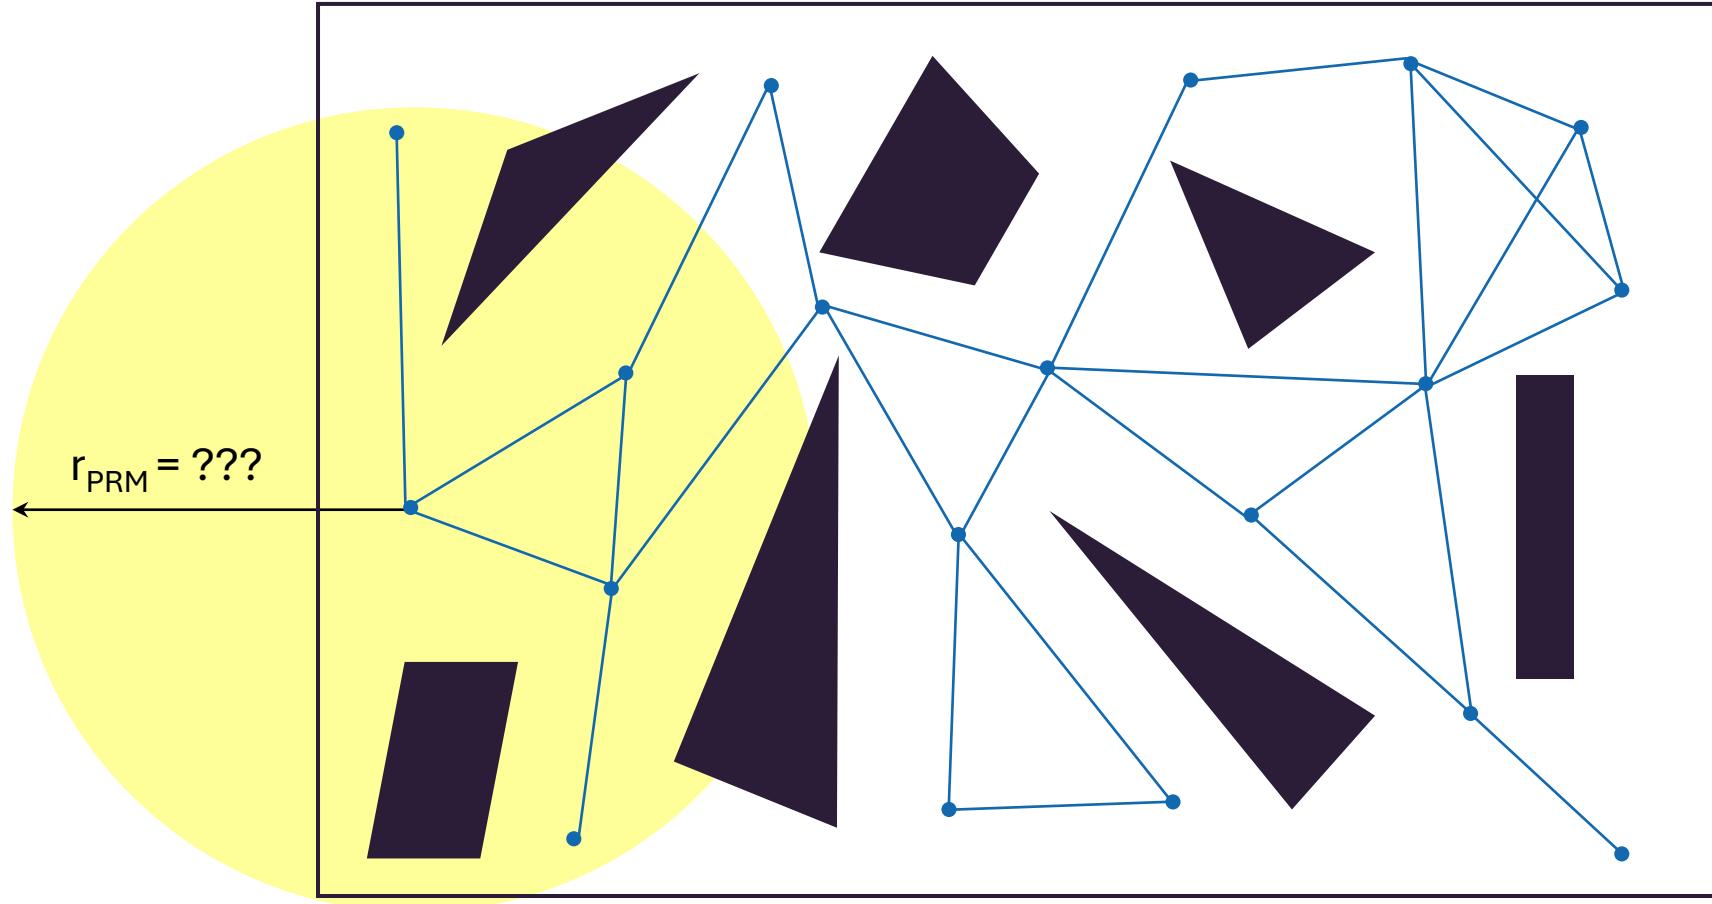
- Sertac Karaman and Emilio Frazzoli worked on incremental versions of PRM and RRT, and developed PRM* and RRT*, **asymptotically optimal** versions of PRM and RRT

“The cost of the returned solution will approach the optimal as the number of samples $\rightarrow \infty$ ”

- ‘Star’ versions use rewiring technique (RRT*) and formal proof to select connection radius (PRM*) as function of sample density

S. Karaman and E. Frazzoli, "Sampling-based algorithms for optimal motion planning." *The International Journal of Robotics Research* 30.7 (2011): 846-894.

PRM*



$$r_{PRM}^* = \gamma_{PRM} \left(\frac{\log n}{n} \right)^{\frac{1}{d}}$$

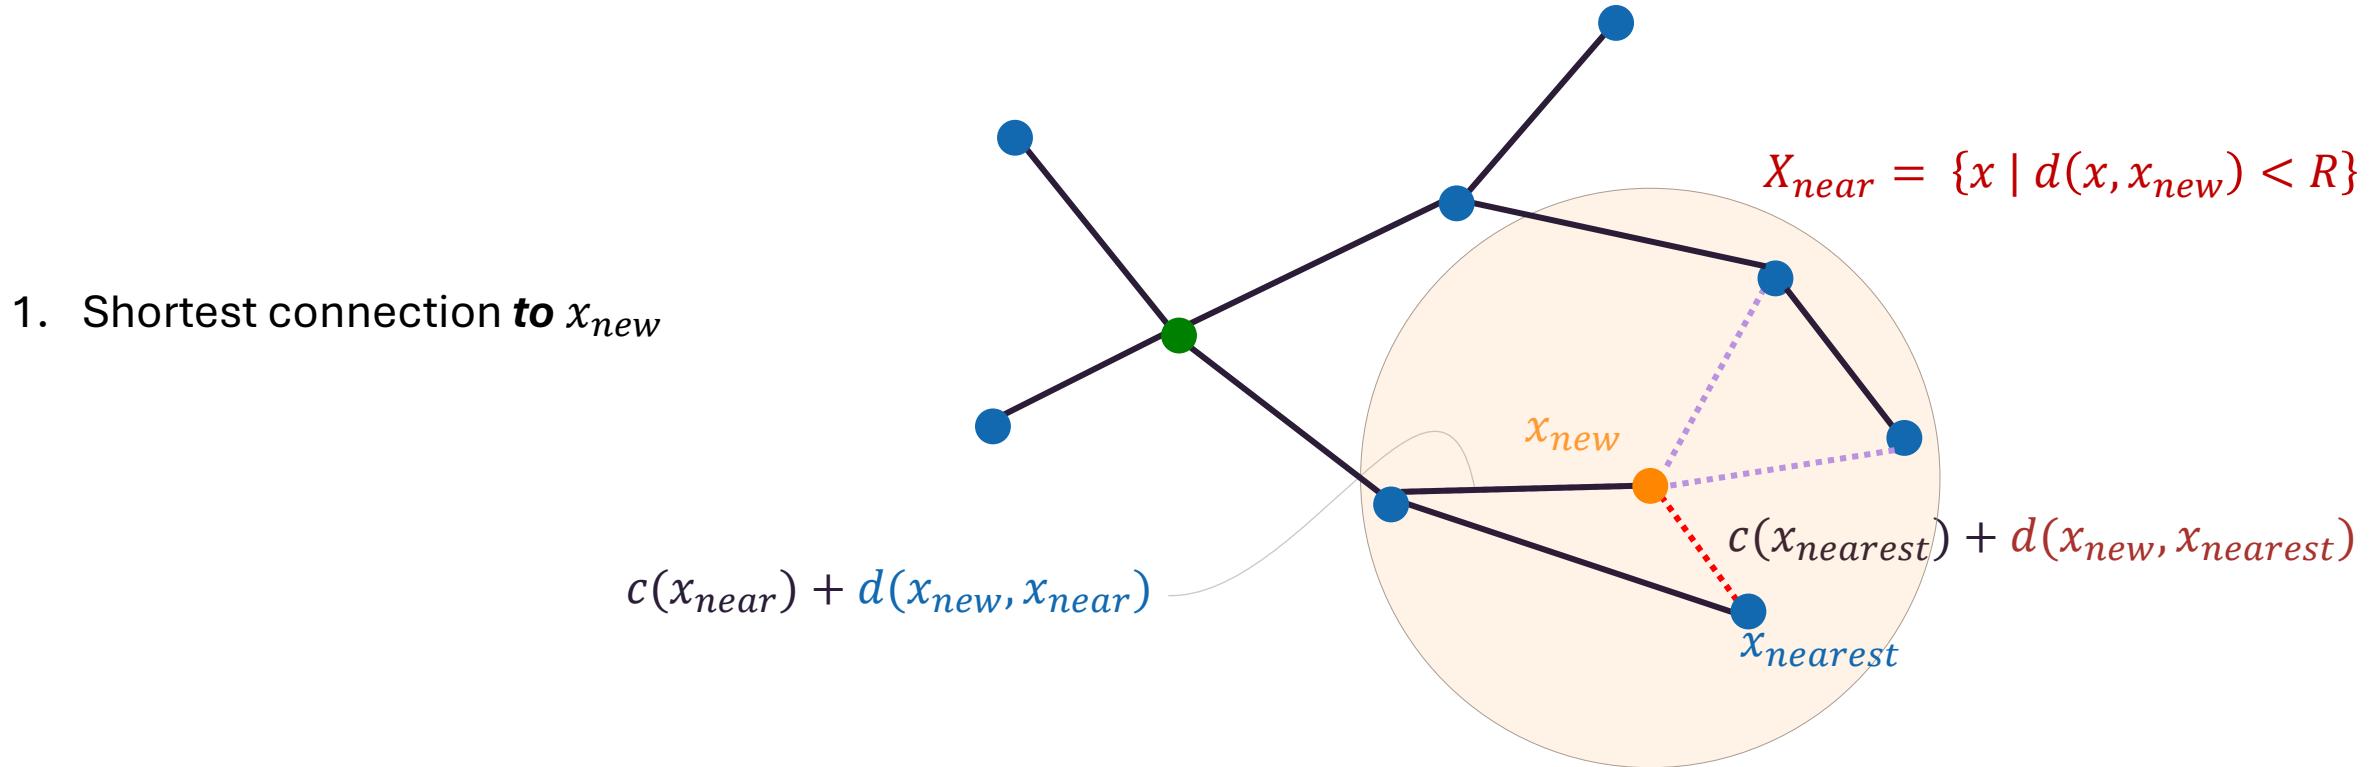
number of samples problem dimensionality

$$\gamma_{PRM} > 2 \left(1 + \frac{1}{d} \right)^{\frac{1}{d}} \left(\frac{\mu(X_{free})}{\zeta_d} \right)^{\frac{1}{d}}$$

total free space area of unit ball

■ Two key features

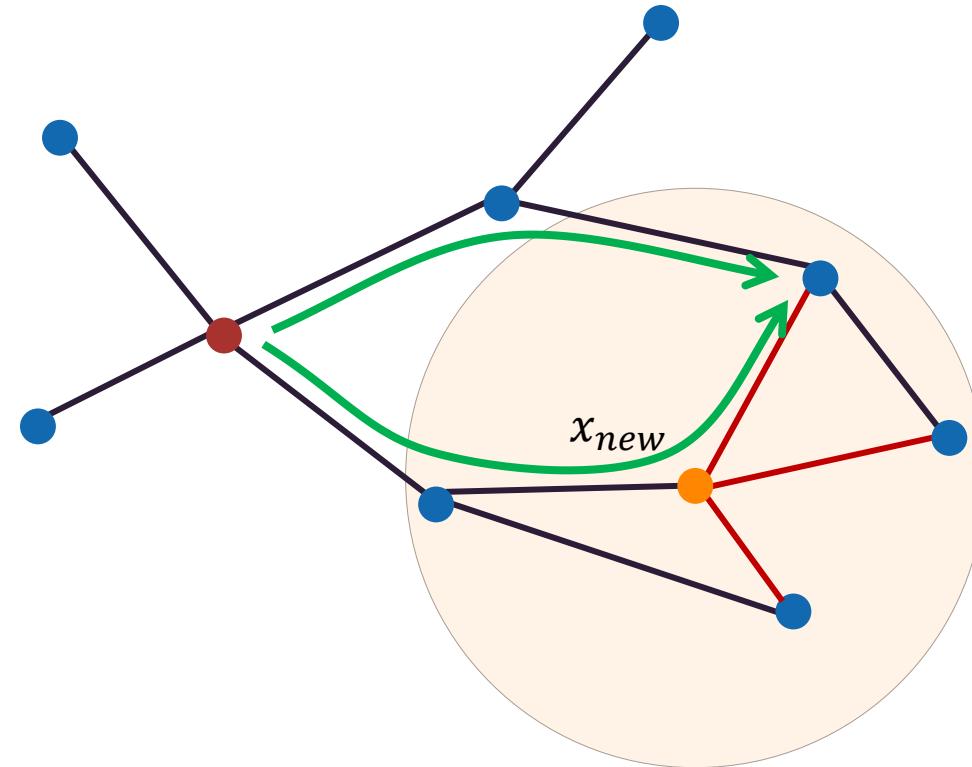
- **Near** operations find the set of nodes that might be affected by a new node addition
- **Rewiring** rebuilds the tree within this set to maintain the minimum cost of arrival tree



- Two key features

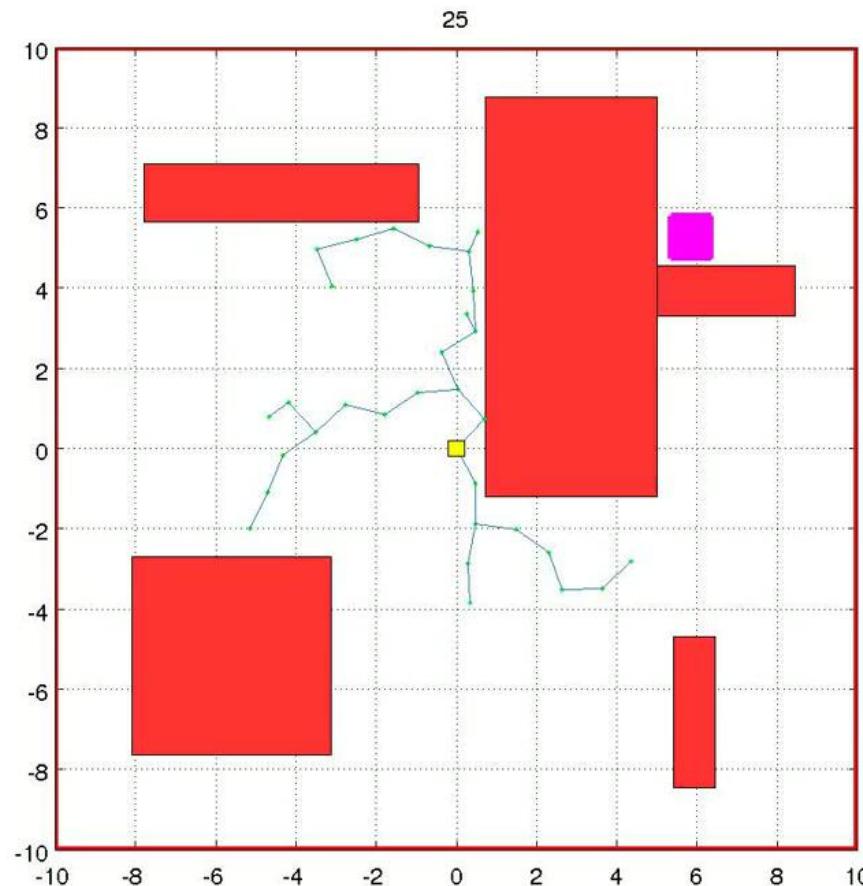
- **Near** operations find the set of nodes that might be affected by a new node addition
- **Rewiring** rebuilds the tree within this set to maintain the minimum cost of arrival tree

1. Shortest connection **to** x_{new}
2. Shortest connections **from** x_{new}



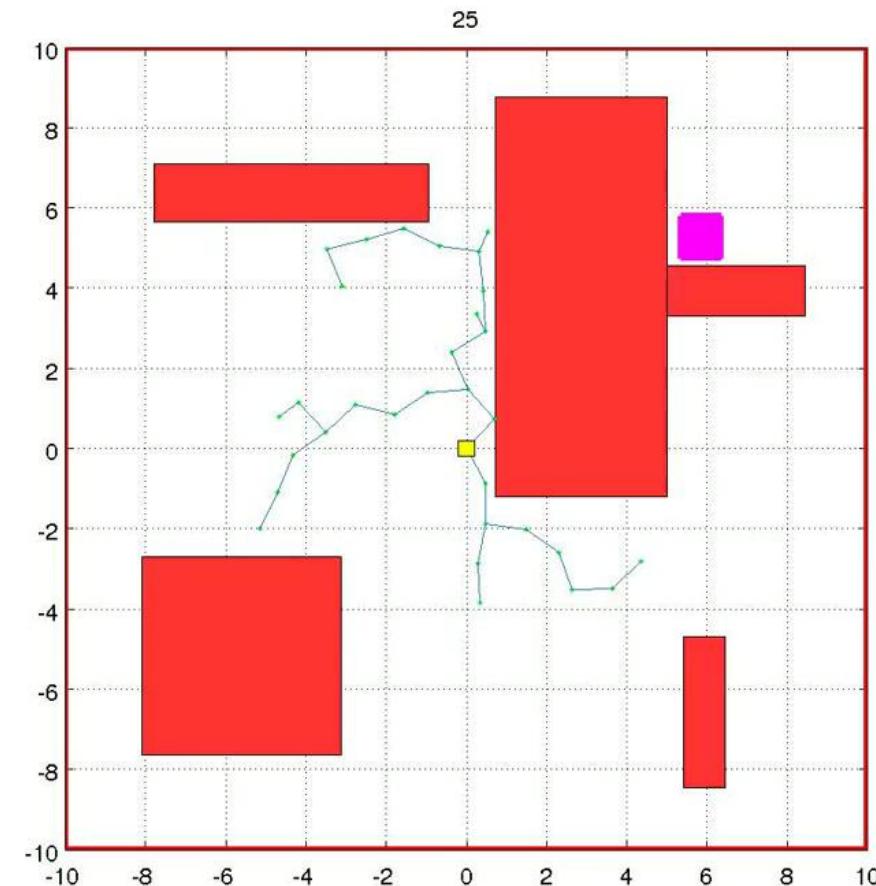
RRT vs RRT*

RRT



<https://www.youtube.com/watch?v=FAFw8DoKvik>

RRT*



<https://www.youtube.com/watch?v=YKiQTJpPFkA>

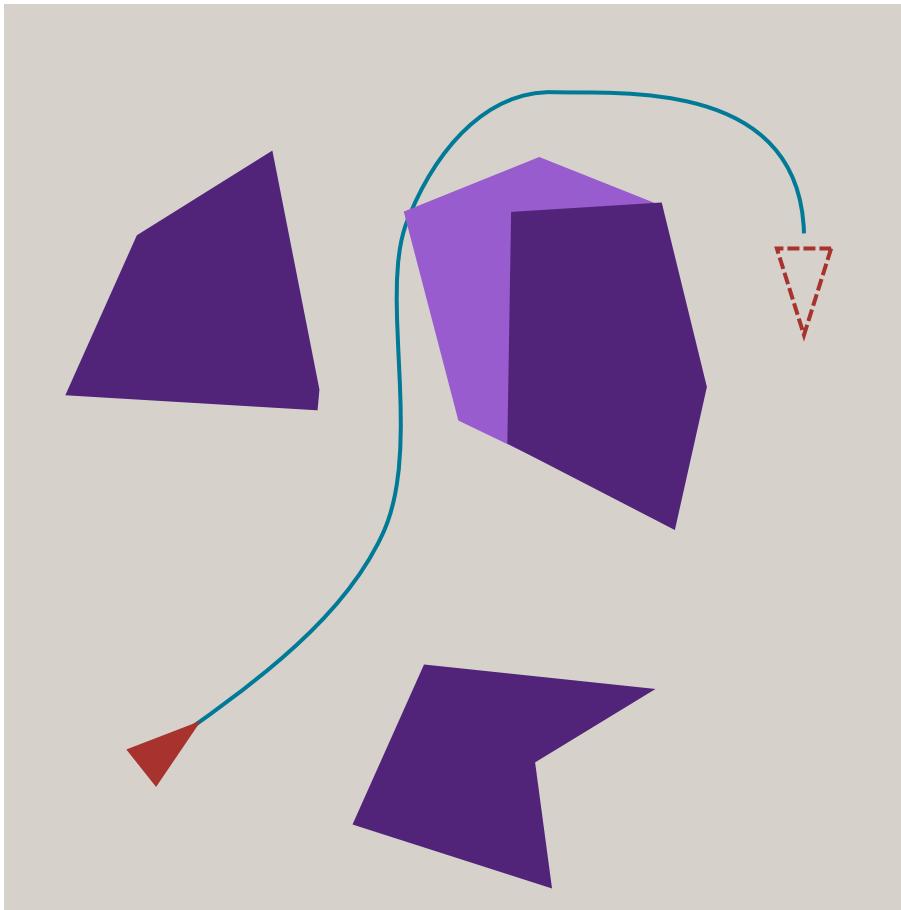
The Many MANY Variants of PRM and RRT

- The popularity of sampling-based planners has led to a huge variety of probabilistic sampling-based planning methods
- Most focus of finding more optimal solutions faster, sometimes at the cost of formal guarantees
- Examples include:
 - Searching from start and goal and meshing the two trees
 - Limiting the search space once a plan has been found
 - Varying sampling density based on domain-knowledge
 - Computational improvements (collision checks, bootstrapping search, efficient data structures)

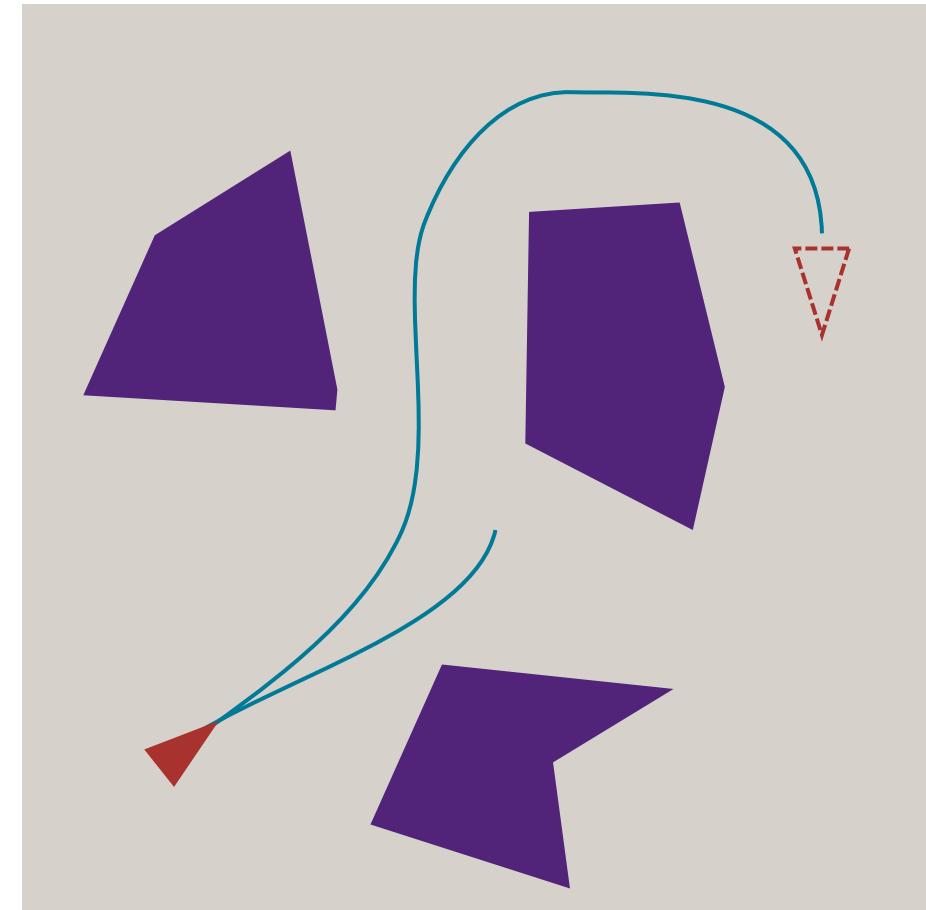
Planning Hierarchies

Sources of Uncertainty

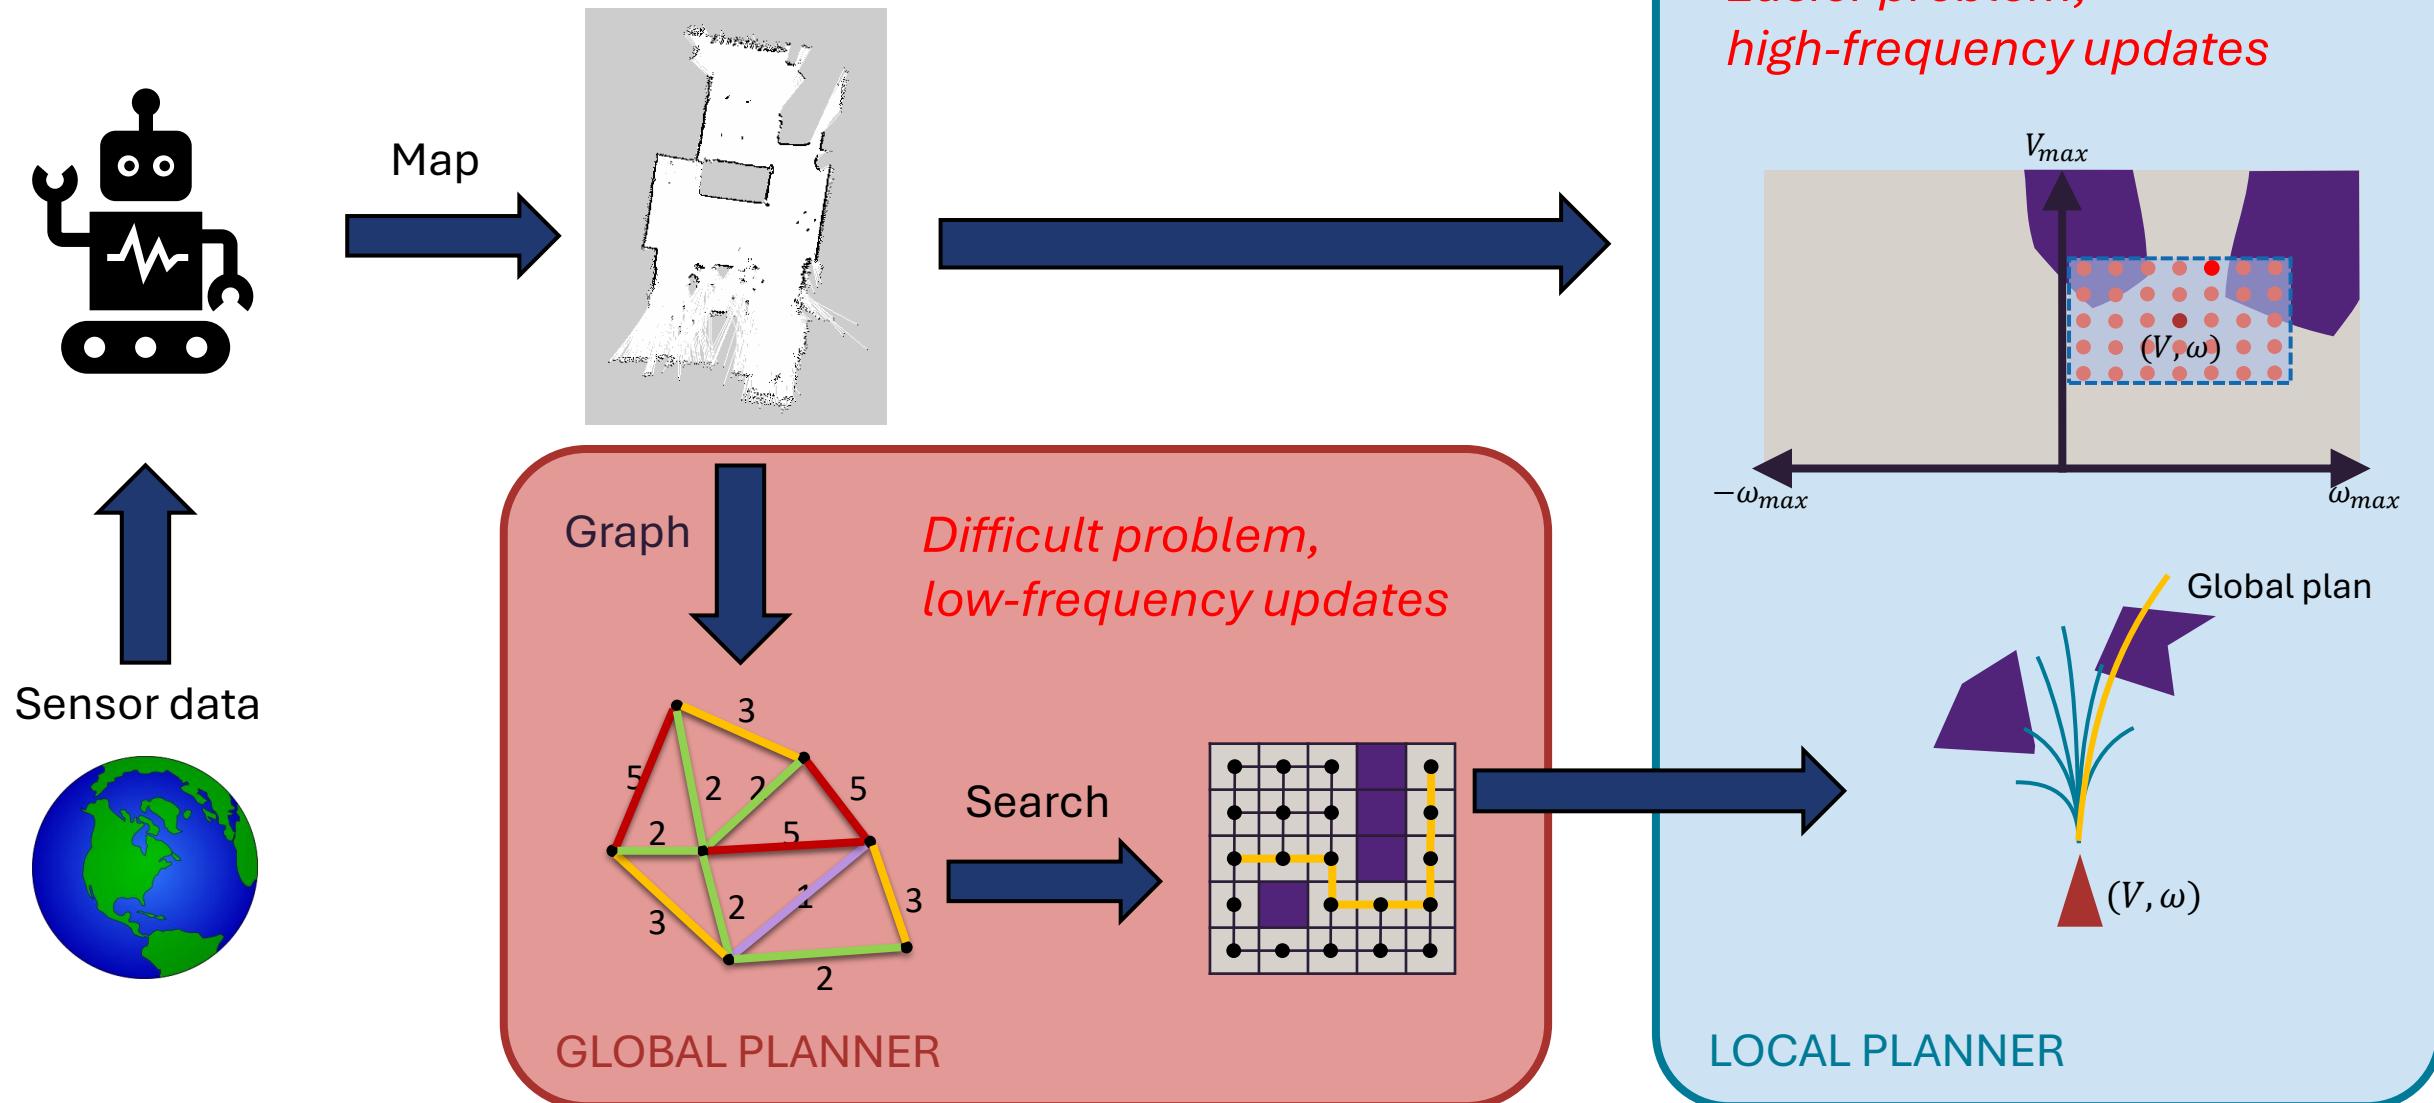
Environment uncertainty



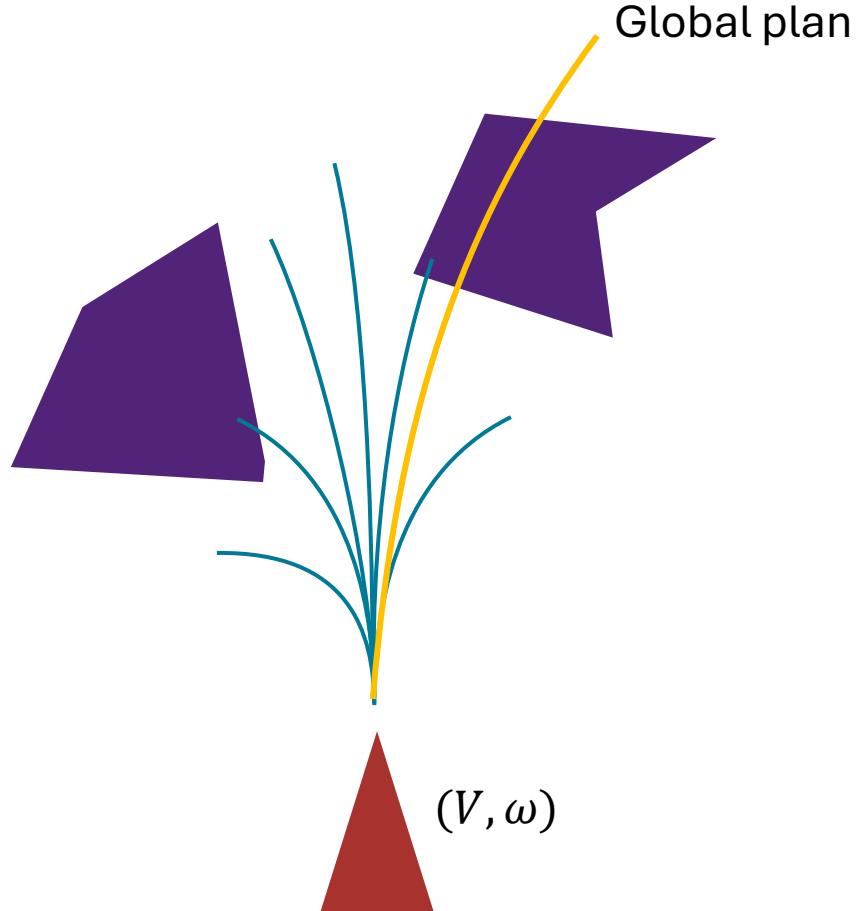
Motion uncertainty



Planning Hierarchy



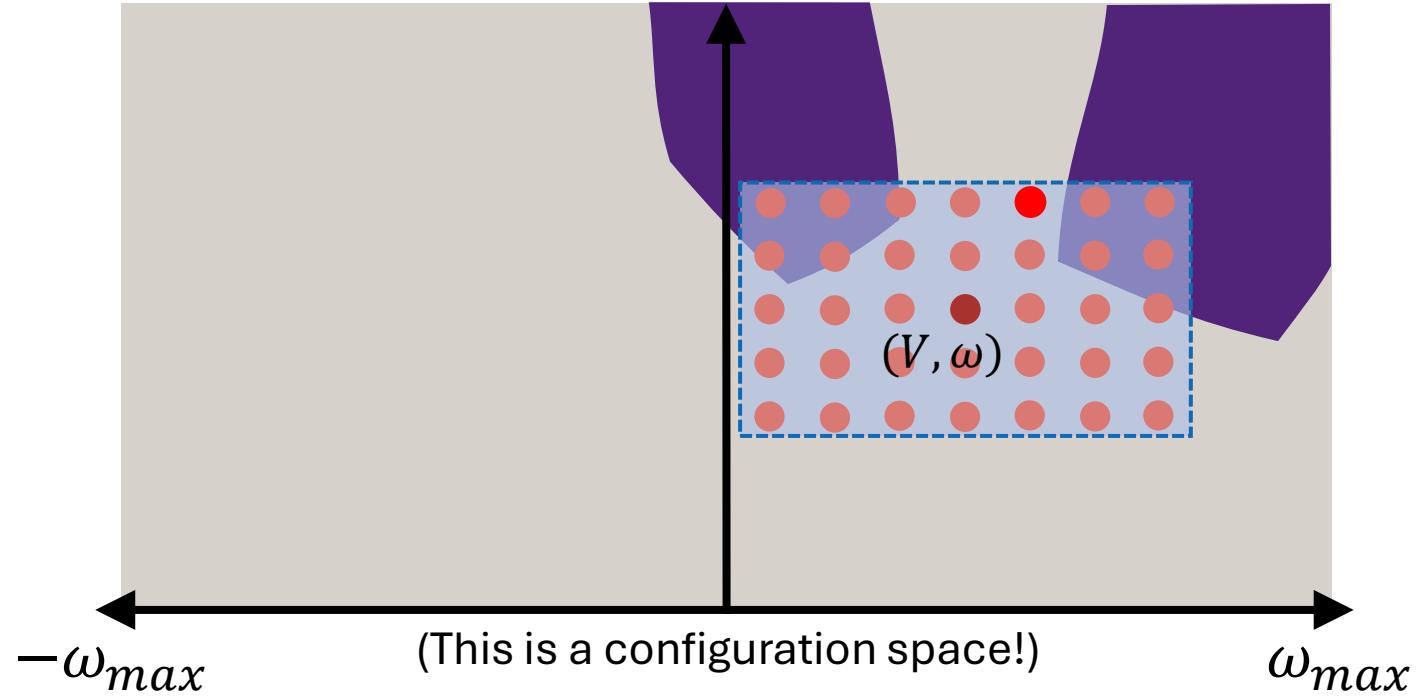
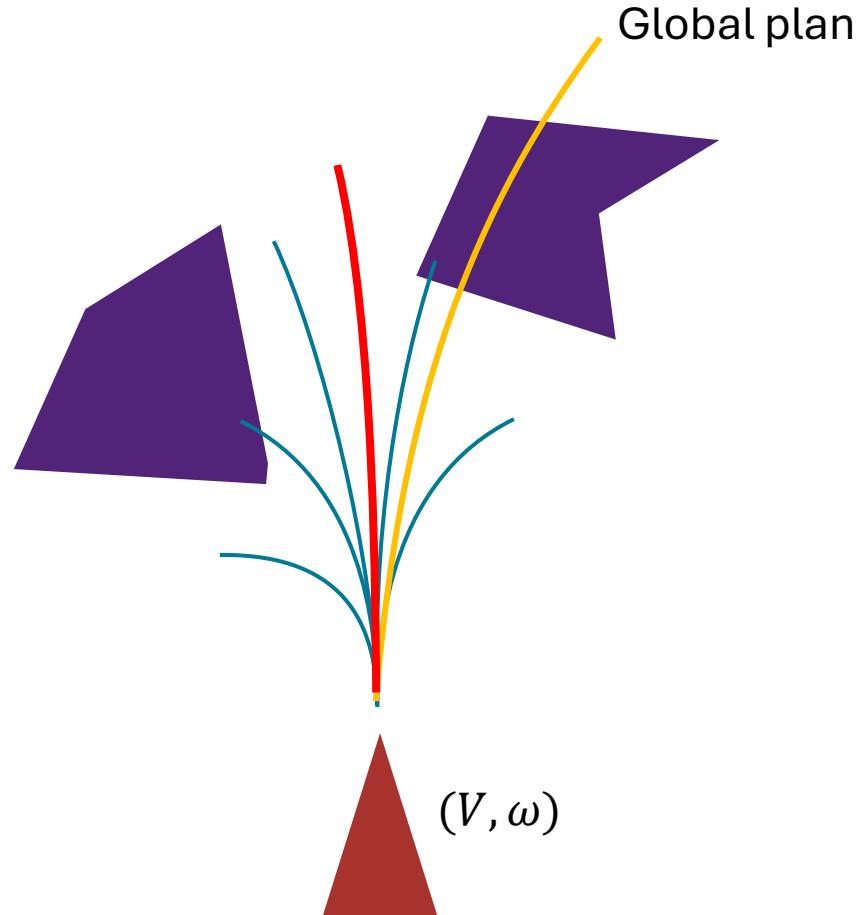
Dynamic Window Collision Avoidance



- Generating full time-varying trajectories for $V(t)$ and $\omega(t)$ is still very challenging
- If we assume (V, ω) are constant for a fixed Δt , each local path in the future is a **circular arc segment**
- This can be easily considered as a type of *velocity configuration space*

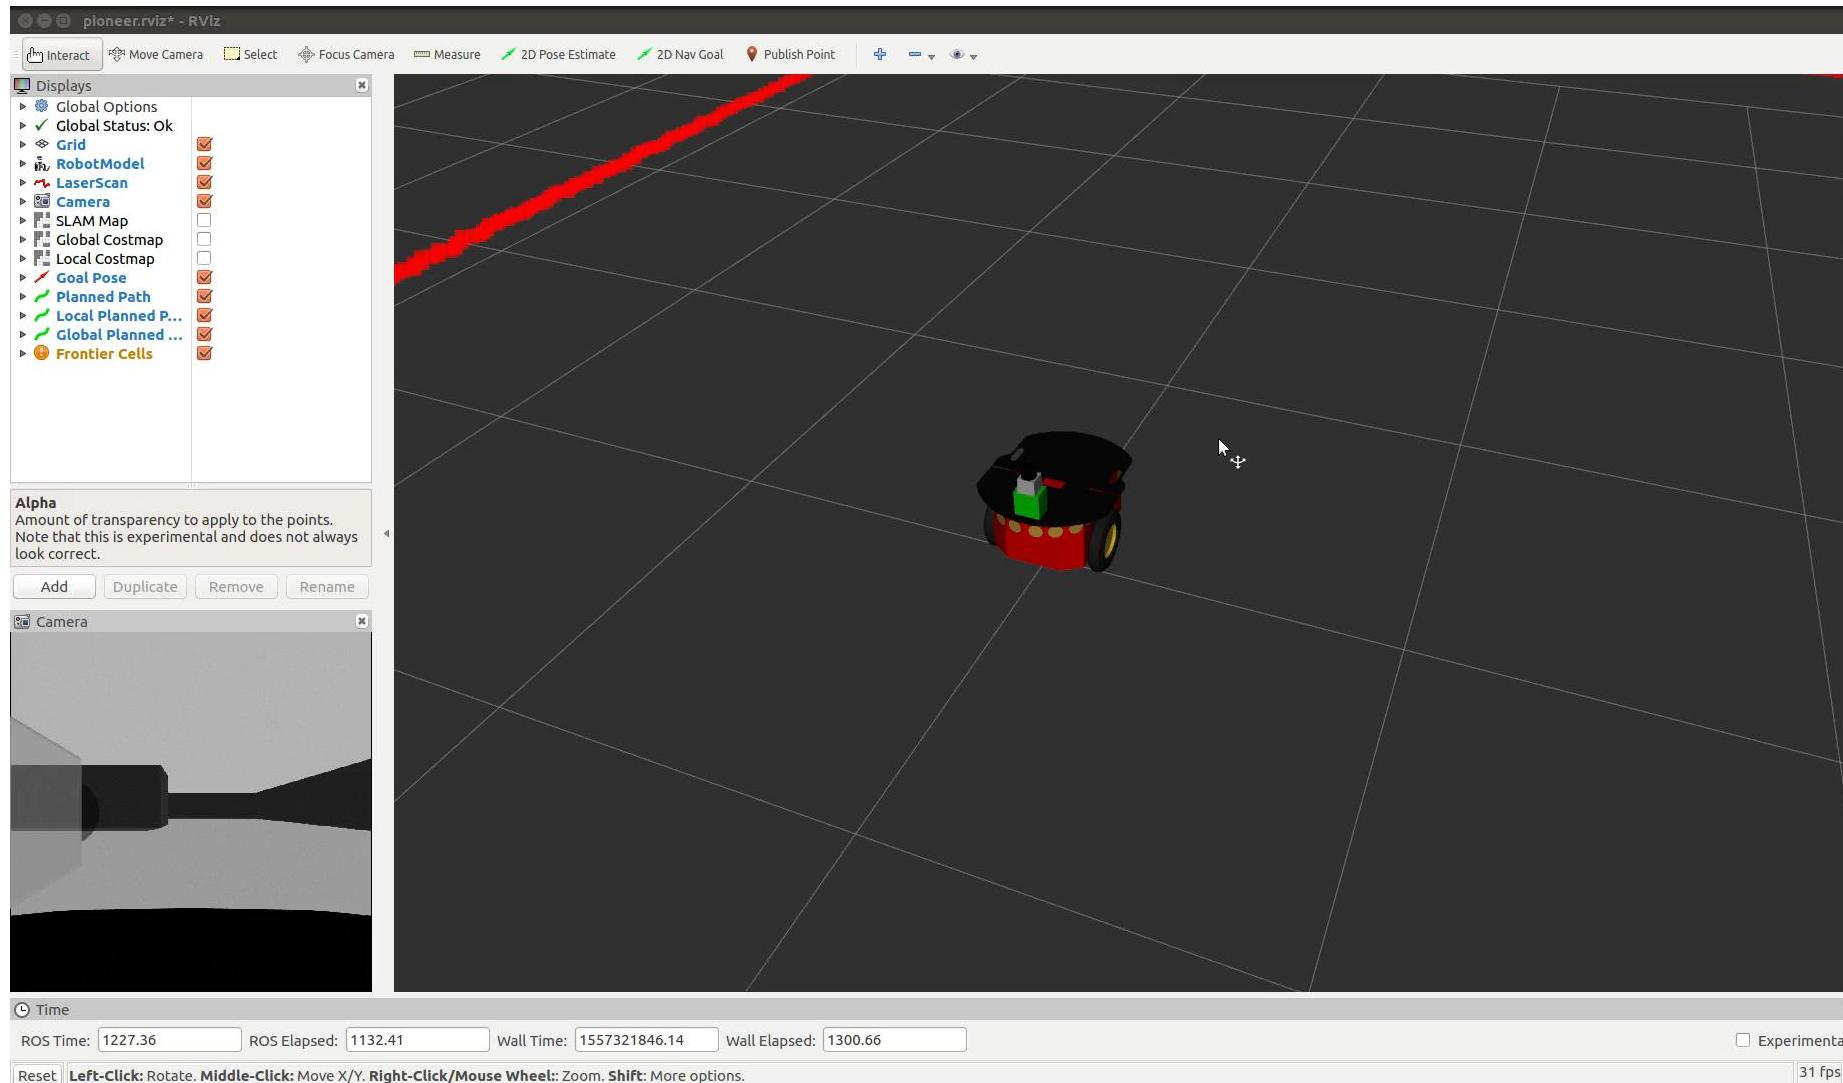
Fox, D., Burgard, W., & Thrun, S. (1997). The dynamic window approach to collision avoidance. *IEEE Robotics & Automation Magazine*, 4(1), 23-33.

Dynamic Window Collision Avoidance



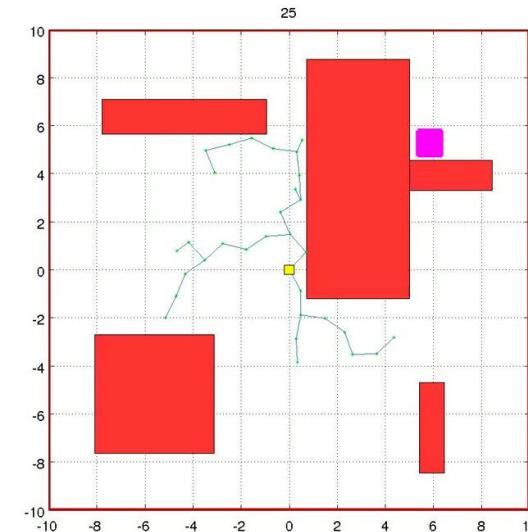
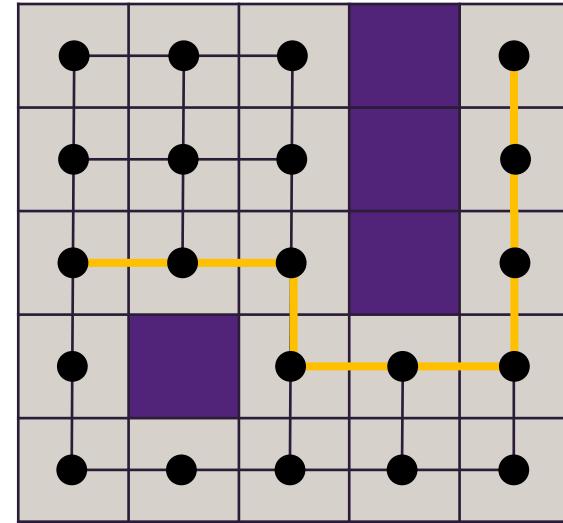
Maximise utility metric (usually maximise speed, minimise distance to goal, maximise distance from obstacles) across configuration samples

ROS Navigation Stack (A* global, DWA local)



Summary: Motion Planning Essentials

- Motion planning representations:
 - Workspace vs. configuration space
- Graphs to discretise the planning space
 - Graph search methods e.g. Dijkstra's algorithm (non-uniform edge weights), A* (cost-to-go heuristic for faster search)
- Sampling-based planning
 - PRM/PRM*, RRT/RRT*
- Hierarchies for global and local planning





AuSRoS

Australian School of Robotic Systems

B1: Motion Planning Essentials

A/Prof. Jen Jen Chung

With slides adapted from Dr Nicholas Lawrance

